

GoDaddy - Microbusiness Density Forecasting

Mohammad Solki

2023-02-28

1. Introduction

1.1. Goal of the Competition

The challenge in this competition is to forecast microbusiness activity across the United States, as measured by the density of microbusinesses in US counties. Microbusinesses are often too small or too new to show up in traditional economic data sources, but microbusiness activity may be correlated with other economic indicators of general interest.

This work will help policymakers gain visibility into microbusinesses, a growing trend of very small entities. Additional information will enable new policies and programs to improve the success and impact of these smallest of businesses.

GoDaddy's Venture Forward team has gathered data on over 20 million microbusinesses in the United States, defined as businesses with an online presence and ten or fewer employees, to help policymakers understand the factors associated with these small businesses. While traditional economic data sources often miss these businesses, GoDaddy's survey data can provide insights into this sector of the economy. The data can be used to improve predictions and inform decision-making to create more inclusive and resilient economies. The competition hosted by GoDaddy aims to empower entrepreneurs by giving them the tools they need to grow online and make a substantial impact on communities across the country.

Model accuracy will be evaluated on SMAPE (Symmetric mean absolute percentage error) between forecasts and actual values. We define SMAPE = 0 when the actual and predicted values are both 0.

SMAPE formula is usually defined as follows:

$$\text{SMAPE} = \frac{100\%}{n} \sum_{t=1}^n \frac{|F_t - A_t|}{(|F_t| + |A_t|)/2}$$

where:

- n is the number of observations in the time series
- F_t is the forecasted value at time t
- A_t is the actual value at time t
- $|x|$ denotes the absolute value of x .

1.2. Datasets

A great deal of data is publicly available about counties and we have not attempted to gather it all here. You are strongly encouraged to use external data sources for features.

train.csv

- `row_id` An ID code for the row.
- `cfips` A unique identifier for each county using the Federal Information Processing System. The first two digits correspond to the state FIPS code, while the following 3 represent the county.
- `county_name` The written name of the county.
- `state_name` The name of the state.
- `first_day_of_month` The date of the first day of the month.
- `microbusiness_density` Microbusinesses per 100 people over the age of 18 in the given county. This is the target variable. The population figures used to calculate the density are on a two-year lag due to the pace of update provided by the U.S. Census Bureau, which provides the underlying population data annually. 2021 density figures are calculated using 2019 population figures, etc.
- `active` The raw count of microbusinesses in the county. Not provided for the test set.

test.csv Metadata for the submission rows. This file will remain unchanged throughout the competition.

- `row_id` An ID code for the row.
- `cfips` A unique identifier for each county using the Federal Information Processing System. The first two digits correspond to the state FIPS code, while the following 3 represent the county.
- `first_day_of_month` The date of the first day of the month.

census_starter.csv Examples of useful columns from the Census Bureau's American Community Survey (ACS) at data.census.gov (<https://data.census.gov/>). The percentage fields were derived from the raw counts provided by the ACS. All fields have a two year lag to match what information was available at the time a given microbusiness data update was published.

- `pct_bb_[year]` The percentage of households in the county with access to broadband of any type. Derived from ACS table B28002: PRESENCE AND TYPES OF INTERNET SUBSCRIPTIONS IN HOUSEHOLD.
- `cfips` The CFIPS code.
- `pct_college_[year]` The percent of the population in the county over age 25 with a 4-year college degree. Derived from ACS table S1501: EDUCATIONAL ATTAINMENT.
- `pct_foreign_born_[year]` The percent of the population in the county born outside of the United States. Derived from ACS table DP02: SELECTED SOCIAL CHARACTERISTICS IN THE UNITED STATES.
- `pct_it_workers_[year]` The percent of the workforce in the county employed in information related industries. Derived from ACS table S2405: INDUSTRY BY OCCUPATION FOR THE CIVILIAN EMPLOYED POPULATION 16 YEARS AND OVER.
- `median_hh_inc_[year]` The median household income in the county. Derived from ACS table S1901: INCOME IN THE PAST 12 MONTHS (IN 2021 INFLATION-ADJUSTED DOLLARS).

2. Setup the Environment

First, we'll set the working directory using `setwd()`, and then import the required libraries. As we proceed through the report the list of libraries might change.

```
# Set the working directory
setwd("/Users/dreamer/Downloads/Godaddy/godaddy_microbusiness_forecasting")

# Importing the libraries

# Recognize package conflicts
library(conflicted)

# Multi-purpose package for data import, tidying, manipulation, visualisation, and
# programming
# Most common packages include: ggplot2, purrr, tibble, dplyr, tidyverse, readr,
# forcats
library(tidyverse)

# Deal with missing data
library(mice)

# Related to plots
library(maps)
#library(ggmap)
library(gridExtra)
library(mapdata)
library(ggcorrplot)
library(corrplot)

# Training
library(caret)
library(gbm)

# Color palette
library(viridis)

# Future Selection
#library(KernSmooth)
library(glmnet)
library(randomForest)

# libraries required for calculating SMAPE
library(forecast)
library(Metrics)
```

3. Exploratory Data Analysis

3.1. Exploring the datasets

Explore the datasets to get a better understanding of the data.

Load the **train**, **test**, and **census_starter** datasets into R dataframes.

```
# Load train.csv into a dataframe
train_df <- read.csv("./datasets/train.csv")

# Load test.csv into a dataframe
test_df <- read.csv("./datasets/test.csv")

# Load census_starter.csv into a dataframe
census_df <- read.csv("./datasets/census_starter.csv")
```

After reading the CSV files into dataframes, we should check whether the data is loaded correctly or not. We can use the `head()` function of R to display the first few rows of the dataframes and `tail()` function to display the last rows. This will display the first and last six rows of the **train**, **test** and **census** dataframes. We can also use other R functions such as `str()` and `summary()` to get more information about the dataframes, such as column names, data types, and summary statistics.

```
# Display the first 6 rows of the dataframes
head(train_df)
```

```
##           row_id cfips      county state first_day_of_month
## 1 1001_2019-08-01  1001 Autauga County Alabama        2019-08-01
## 2 1001_2019-09-01  1001 Autauga County Alabama        2019-09-01
## 3 1001_2019-10-01  1001 Autauga County Alabama        2019-10-01
## 4 1001_2019-11-01  1001 Autauga County Alabama        2019-11-01
## 5 1001_2019-12-01  1001 Autauga County Alabama        2019-12-01
## 6 1001_2020-01-01  1001 Autauga County Alabama        2020-01-01
##   microbusiness_density active
## 1             3.007682    1249
## 2             2.884870    1198
## 3             3.055843    1269
## 4             2.993233    1243
## 5             2.993233    1243
## 6             2.969090    1242
```

```
head(test_df)
```

```
##           row_id cfips first_day_of_month
## 1 1001_2022-11-01  1001        2022-11-01
## 2 1003_2022-11-01  1003        2022-11-01
## 3 1005_2022-11-01  1005        2022-11-01
## 4 1007_2022-11-01  1007        2022-11-01
## 5 1009_2022-11-01  1009        2022-11-01
## 6 1011_2022-11-01  1011        2022-11-01
```

```
head(census_df)
```

```

##   pct_bb_2017 pct_bb_2018 pct_bb_2019 pct_bb_2020 pct_bb_2021 cfips
## 1      76.6     78.9    80.6     82.7     85.5    1001
## 2      74.5     78.1    81.8     85.1     87.9    1003
## 3      57.2     60.4    60.5     64.6     64.6    1005
## 4      62.0     66.1    69.2     76.1     74.6    1007
## 5      65.8     68.5    73.0     79.6     81.0    1009
## 6      49.4     58.9    60.1     60.6     59.4    1011
##   pct_college_2017 pct_college_2018 pct_college_2019 pct_college_2020
## 1          14.5       15.9     16.1     16.7
## 2          20.4       20.7     21.0     20.2
## 3          7.6        7.8      7.6      7.3
## 4          8.1        7.6      6.5      7.4
## 5          8.7        8.1      8.6      8.9
## 6          6.6        7.4      7.4      6.1
##   pct_college_2021 pct_foreign_born_2017 pct_foreign_born_2018
## 1          16.4       2.1      2.0
## 2          20.6       3.2      3.4
## 3          6.7        2.7      2.5
## 4          7.9        1.0      1.4
## 5          9.3        4.5      4.4
## 6          8.1        1.8      0.9
##   pct_foreign_born_2019 pct_foreign_born_2020 pct_foreign_born_2021
## 1          2.3        2.3      2.1
## 2          3.7        3.4      3.5
## 3          2.7        2.6      2.6
## 4          1.5        1.6      1.1
## 5          4.5        4.4      4.5
## 6          0.7        1.5      1.2
##   pct_it_workers_2017 pct_it_workers_2018 pct_it_workers_2019
## 1          1.3        1.1      0.7
## 2          1.4        1.3      1.4
## 3          0.5        0.3      0.8
## 4          1.2        1.4      1.6
## 5          1.3        1.4      0.9
## 6          0.4        0.3      0.5
##   pct_it_workers_2020 pct_it_workers_2021 median_hh_inc_2017 median_hh_inc_2018
## 1          0.6        1.1      55317    58786
## 2          1.0        1.3      52562    55962
## 3          1.1        0.8      33368    34186
## 4          1.7        2.1      43404    45340
## 5          1.1        0.9      47412    48695
## 6          0.3        0.2      29655    32152
##   median_hh_inc_2019 median_hh_inc_2020 median_hh_inc_2021
## 1          58731      57982     62660
## 2          58320      61756     64346
## 3          32525      34990     36422
## 4          47542      51721     54277
## 5          49358      48922     52830
## 6          37785      33866     29063

```

```

# Display the last 6 rows of the dataframes
tail(train_df)

```

```

##           row_id cfips      county state first_day_of_month
## 122260 56045_2022-05-01 56045 Weston County Wyoming          2022-05-01
## 122261 56045_2022-06-01 56045 Weston County Wyoming          2022-06-01
## 122262 56045_2022-07-01 56045 Weston County Wyoming          2022-07-01
## 122263 56045_2022-08-01 56045 Weston County Wyoming          2022-08-01
## 122264 56045_2022-09-01 56045 Weston County Wyoming          2022-09-01
## 122265 56045_2022-10-01 56045 Weston County Wyoming          2022-10-01
##           microbusiness_density active
## 122260             1.803249    101
## 122261             1.803249    101
## 122262             1.803249    101
## 122263             1.785395    100
## 122264             1.785395    100
## 122265             1.785395    100

```

```
tail(test_df)
```

```

##           row_id cfips first_day_of_month
## 25075 56035_2023-06-01 56035          2023-06-01
## 25076 56037_2023-06-01 56037          2023-06-01
## 25077 56039_2023-06-01 56039          2023-06-01
## 25078 56041_2023-06-01 56041          2023-06-01
## 25079 56043_2023-06-01 56043          2023-06-01
## 25080 56045_2023-06-01 56045          2023-06-01

```

```
tail(census_df)
```

```

##           pct_bb_2017 pct_bb_2018 pct_bb_2019 pct_bb_2020 pct_bb_2021 cfips
## 3137        82.9       81.7       85.6       88.1       89.8 56035
## 3138        82.2       82.4       84.0       86.7       88.4 56037
## 3139        83.5       85.9       87.1       89.1       90.5 56039
## 3140        83.8       88.2       89.5       91.4       90.6 56041
## 3141        76.4       78.3       78.2       82.8       85.4 56043
## 3142        71.1       73.3       76.8       79.7       81.3 56045
##           pct_college_2017 pct_college_2018 pct_college_2019 pct_college_2020
## 3137         19.2        19.0        16.7        21.7
## 3138         15.3        15.2        14.8        13.7
## 3139         37.7        37.8        38.9        37.2
## 3140         11.9        10.5        11.1        12.6
## 3141         15.4        15.0        15.4        15.0
## 3142         14.1        13.5        13.4        12.7
##           pct_college_2021 pct_foreign_born_2017 pct_foreign_born_2018
## 3137         20.9          3.9          3.1
## 3138         12.4          5.0          5.3
## 3139         38.3         10.8         11.2
## 3140         12.3          2.9          3.1
## 3141         17.2          2.3          1.4
## 3142         13.9          3.8          4.1
##           pct_foreign_born_2019 pct_foreign_born_2020 pct_foreign_born_2021
## 3137            4.4            5.1            5.1

```

```

## 3138           4.7           5.2           5.5
## 3139          11.8          11.4          11.1
## 3140           2.9           2.9           2.9
## 3141           1.6           2.2           1.0
## 3142           1.7           2.3           1.6
##   pct_it_workers_2017 pct_it_workers_2018 pct_it_workers_2019
## 3137          0.1           0.0           0.0
## 3138          0.6           0.6           1.0
## 3139          0.7           1.2           1.4
## 3140          1.2           1.2           1.4
## 3141          1.3           1.0           0.9
## 3142          0.6           0.6           0.0
##   pct_it_workers_2020 pct_it_workers_2021 median_hh_inc_2017
## 3137          0.0           0.0          84911
## 3138          0.9           1.0          71083
## 3139          1.5           2.0          80049
## 3140          1.7           0.9          54672
## 3141          0.9           1.1          51362
## 3142          0.0           0.0          59605
##   median_hh_inc_2018 median_hh_inc_2019 median_hh_inc_2020
## 3137         78680          77403          78655
## 3138         73008          74843          73384
## 3139         83831          84678          87053
## 3140         58235          63403          72458
## 3141         53426          54158          57306
## 3142         52867          57031          53333
##   median_hh_inc_2021
## 3137         82342
## 3138         76668
## 3139         94498
## 3140         75106
## 3141         62271
## 3142         65566

```

```

# Display information about the train dataframe
str(train_df)

```

```

## 'data.frame':    122265 obs. of  7 variables:
## $ row_id          : chr  "1001_2019-08-01" "1001_2019-09-01" "1001_2019-09-01" ...
## $ cfips          : int  1001 1001 1001 1001 1001 1001 1001 1001 1001 1001 ...
## $ county         : chr  "Autauga County" "Autauga County" "Autauga County" ...
## $ state          : chr  "Alabama" "Alabama" "Alabama" "Alabama" ...
## $ first_day_of_month: chr  "2019-08-01" "2019-09-01" "2019-10-01" "2019-11-01" ...
## $ microbusiness_density: num  3.01 2.88 3.06 2.99 2.99 ...
## $ active          : int  1249 1198 1269 1243 1243 1242 1217 1227 1255 1257 ...

```

```
#cat(rep("=", 40), "\n") # Print a line of 40 equal signs
summary(train_df)
```

```
##      row_id          cfips        county       state
## Length:122265    Min.   : 1001    Length:122265    Length:122265
## Class :character  1st Qu.:18177   Class :character  Class :character
## Mode  :character  Median :29173    Mode  :character  Mode  :character
##                   Mean   :30376
##                   3rd Qu.:45077
##                   Max.  :56045
## first_day_of_month microbusiness_density     active
## Length:122265    Min.   : 0.000    Min.   :      0
## Class :character  1st Qu.: 1.639   1st Qu.: 145
## Mode  :character  Median : 2.587   Median : 488
##                   Mean   : 3.818   Mean   : 6443
##                   3rd Qu.: 4.519   3rd Qu.: 2124
##                   Max.  :284.340   Max.  :1167744
```

```
# Display information about the test dataframe
str(test_df)
```

```
## 'data.frame': 25080 obs. of 3 variables:
## $ row_id           : chr "1001_2022-11-01" "1003_2022-11-01" "1005_2022-11-01"
## "1007_2022-11-01" ...
## $ cfips            : int 1001 1003 1005 1007 1009 1011 1013 1015 1017 1019 .
## ..
## $ first_day_of_month: chr "2022-11-01" "2022-11-01" "2022-11-01" "2022-11-01"
## ...
```

```
#cat(rep("=", 40), "\n") # Print a line of 40 equal signs
summary(test_df)
```

```
##      row_id          cfips      first_day_of_month
## Length:25080    Min.   : 1001    Length:25080
## Class :character  1st Qu.:18177   Class :character
## Mode  :character  Median :29173    Mode  :character
##                   Mean   :30376
##                   3rd Qu.:45077
##                   Max.  :56045
```

```
# Display information about the census dataframe
str(census_df)
```

```

## 'data.frame': 3142 obs. of 26 variables:
## $ pct_bb_2017 : num 76.6 74.5 57.2 62 65.8 49.4 58.2 71 62.8 67.5 ...
## $ pct_bb_2018 : num 78.9 78.1 60.4 66.1 68.5 58.9 62.1 73 66.5 68.6 ...
## $ pct_bb_2019 : num 80.6 81.8 60.5 69.2 73 60.1 64.6 75.1 69.4 70.7 ...
## $ pct_bb_2020 : num 82.7 85.1 64.6 76.1 79.6 60.6 73.6 79.8 74.5 75 ...
## $ pct_bb_2021 : num 85.5 87.9 64.6 74.6 81 59.4 76.3 81.6 77.1 76.7 ...
## $ cfips : int 1001 1003 1005 1007 1009 1011 1013 1015 1017 1019 ...
## $ pct_college_2017 : num 14.5 20.4 7.6 8.1 8.7 6.6 9.6 10.2 9 6.6 ...
## $ pct_college_2018 : num 15.9 20.7 7.8 7.6 8.1 7.4 9.7 10.2 9.3 6.8 ...
## $ pct_college_2019 : num 16.1 21 7.6 6.5 8.6 7.4 9.7 10.5 9.5 6.6 ...
## $ pct_college_2020 : num 16.7 20.2 7.3 7.4 8.9 6.1 10.1 10.5 10.5 6.3 ...
## $ pct_college_2021 : num 16.4 20.6 6.7 7.9 9.3 8.1 8.1 11.4 9.6 6.2 ...
## $ pct_foreign_born_2017: num 2.1 3.2 2.7 1 4.5 1.8 1 2.6 1.3 0.7 ...
## $ pct_foreign_born_2018: num 2 3.4 2.5 1.4 4.4 0.9 1.4 2.7 1.4 0.8 ...
## $ pct_foreign_born_2019: num 2.3 3.7 2.7 1.5 4.5 0.7 0.8 2.7 1.8 0.9 ...
## $ pct_foreign_born_2020: num 2.3 3.4 2.6 1.6 4.4 1.5 1.9 2.5 1.9 1.9 ...
## $ pct_foreign_born_2021: num 2.1 3.5 2.6 1.1 4.5 1.2 1.7 2.5 2 2 ...
## $ pct_it_workers_2017 : num 1.3 1.4 0.5 1.2 1.3 0.4 1.1 1.4 2.4 1.4 ...
## $ pct_it_workers_2018 : num 1.1 1.3 0.3 1.4 1.4 0.3 1.4 1.4 2.1 1.3 ...
## $ pct_it_workers_2019 : num 0.7 1.4 0.8 1.6 0.9 0.5 1.7 1.2 2.1 1.2 ...
## $ pct_it_workers_2020 : num 0.6 1 1.1 1.7 1.1 0.3 1.3 1 2.3 0.9 ...
## $ pct_it_workers_2021 : num 1.1 1.3 0.8 2.1 0.9 0.2 1.4 1 1.8 0.4 ...
## $ median_hh_inc_2017 : int 55317 52562 33368 43404 47412 29655 36326 43686
37342 40041 ...
## $ median_hh_inc_2018 : num 58786 55962 34186 45340 48695 ...
## $ median_hh_inc_2019 : int 58731 58320 32525 47542 49358 37785 40688 47255
42289 41919 ...
## $ median_hh_inc_2020 : num 57982 61756 34990 51721 48922 ...
## $ median_hh_inc_2021 : num 62660 64346 36422 54277 52830 ...

```

```

#cat(rep("=", 40), "\n") # Print a line of 40 equal signs
summary(census_df)

```

	pct_bb_2017	pct_bb_2018	pct_bb_2019	pct_bb_2020
## Min.	:24.50	:25.70	:34.80	:33.30
## 1st Qu.	:64.20	:67.42	:70.50	:74.10
## Median	:70.70	:73.60	:76.45	:79.60
## Mean	:69.92	:72.69	:75.40	:78.54
## 3rd Qu.	:76.40	:78.80	:81.40	:84.10
## Max.	:94.60	:95.50	:96.00	:97.10
##			NA's	:1
	pct_bb_2021	cfips	pct_college_2017	pct_college_2018
## Min.	:37.00	: 1001	: 2.40	: 0.00
## 1st Qu.	:76.40	:18178	: 9.70	: 9.90
## Median	:81.70	:29176	:12.80	:13.00
## Mean	:80.54	:30384	:13.81	:14.01

```

## 3rd Qu.:85.90   3rd Qu.:45080    3rd Qu.:16.80    3rd Qu.:17.10
## Max.     :97.60   Max.      :56045    Max.      :43.70    Max.      :48.00
## NA's     :1

## pct_college_2019 pct_college_2020 pct_college_2021 pct_foreign_born_2017
## Min.     : 0.00   Min.     : 0.00   Min.     : 0.00   Min.     : 0.000
## 1st Qu.:10.10   1st Qu.:10.50   1st Qu.:10.60   1st Qu.: 1.400
## Median  :13.25   Median  :13.60   Median  :13.80   Median  : 2.700
## Mean    :14.24   Mean    :14.63   Mean    :14.85   Mean    : 4.702
## 3rd Qu.:17.30   3rd Qu.:17.90   3rd Qu.:18.00   3rd Qu.: 5.700
## Max.    :45.40   Max.    :43.00   Max.    :43.70   Max.    :52.900
## NA's     :1       NA's     :1       NA's     :1

## pct_foreign_born_2018 pct_foreign_born_2019 pct_foreign_born_2020
## Min.     : 0.000   Min.     : 0.000   Min.     : 0.000
## 1st Qu.: 1.400   1st Qu.: 1.400   1st Qu.: 1.400
## Median  : 2.700   Median  : 2.700   Median  : 2.800
## Mean    : 4.725   Mean    : 4.769   Mean    : 4.749
## 3rd Qu.: 5.700   3rd Qu.: 5.700   3rd Qu.: 5.700
## Max.    :53.300   Max.    :53.700   Max.    :54.000
## NA's     :1

## pct_foreign_born_2021 pct_it_workers_2017 pct_it_workers_2018
## Min.     : 0.000   Min.     : 0.000   Min.     : 0.000
## 1st Qu.: 1.400   1st Qu.: 0.800   1st Qu.: 0.800
## Median  : 2.700   Median  : 1.300   Median  : 1.300
## Mean    : 4.744   Mean    : 1.427   Mean    : 1.382
## 3rd Qu.: 5.700   3rd Qu.: 1.900   3rd Qu.: 1.800
## Max.    :54.000   Max.    :17.400   Max.    :11.700
## NA's     :1       NA's     :1

## pct_it_workers_2019 pct_it_workers_2020 pct_it_workers_2021 median_hh_inc_2017
## Min.     : 0.000   Min.     : 0.000   Min.     : 0.000   Min.     : 19264
## 1st Qu.: 0.700   1st Qu.: 0.700   1st Qu.: 0.600   1st Qu.: 41123
## Median  : 1.200   Median  : 1.200   Median  : 1.100   Median  : 48066
## Mean    : 1.339   Mean    : 1.309   Mean    : 1.273   Mean    : 49754
## 3rd Qu.: 1.800   3rd Qu.: 1.800   3rd Qu.: 1.700   3rd Qu.: 55764
## Max.    :10.500   Max.    :15.200   Max.    :15.200   Max.    :129588
## NA's     :1       NA's     :1

## median_hh_inc_2018 median_hh_inc_2019 median_hh_inc_2020 median_hh_inc_2021
## Min.     : 20188   Min.     : 21504   Min.     : 22292   Min.     : 17109
## 1st Qu.: 42480   1st Qu.: 44155   1st Qu.: 45653   1st Qu.: 48180
## Median  : 49888   Median  : 51758   Median  : 52842   Median  : 55907
## Mean    : 51583   Mean    : 53476   Mean    : 55012   Mean    : 58223
## 3rd Qu.: 57611   3rd Qu.: 59867   3rd Qu.: 61501   3rd Qu.: 64930
## Max.    :136268   Max.    :142299   Max.    :147111   Max.    :156821
## NA's     :1       NA's     :2       NA's     :2

```

The data type of `first_day_of_month` column in `train_df` and `test_df` is `character`. We will convert the character to `Date` format.

```

# Change first_day_of_month data type to Date
train_df$first_day_of_month <- as.Date(train_df$first_day_of_month)
test_df$first_day_of_month <- as.Date(test_df$first_day_of_month)
str(train_df$first_day_of_month)

```

```
## Date[1:122265], format: "2019-08-01" "2019-09-01" "2019-10-01" "2019-11-01" "2  
019-12-01" ...
```

3.2. Checking the Dataframes for Missing Values

The `is.na()` function is used to create a logical matrix where *TRUE* represents a missing value and *FALSE* represents a non-missing value. The `colSums()` function is then used to count the number of missing values in each column of the data frame. If the sum of a column is greater than 0, it means that there is at least one missing value in that column.

```
# Check for missing values in the train data frame  
colSums(is.na(train_df))
```

```
##           row_id          cfips        county  
##             0              0              0  
##           state first_day_of_month microbusiness_density  
##             0                  0                  0  
##         active  
##             0
```

```
# Check for missing values in the test data frame  
colSums(is.na(test_df))
```

```
##           row_id          cfips first_day_of_month  
##             0              0                  0
```

```
# Check for missing values in the census data frame  
colSums(is.na(census_df))
```

```

##          pct_bb_2017          pct_bb_2018          pct_bb_2019
##                0                      0                      0
##          pct_bb_2020          pct_bb_2021           cfips
##                1                      1                      0
##          pct_college_2017      pct_college_2018      pct_college_2019
##                0                      0                      0
##          pct_college_2020      pct_college_2021 pct_foreign_born_2017
##                1                      1                      0
##          pct_foreign_born_2018 pct_foreign_born_2019 pct_foreign_born_2020
##                0                      0                      1
##          pct_foreign_born_2021      pct_it_workers_2017      pct_it_workers_2018
##                1                      0                      1
##          pct_it_workers_2019      pct_it_workers_2020      pct_it_workers_2021
##                0                      1                      1
##          median_hh_inc_2017      median_hh_inc_2018      median_hh_inc_2019
##                0                      1                      0
##          median_hh_inc_2020      median_hh_inc_2021
##                2                      2

```

```

#{r fig.width=7, fig.align='center', fig.height=4, out.width='100%'}
# Calculate the number and percentage of missing values for each column
missing_data <- census_df %>%
  summarise_all(~ sum(is.na(.))) %>%
  gather(variable, missing_count) %>%
  mutate(missing_percent = missing_count/nrow(census_df)*100)

# Create two plots side by side
plot1 <- ggplot(missing_data, aes(x = missing_count, y = variable)) +
  geom_bar(stat = "identity", fill = "steelblue") +
  labs(x = "Number of missing values", y = "") +
  ggtitle("Number of missing values in census_df") +
  theme_gray()

plot2 <- ggplot(missing_data, aes(x = missing_percent, y = variable)) +
  geom_bar(stat = "identity", fill = "lightblue") +
  labs(x = "Percentage of missing values", y = "") +
  ggtitle("Percentage of missing values in census_df") +
  theme_gray()

# Arrange the two plots side by side
grid.arrange(plot1, plot2, ncol = 2)

```



We use the `complete.cases()` function to determine which rows have complete data and which rows have missing values. This function returns a logical vector indicating which rows have no missing values. Therefore, to identify the rows with missing values, we use the `!` operator to negate the logical vector returned by `complete.cases()`. Then, we use the `is.na()` function to identify which columns have missing values for each missing row:

```
# Identify rows with missing values in census_df
missing_rows <- which(!complete.cases(census_df))

# Identify columns with missing values for each missing row
for (i in missing_rows) {
  cat("Row", i, "has missing values in columns:",
      paste(names(census_df)[is.na(census_df[i,])], collapse = ", "), "\n")
}
```

```

## Row 93 has missing values in columns: pct_bb_2020, pct_bb_2021, pct_college_202
0, pct_college_2021, pct_foreign_born_2020, pct_foreign_born_2021, pct_it_workers_
2020, pct_it_workers_2021, median_hh_inc_2020, median_hh_inc_2021
## Row 1817 has missing values in columns: pct_it_workers_2018, median_hh_inc_2018
## Row 2645 has missing values in columns: median_hh_inc_2020
## Row 2674 has missing values in columns: median_hh_inc_2021

```

```
print(census_df[missing_rows,])
```

	pct_bb_2017	pct_bb_2018	pct_bb_2019	pct_bb_2020	pct_bb_2021	cfips
## 93	80.5	79.1	80.4	NA	NA	2261
## 1817	49.1	52.1	57.6	60.7	63.5	35039
## 2645	66.3	66.6	61.2	63.2	70.1	48243
## 2674	64.5	72.7	73.3	96.8	97.0	48301
	pct_college_2017	pct_college_2018	pct_college_2019	pct_college_2020		
## 93	23.1	19.0		16.5		NA
## 1817		12.0	12.5	12.6		10.6
## 2645		18.4	16.0	10.8		14.3
## 2674		4.7	0.0	0.0		0.0
	pct_college_2021	pct_foreign_born_2017	pct_foreign_born_2018			
## 93		NA	4.9		6.3	
## 1817		10.1	4.5		3.7	
## 2645		10.9	22.4		14.9	
## 2674		0.0	10.8		15.7	
	pct_foreign_born_2019	pct_foreign_born_2020	pct_foreign_born_2021			
## 93		6.6	NA		NA	
## 1817		4.2	4.5		4.8	
## 2645		20.9	10.1		12.7	
## 2674		12.2	0.0		1.2	
	pct_it_workers_2017	pct_it_workers_2018	pct_it_workers_2019			
## 93		3.3	3.9		5.3	
## 1817		0.8	NA		0.8	
## 2645		0.0	0.0		0.0	
## 2674		0.0	0.0		0.0	
	pct_it_workers_2020	pct_it_workers_2021	median_hh_inc_2017			
## 93		NA	NA		86019	
## 1817		0.4	0.7		33422	
## 2645		0.0	0.0		46534	
## 2674		0.0	0.0		80938	
	median_hh_inc_2018	median_hh_inc_2019	median_hh_inc_2020			
## 93	82306	79867			NA	
## 1817		NA	39952		42264	
## 2645	53194	53088			NA	
## 2674	81875	83750			44076	
	median_hh_inc_2021					
## 93		NA				
## 1817		46994				
## 2645		38659				
## 2674		NA				

3.3. Dealing with Missing Values

The `mice` package implements a method to deal with missing data. The package creates multiple imputations (replacement values) for multivariate missing data.

We'll use the `mice` package to impute missing values in the `census_df` dataframe with below arguments:

- *m*: The number of imputations to generate was set to 5, because, generally, *m* should be set to at least 5 for good imputation performance. Creating too many datasets will increase the computational load and may not necessarily lead to better results.
- *maxit*: The *maxit* value was set to 50 to allow for a larger number of iterations to ensure that the imputation algorithm converges and fills in missing values as accurately as possible.
- *method*: In this case, we are using “*pmm*” which stands for *Predictive Mean Matching*, because it is a flexible and widely used imputation method that works well with continuous variables. The method estimates the missing values by drawing from a set of observed values that have similar characteristics to the missing values.
- *print*: The print value is set to *FALSE* because this function prints a huge log output to console.

```
# Impute missing data using mice
imputed_data <- mice(census_df, m = 5, maxit = 50, method = "pmm", print = FALSE)
# Extract imputed data
imputed_data <- complete(imputed_data)
```

```
# Check for missing values in imputed data
colSums(is.na(imputed_data))
```

```
##          pct_bb_2017      pct_bb_2018      pct_bb_2019
##                 0                  0                  0
##          pct_bb_2020      pct_bb_2021      cfips
##                 0                  0                  0
##          pct_college_2017    pct_college_2018    pct_college_2019
##                 0                  0                  0
##          pct_college_2020    pct_college_2021 pct_foreign_born_2017
##                 0                  0                  0
##          pct_foreign_born_2018 pct_foreign_born_2019 pct_foreign_born_2020
##                 0                  0                  0
##          pct_foreign_born_2021    pct_it_workers_2017    pct_it_workers_2018
##                 0                  0                  0
##          pct_it_workers_2019    pct_it_workers_2020    pct_it_workers_2021
##                 0                  0                  0
##          median_hh_inc_2017    median_hh_inc_2018    median_hh_inc_2019
##                 0                  0                  0
##          median_hh_inc_2020    median_hh_inc_2021
##                 0                  0
```

```
# Check the filled missing values
print(imputed_data[missing_rows,])
```

```

##      pct_bb_2017 pct_bb_2018 pct_bb_2019 pct_bb_2020 pct_bb_2021 cfips
## 93          80.5       79.1     80.4      81.0      83.5    2261
## 1817        49.1       52.1     57.6      60.7      63.5  35039
## 2645        66.3       66.6     61.2      63.2      70.1  48243
## 2674        64.5       72.7     73.3      96.8      97.0  48301
##      pct_college_2017 pct_college_2018 pct_college_2019 pct_college_2020
## 93            23.1       19.0     16.5      15.1
## 1817          12.0       12.5     12.6      10.6
## 2645          18.4       16.0     10.8      14.3
## 2674          4.7        0.0      0.0      0.0
##      pct_college_2021 pct_foreign_born_2017 pct_foreign_born_2018
## 93            15.1       4.9      6.3
## 1817          10.1       4.5      3.7
## 2645          10.9       22.4     14.9
## 2674          0.0        10.8     15.7
##      pct_foreign_born_2019 pct_foreign_born_2020 pct_foreign_born_2021
## 93            6.6        6.0      7.6
## 1817          4.2        4.5      4.8
## 2645          20.9       10.1     12.7
## 2674          12.2       0.0      1.2
##      pct_it_workers_2017 pct_it_workers_2018 pct_it_workers_2019
## 93            3.3        3.9      5.3
## 1817          0.8        0.5      0.8
## 2645          0.0        0.0      0.0
## 2674          0.0        0.0      0.0
##      pct_it_workers_2020 pct_it_workers_2021 median_hh_inc_2017
## 93            5.6        1.0     86019
## 1817          0.4        0.7     33422
## 2645          0.0        0.0     46534
## 2674          0.0        0.0     80938
##      median_hh_inc_2018 median_hh_inc_2019 median_hh_inc_2020
## 93           82306       79867     78889
## 1817         36148       39952     42264
## 2645         53194       53088     42285
## 2674         81875       83750     44076
##      median_hh_inc_2021
## 93           81817
## 1817         46994
## 2645         38659
## 2674         52842

```

3.4. Checking the Time Frame of *train* and *test* Dataframes

After dealing with the missing values, we have to check the time frame provided in the **train** and **test** datasets.

```

index <- unique(train_df$first_day_of_month)
print(index)

```

```

## [1] "2019-08-01" "2019-09-01" "2019-10-01" "2019-11-01" "2019-12-01"
## [6] "2020-01-01" "2020-02-01" "2020-03-01" "2020-04-01" "2020-05-01"
## [11] "2020-06-01" "2020-07-01" "2020-08-01" "2020-09-01" "2020-10-01"
## [16] "2020-11-01" "2020-12-01" "2021-01-01" "2021-02-01" "2021-03-01"
## [21] "2021-04-01" "2021-05-01" "2021-06-01" "2021-07-01" "2021-08-01"
## [26] "2021-09-01" "2021-10-01" "2021-11-01" "2021-12-01" "2022-01-01"
## [31] "2022-02-01" "2022-03-01" "2022-04-01" "2022-05-01" "2022-06-01"
## [36] "2022-07-01" "2022-08-01" "2022-09-01" "2022-10-01"

```

The training data time frame includes 08/2019 to 10/2022

```

index <- unique(test_df$first_day_of_month)
print(index)

```

```

## [1] "2022-11-01" "2022-12-01" "2023-01-01" "2023-02-01" "2023-03-01"
## [6] "2023-04-01" "2023-05-01" "2023-06-01"

```

The prediction dates provided include 11/2022 to 06/2023

3.5. Adding New Columns to *train* and *test*

To make analysis easier and be able to group the data by year and month, we will use `substr()` function to extract the relevant characters of the `first_day_of_month` column, which is a string that contains the date in the format “YYYY-MM-DD”. Then, `as.integer()` function is used to convert the extracted year and month values from character strings to integers.

```

# Add year, month and year_month columns to train_df
train_df$year <- as.integer(substr(train_df$first_day_of_month, 1, 4))
train_df$month <- as.integer(substr(train_df$first_day_of_month, 6, 7))
train_df$year_month <- substr(train_df$first_day_of_month, 1, 7)
str(train_df)

```

```

## 'data.frame': 122265 obs. of 10 variables:
## $ row_id           : chr  "1001_2019-08-01" "1001_2019-09-01" "1001_2019-1
0-01" "1001_2019-11-01" ...
## $ cfips            : int  1001 1001 1001 1001 1001 1001 1001 1001 100
1 ...
## $ county           : chr  "Autauga County" "Autauga County" "Autauga Count
y" "Autauga County" ...
## $ state             : chr  "Alabama" "Alabama" "Alabama" "Alabama" ...
## $ first_day_of_month: Date, format: "2019-08-01" "2019-09-01" ...
## $ microbusiness_density: num  3.01 2.88 3.06 2.99 2.99 ...
## $ active            : int  1249 1198 1269 1243 1243 1242 1217 1227 1255 125
7 ...
## $ year              : int  2019 2019 2019 2019 2019 2020 2020 2020 202
0 ...
## $ month             : int  8 9 10 11 12 1 2 3 4 5 ...
## $ year_month        : chr  "2019-08" "2019-09" "2019-10" "2019-11" ...

```

```
# Add year, month and year_month columns to test_df
test_df$year <- as.integer(substr(test_df$first_day_of_month, 1, 4))
test_df$month <- as.integer(substr(test_df$first_day_of_month, 6, 7))
test_df$year_month <- substr(test_df$first_day_of_month, 1, 7)
```

3.6. Merging the *train* and *imputed_data* datasets

The merging process is challenging because all data fields provided in the **imputed_data** (formerly **census_df**) dataframe have a two-year lag to match the data in the **train_df** dataframe. Also, the data provided in the **imputed_data** is on a yearly basis, but the data in the **train_df** dataframe is on a monthly basis. To merge these two dataframes, it is assumed that the yearly data provided is valid for all the months of the corresponding year. For example, data provided in the `pct_bb_2017` is valid for all the months of 2019 in the **train_df**.

```

# Set variables of interest
vars <- c("pct_bb", "pct_college", "pct_foreign_born", "pct_it_workers", "median_hh_inc")

# Loop through variables and merge with train_df
merged_df <- train_df

for (var in vars) {
  # Select columns and pivot longer
  merged_df <- imputed_data %>%
    select(cfips, paste0(var, "_2017"):paste0(var, "_2020")) %>%
    pivot_longer(cols = starts_with(var),
                 names_to = "year",
                 values_to = var) %>%
  # Modify year and month columns
  mutate(year = as.integer(str_sub(year, -4)) + 2) %>%
  uncount(12, .id = "month") %>%
  mutate(month = month) %>%
  # Merge with merged_df
  merge(merged_df, by = c("cfips", "year", "month"), all.x = TRUE) %>%
  arrange(cfips, row_id)
}

merged_df <- merged_df %>%
  select(row_id, cfips, county, state, first_day_of_month, microbusiness_density,
active, year_month, year, month, pct_bb, pct_college, pct_foreign_born, pct_it_workers, median_hh_inc)

merged_test <- test_df
for (var in vars) {
  # Select columns and pivot longer
  merged_test <- imputed_data %>%
    select(cfips, paste0(var, "_2020"):paste0(var, "_2021")) %>%
    pivot_longer(cols = starts_with(var),
                 names_to = "year",
                 values_to = var) %>%
  # Modify year and month columns
  mutate(year = as.integer(str_sub(year, -4)) + 2) %>%
  uncount(12, .id = "month") %>%
  mutate(month = month) %>%
  # Merge with merged_df
  merge(merged_test, by = c("cfips", "year", "month"), all.x = TRUE) %>%
  arrange(cfips, row_id)
}

merged_test <- merged_test %>%
  select(row_id, cfips, first_day_of_month, year_month, year, month, pct_bb, pct_college, pct_foreign_born, pct_it_workers, median_hh_inc)

```

```
colSums(is.na(merged_df))
```

```

##           row_id          cfips        county
##      28551              0            28551
##       state   first_day_of_month microbusiness_density
##      28551              28551            28551
##      active      year_month           year
##      28551              28551            0
##      month          pct_bb      pct_college
##          0                  0            0
##      pct_foreign_born      pct_it_workers median_hh_inc
##          0                      0            0

```

```
colSums(is.na(merged_test))
```

```

##           row_id          cfips first_day_of_month year_month
##      50328              0            50328            50328
##       year          month          pct_bb      pct_college
##          0                  0                  0            0
##      pct_foreign_born      pct_it_workers median_hh_inc
##          0                      0            0

```

Since the data from 1/2019 to 7/2019 and 11/2022 to 12/2022 is not available in **train_df** merging the data has created NA values in **merged_df** for those months. Now we have to remove the rows with missing values.

```

# remove NA values created in merged_df
merged_df <- merged_df %>%
  na.omit(merged_df)

merged_test <- merged_test %>%
  na.omit(merged_test)

```

```
colSums(is.na(merged_df))
```

```

##           row_id          cfips        county
##          0                  0            0
##       state   first_day_of_month microbusiness_density
##          0                  0            0
##      active      year_month           year
##          0                  0            0
##      month          pct_bb      pct_college
##          0                  0            0
##      pct_foreign_born      pct_it_workers median_hh_inc
##          0                      0            0

```

```
colSums(is.na(merged_test))
```

```

##           row_id          cfips first_day_of_month      year_month
##             0                  0                      0                      0
##       year        month          pct_bb      pct_college
##             0                  0                      0                      0
##   pct_foreign_born    pct_it_workers median_hh_inc
##                     0                          0                      0

```

```
summary(merged_df)
```

```

##       row_id          cfips      county      state
## Length:122265  Min.   : 1001  Length:122265  Length:122265
## Class  :character 1st Qu.:18177  Class  :character  Class  :character
## Mode   :character Median :29173  Mode   :character  Mode   :character
##                   Mean   :30376
##                   3rd Qu.:45077
##                   Max.  :56045
##   first_day_of_month microbusiness_density      active
##                   Min.   : 0.000  Min.   :     0
## 1st Qu.: 1.639  1st Qu.: 145
## Median : 2.587  Median : 488
## Mean   : 3.818  Mean   : 6443
## 3rd Qu.: 4.519  3rd Qu.: 2124
## Max.   :284.340  Max.   :1167744
##       year_month        year        month      pct_bb
## Length:122265  Min.   :2019  Min.   : 1.000  Min.   :24.50
## Class  :character 1st Qu.:2020  1st Qu.: 4.000  1st Qu.:69.30
## Mode   :character Median :2021  Median : 7.000  Median :75.80
##                   Mean   :2021  Mean   : 6.692  Mean   :74.69
##                   3rd Qu.:2022  3rd Qu.:10.000 3rd Qu.:81.20
##                   Max.   :2022  Max.   :12.000  Max.   :97.10
##       pct_college    pct_foreign_born pct_it_workers median_hh_inc
##       Min.   : 0.00  Min.   : 0.000  Min.   : 0.000  Min.   :19264
## 1st Qu.:10.10  1st Qu.: 1.400  1st Qu.: 0.700  1st Qu.:43505
## Median :13.20  Median : 2.700  Median : 1.200  Median :51094
## Mean   :14.22  Mean   : 4.745  Mean   : 1.356  Mean   :52830
## 3rd Qu.:17.30  3rd Qu.: 5.700  3rd Qu.: 1.800  3rd Qu.:59230
## Max.   :48.00  Max.   :54.000  Max.   :17.400  Max.   :147111

```

3.7. Data Visualization

The main feature in this project is `microbusiness_density` provided in the `merged_df`. Also, the number of active microbusinesses is provided in the `active` column.

3.7.1. Overall Microbusiness Density and Count

First, we will plot overall microbusiness density and count of active microbusiness in the United States:

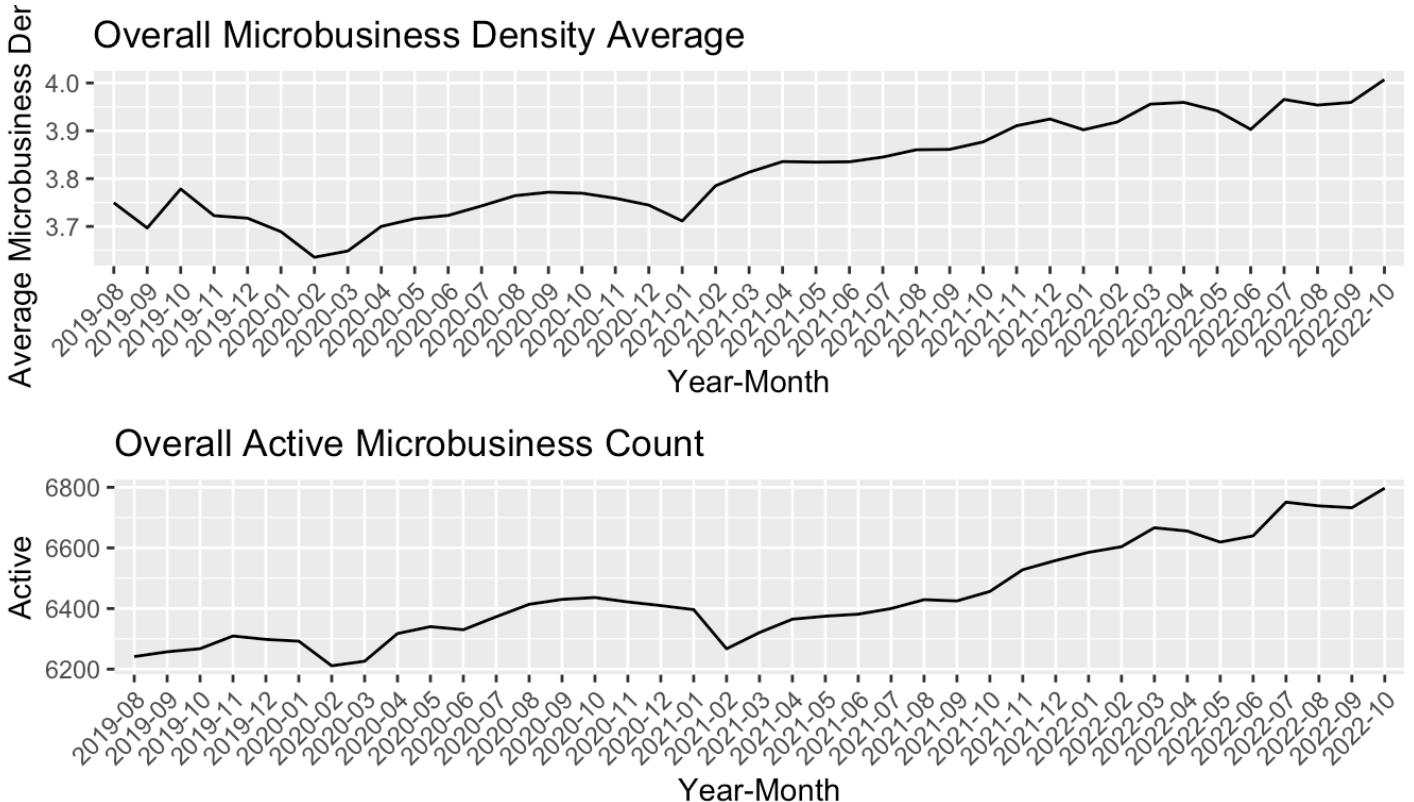
```

# Create plots
p1 <- merged_df %>%
  # Group merged_df by year_month
  group_by(year_month) %>%
  # calculate the mean value of microbusiness_density for each group
  summarise(mean_microbusiness_density = mean(microbusiness_density)) %>%
  ggplot(aes(x = year_month, y = mean_microbusiness_density, group = 1)) +
  geom_line() +
  labs(title = "Overall Microbusiness Density Average",
       x = "Year-Month",
       y = "Average Microbusiness Density") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

p2 <- merged_df %>%
  group_by(year_month) %>%
  summarize(avg_active = mean(active)) %>%
  ggplot(aes(x = year_month, y = avg_active)) +
  geom_line(group = 1) +
  labs(title = "Overall Active Microbusiness Count",
       x = "Year-Month",
       y = "Active") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

# Display the plots
grid.arrange(p1, p2, nrow = 2)

```



As expected, these two graphs show almost similar behavior. If we ignore the slight fluctuations of the two graphs, the general microbusiness density and count are growing over the whole time frame.

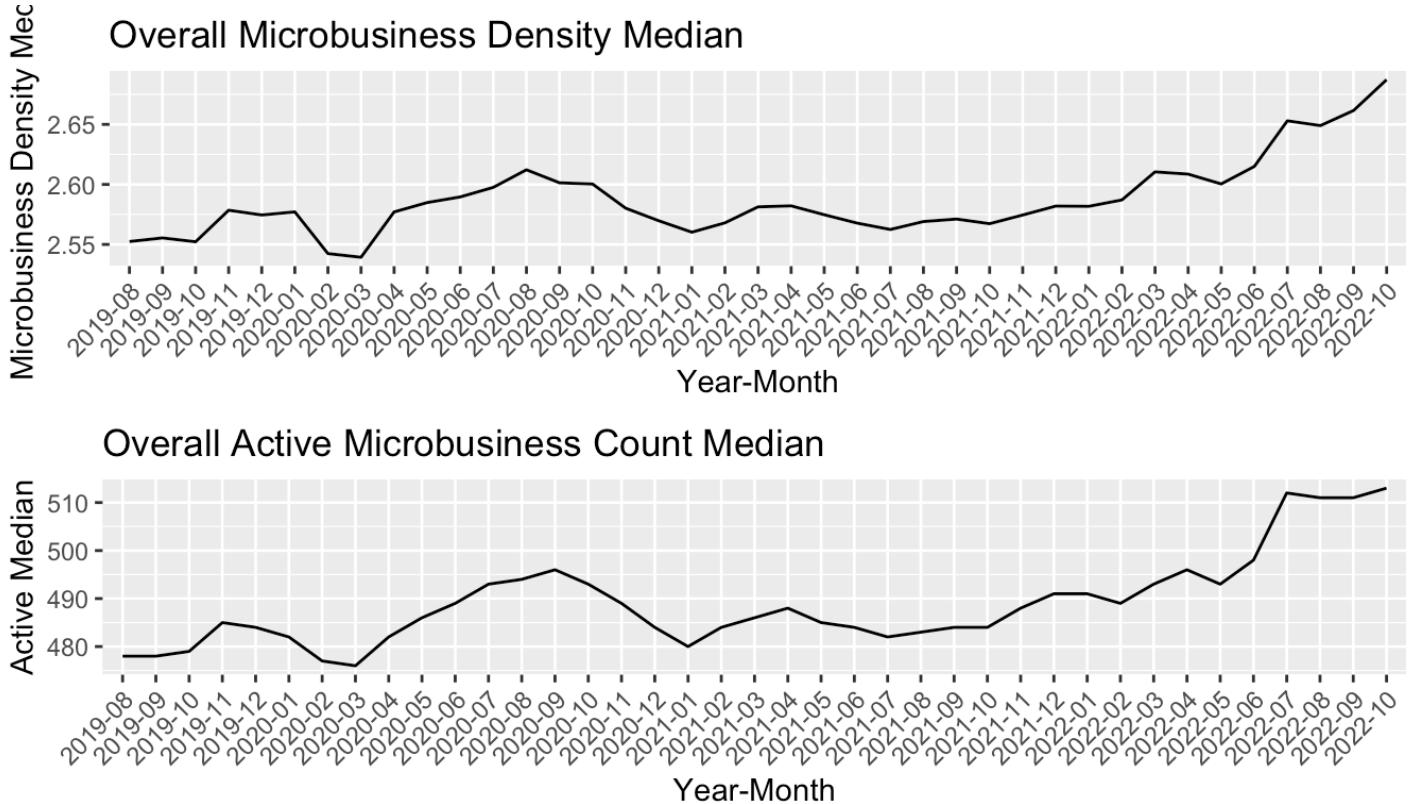
```

# Create plots
p1 <- merged_df %>%
  # Group merged_df by year_month
  group_by(year_month) %>%
  # calculate the median value of microbusiness_density for each group
  summarise(median_microbusiness_density = median(microbusiness_density)) %>%
  ggplot(aes(x = year_month, y = median_microbusiness_density, group = 1)) +
  geom_line() +
  labs(title = "Overall Microbusiness Density Median",
       x = "Year-Month",
       y = "Microbusiness Density Median") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

p2 <- merged_df %>%
  group_by(year_month) %>%
  summarize(median_active = median(active)) %>%
  ggplot(aes(x = year_month, y = median_active)) +
  geom_line(group = 1) +
  labs(title = "Overall Active Microbusiness Count Median",
       x = "Year-Month",
       y = "Active Median") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

# Display the plots
grid.arrange(p1, p2, nrow = 2)

```



Then, we will examine the behavior of these two variables (*microbusiness density* and *active*) while grouping the data by month and year:

```

# Group merged_df by year and calculate the mean value of microbusiness_density and active for each group
merged_df_mean_year <- merged_df %>%
  group_by(year) %>%
  summarize(avg_microbusiness_density = mean(microbusiness_density),
            avg_active = mean(active))

# Group merged_df by month and calculate the mean value of the microbusiness_density for each group
merged_df_mean_month <- merged_df %>%
  group_by(month) %>%
  summarize(avg_microbusiness_density = mean(microbusiness_density),
            avg_active = mean(active))

# Plot the monthly mean values for microbusiness density
p1 <-
  ggplot(merged_df_mean_month, aes(x = month, y = avg_microbusiness_density)) +
  geom_line() +
  ggtitle("Average Monthly Microbusiness Density") +
  xlab("Month") +
  ylab("Avg Microbusiness Density") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

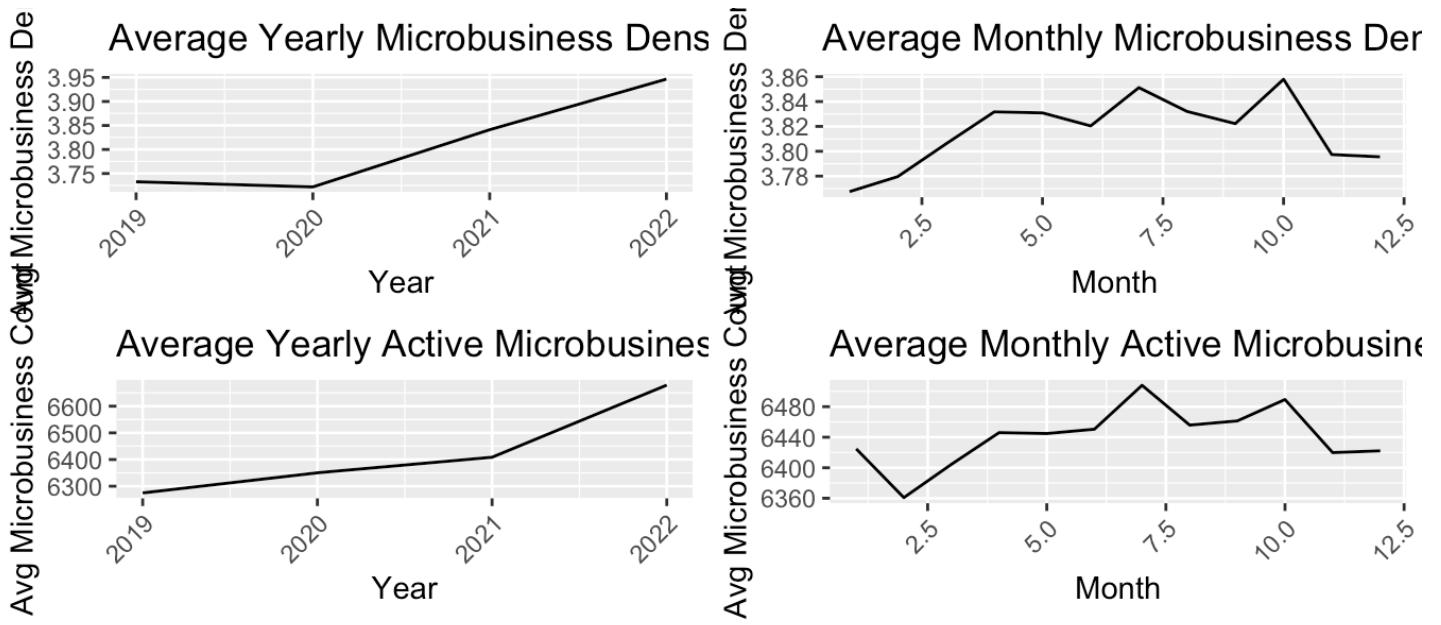
# Plot the yearly mean values for microbusiness density
p2 <-
  ggplot(merged_df_mean_year, aes(x = year, y = avg_microbusiness_density)) +
  geom_line() +
  ggtitle("Average Yearly Microbusiness Density") +
  xlab("Year") +
  ylab("Avg Microbusiness Density") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

# Plot the monthly mean values for active
p3 <-
  ggplot(merged_df_mean_month, aes(x = month, y = avg_active)) +
  geom_line() +
  ggtitle("Average Monthly Active Microbusiness Count") +
  xlab("Month") +
  ylab("Avg Microbusiness Count") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

# Plot the yearly mean values for active
p4 <-
  ggplot(merged_df_mean_year, aes(x = year, y = avg_active)) +
  geom_line() +
  ggtitle("Average Yearly Active Microbusiness Count") +
  xlab("Year") +
  ylab("Avg Microbusiness Count") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))

# Display the plots side by side
grid.arrange(p2, p1, p4, p3, nrow = 2, ncol = 2)

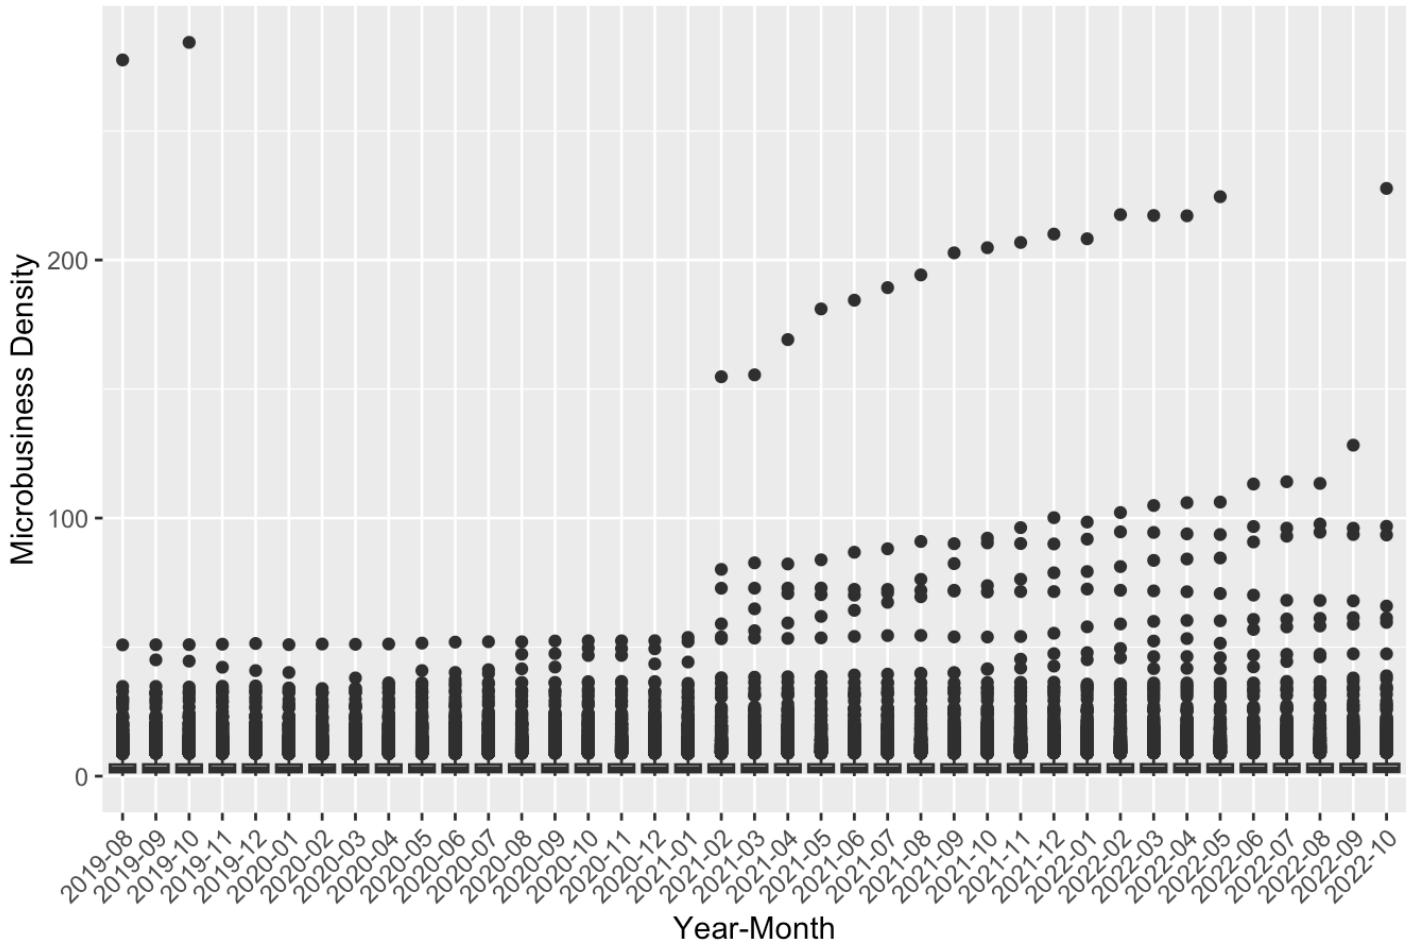
```



The left plots show that the average microbusiness density has increased slightly over the years, starting at approximately 3.73 in 2019 and reaching 3.94 in 2022. On the other hand, the average active count has also increased, starting at approximately 6274 in 2019 and reaching 6679 in 2022. In comparison, the right plots show fluctuations in the monthly averages for both variables. Generally, it follows a slightly upward trend over the year, with some peak values observed in July and October for the microbusiness density and active count, respectively. These peak values may represent seasonal variations, indicating that microbusinesses are more active during certain months. Overall, the plot shows some correlation between the monthly average values of microbusiness_density and active count, indicating that common factors may influence both variables.

```
# Create a grid of box plots
ggplot(merged_df, aes(x=year_month, y=microbusiness_density)) +
  geom_boxplot() +
  labs(x="Year-Month", y="Microbusiness Density") +
  ggtitle("Overall Microbusiness Density Average") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

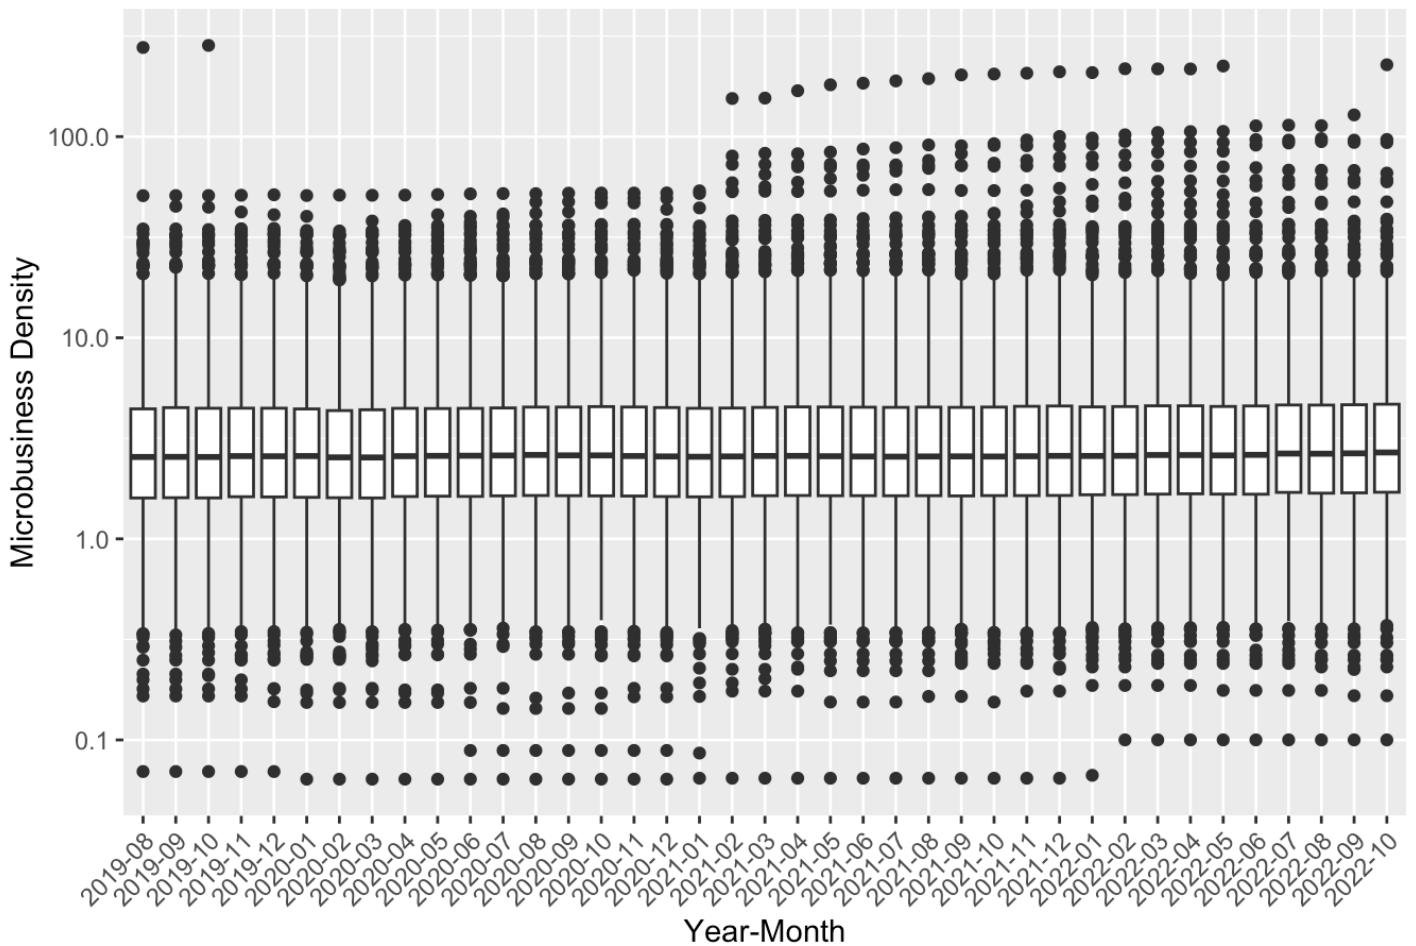
Overall Microbusiness Density Average



Above plot is not very informative because the small values are obscured by the larger ones. Therefore, using a logarithmic scale on the Y-axis can help to reduce this distortion and provide a more informative visualization of the data.

```
# Create a grid of box plots
ggplot(merged_df, aes(x=year_month, y=microbusiness_density)) +
  geom_boxplot() +
  scale_y_log10() +
  labs(x="Year-Month", y="Microbusiness Density") +
  ggtitle("Overall Microbusiness Density Average") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

Overall Microbusiness Density Average



3.7.2. Economic Regional Divisions

The Bureau of Economic Analysis (BEA) divides the United States into eight distinct economic regions¹.

These regions are based on similarities in economic characteristics such as industry composition, income levels, and employment patterns. The eight regions are:

1. **New England:** Connecticut, Maine, Massachusetts, New Hampshire, Rhode Island, and Vermont.
The economy in this region is largely based on manufacturing, healthcare, education, and finance.
2. **Mideast:** Delaware, Maryland, New Jersey, New York, Pennsylvania, and the District of Columbia.
The region has a diverse economy, with a mix of manufacturing, finance, healthcare, and professional services.
3. **Great Lakes:** Illinois, Indiana, Michigan, Ohio, and Wisconsin.
The region has a strong manufacturing base, particularly in the automotive industry, and also has a significant healthcare sector.
4. **Plains:** Iowa, Kansas, Minnesota, Missouri, Nebraska, North Dakota, and South Dakota.
Agriculture and energy production are major industries in this region, along with manufacturing and healthcare.
5. **Southeast:** Alabama, Arkansas, Florida, Georgia, Kentucky, Louisiana, Mississippi, North Carolina, South Carolina, Tennessee, Virginia, and West Virginia.

The Southeast has a diverse economy, with significant industries in healthcare, finance, and manufacturing, as well as tourism and agriculture.

6. **Southwest:** Arizona, New Mexico, Oklahoma, and Texas.

The region has a strong energy sector, particularly in oil and gas production, and also has significant industries in manufacturing, healthcare, and finance.

7. **Rocky Mountain:** Colorado, Idaho, Montana, Utah, and Wyoming.

The region is known for its natural resources, particularly in mining and energy production, as well as tourism, healthcare, and manufacturing.

8. **Far West:** Alaska, California, Hawaii, Nevada, Oregon, and Washington.

This region has a diverse economy, with significant industries in technology, finance, healthcare, and manufacturing, as well as tourism and agriculture.

To draw the map for the BEA regions, first, we need to convert state and county columns in `merged_df` to lowercase letters. Merging two dataframes will cause problems because the data from `map_data()` will be in lowercase letters.

```
# Convert state and county columns in merged_df to lowercase
merged_df <- merged_df %>%
  mutate(state = tolower(state)) %>%
  mutate(county = tolower(county))
```

Then, we'll create a new column in `merged_df` named `region` and assign region values based on `state` column:

```

# Create a new column named region and initialize all values as NA
merged_df$region <- NA

# Assign region values based on state column
for (i in 1:nrow(merged_df)) {
  if (merged_df$state[i] %in% c("connecticut", "maine", "massachusetts", "new hampshire", "rhode island", "vermont")) {
    merged_df$region[i] <- "new england"
  } else if (merged_df$state[i] %in% c("delaware", "maryland", "new jersey", "new york", "pennsylvania", "district of columbia")) {
    merged_df$region[i] <- "mideast"
  } else if (merged_df$state[i] %in% c("illinois", "indiana", "michigan", "ohio", "wisconsin")) {
    merged_df$region[i] <- "great lakes"
  } else if (merged_df$state[i] %in% c("iowa", "kansas", "minnesota", "missouri", "nebraska", "north dakota", "south dakota")) {
    merged_df$region[i] <- "plains"
  } else if (merged_df$state[i] %in% c("alabama", "arkansas", "florida", "georgia", "kentucky", "louisiana", "mississippi", "north carolina", "south carolina", "tennessee", "virginia", "west virginia")) {
    merged_df$region[i] <- "southeast"
  } else if (merged_df$state[i] %in% c("arizona", "new mexico", "oklahoma", "texas")) {
    merged_df$region[i] <- "southwest"
  } else if (merged_df$state[i] %in% c("colorado", "idaho", "montana", "utah", "wyoming")) {
    merged_df$region[i] <- "rocky mountain"
  } else if (merged_df$state[i] %in% c("alaska", "california", "hawaii", "nevada", "oregon", "washington")) {
    merged_df$region[i] <- "far west"
  } else {
    merged_df$region[i] <- "other"
  }
}

```

```

# Print all the unique values in the region column
unique(merged_df$region)

```

```

## [1] "southeast"      "far west"        "southwest"       "rocky mountain"
## [5] "new england"     "mideast"         "great lakes"     "plains"

```

Now that the data in the dataframe matches the `map_data()` output, we appoint each state to the region it belongs to and then use `ggplot()` to draw the map:

```

# Get the map of the United States
us_map <- map_data("state")

# Create a lookup table for state abbreviations and their corresponding full names
state_names <- data.frame(state = state.abb, name = tolower(state.name))

# Map the regions to the states

```

```

region_map <- us_map %>%
#left_join(state_names, by = c("region" = "state")) %>%
  left_join(state_names, by = c("region" = "name")) %>%
# merge(us_map, state_names, by.x=c("region"), by.y=c("name")) %>%
  mutate(region =
    ifelse(region %in% c("connecticut", "maine", "massachusetts", "new hampshire", "rhode island", "vermont"), "New England",
          ifelse(region %in% c("delaware", "maryland", "new jersey", "new york", "pennsylvania", "district of columbia"), "Mideast",
                ifelse(region %in% c("illinois", "indiana", "michigan", "ohio", "wisconsin"), "Great Lakes",
                      ifelse(region %in% c("iowa", "kansas", "minnesota", "missouri", "nebraska", "north dakota", "south dakota"), "Plains",
                            ifelse(region %in% c("alabama", "arkansas", "florida", "georgia", "kentucky", "louisiana", "mississippi", "north carolina", "south carolina", "tennessee", "virginia", "west virginia"), "Southeast",
                                  ifelse(region %in% c("arizona", "new mexico", "oklahoma", "texas"), "Southwest",
                                        ifelse(region %in% c("colorado", "idaho", "montana", "utah", "wyoming"), "Rocky Mountain",
                                              ifelse(region %in% c("alaska", "california", "hawaii", "nevada", "oregon", "washington"), "Far West", NA
                                                )))))))))

```

```

# Summarize the data to get the center coordinates of each state
#state_centers <- region_map %>%
#  group_by(state) %>%
#  summarise(long = mean(long), lat = mean(lat))
# add labels
states <- aggregate(cbind(long, lat) ~ region, data=us_map,
                     FUN=function(x)mean(range(x)))
states$group <- c("AL", "AR", "AZ", "CA", "CO", "CT", "DE", "DC", "FL", "GA", "IA",
                 "ID", "IL", "IN", "KS", "KY", "LA", "MA", "MD", "ME", "MI", "MN",
                 "MO", "MS", "MT", "NC", "ND", "NE", "NH", "NJ", "NM", "NV", "NY",
                 "OH", "OK", "OR", "PA", "RI", "SC", "SD", "TN", "TX", "UT", "VA",
                 "VT", "WA", "WI", "WV", "WY")

# names(states)[names(states) == "region"] <- "group"

#Plot the map
ggplot(region_map, aes(x = long, y = lat, group = group, fill = region)) +
  geom_polygon(color = "black", show.legend = TRUE) +
#  geom_text(aes(label = state), data = region_map, size = 3, vjust = 2, hjust = 2) +
#  geom_text(aes(label = state), data = state_centers, size = 2, vjust = 2, hjust = 2) +
  geom_text(data = states, aes(long, lat, label = group), size = 2.5, inherit.aes = FALSE, color = "white", fontface = "bold") +
#  scale_fill_gradient(low = "white", high = "darkred") +
#  scale_fill_manual(values = viridis(n = 60), na.value = "gray") +

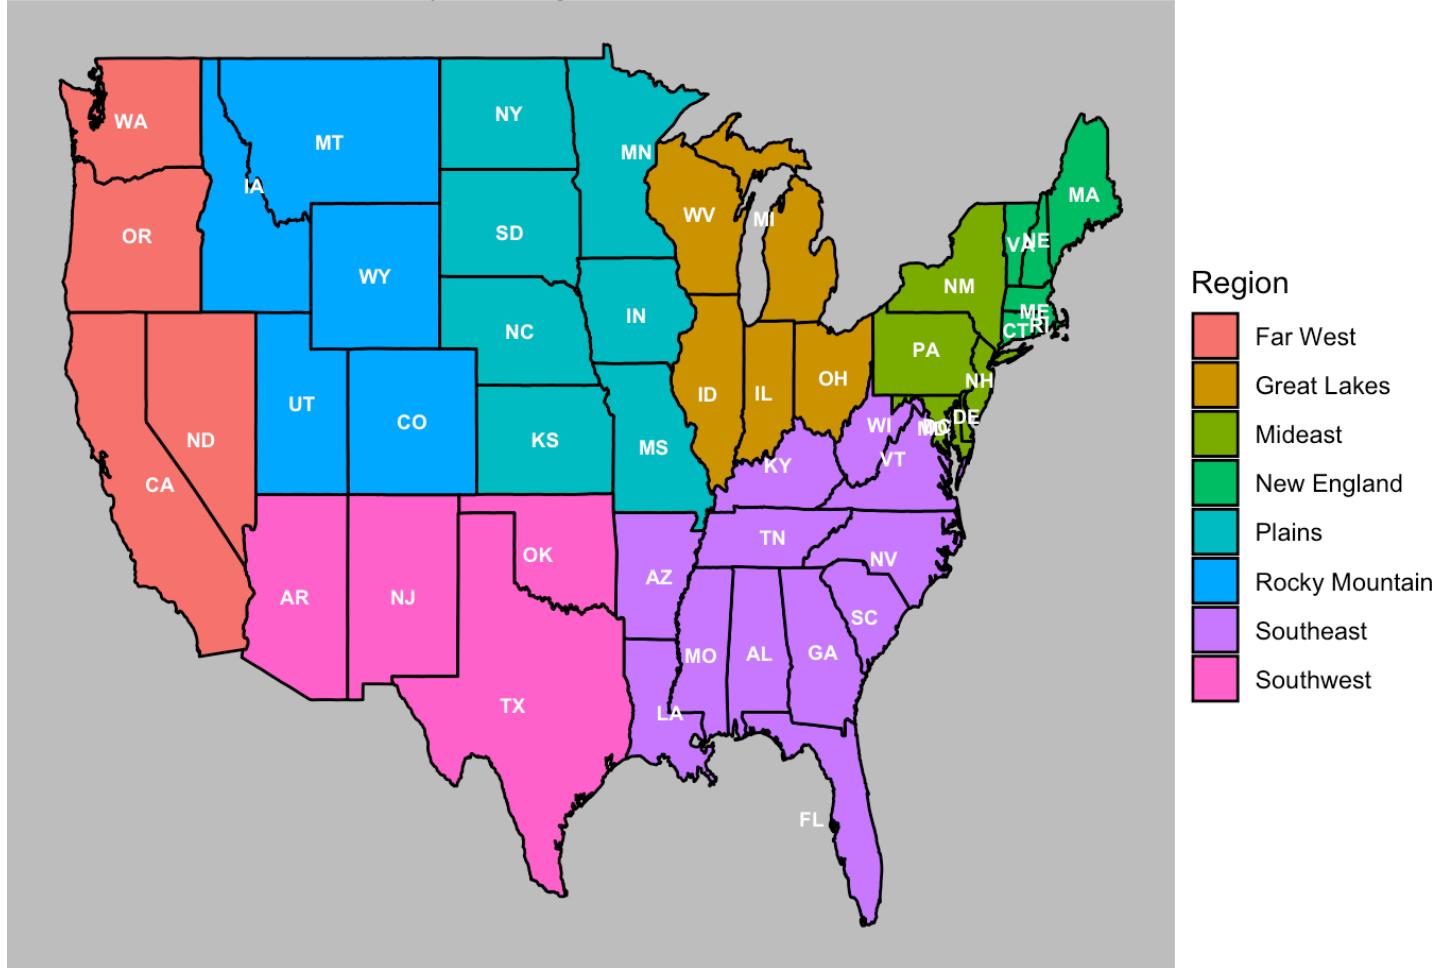
```

```

  labs(title = "Bureau of Economic Analysis Regional Divisions Map", fill = "Region")
  +
  # geom_text(aes(x = long, y = lat, label = state), data = state_centers, size = 3
  , color = "white") +
  theme_void() +
  theme(panel.background = element_rect(fill = "gray75", color = NA))

```

Bureau of Economic Analysis Regional Divisions Map



```

# Group merged_df by region and calculate average microbusiness density
merged_df %>%
  group_by(region) %>%
  summarize(avg_density = mean(microbusiness_density)) %>%

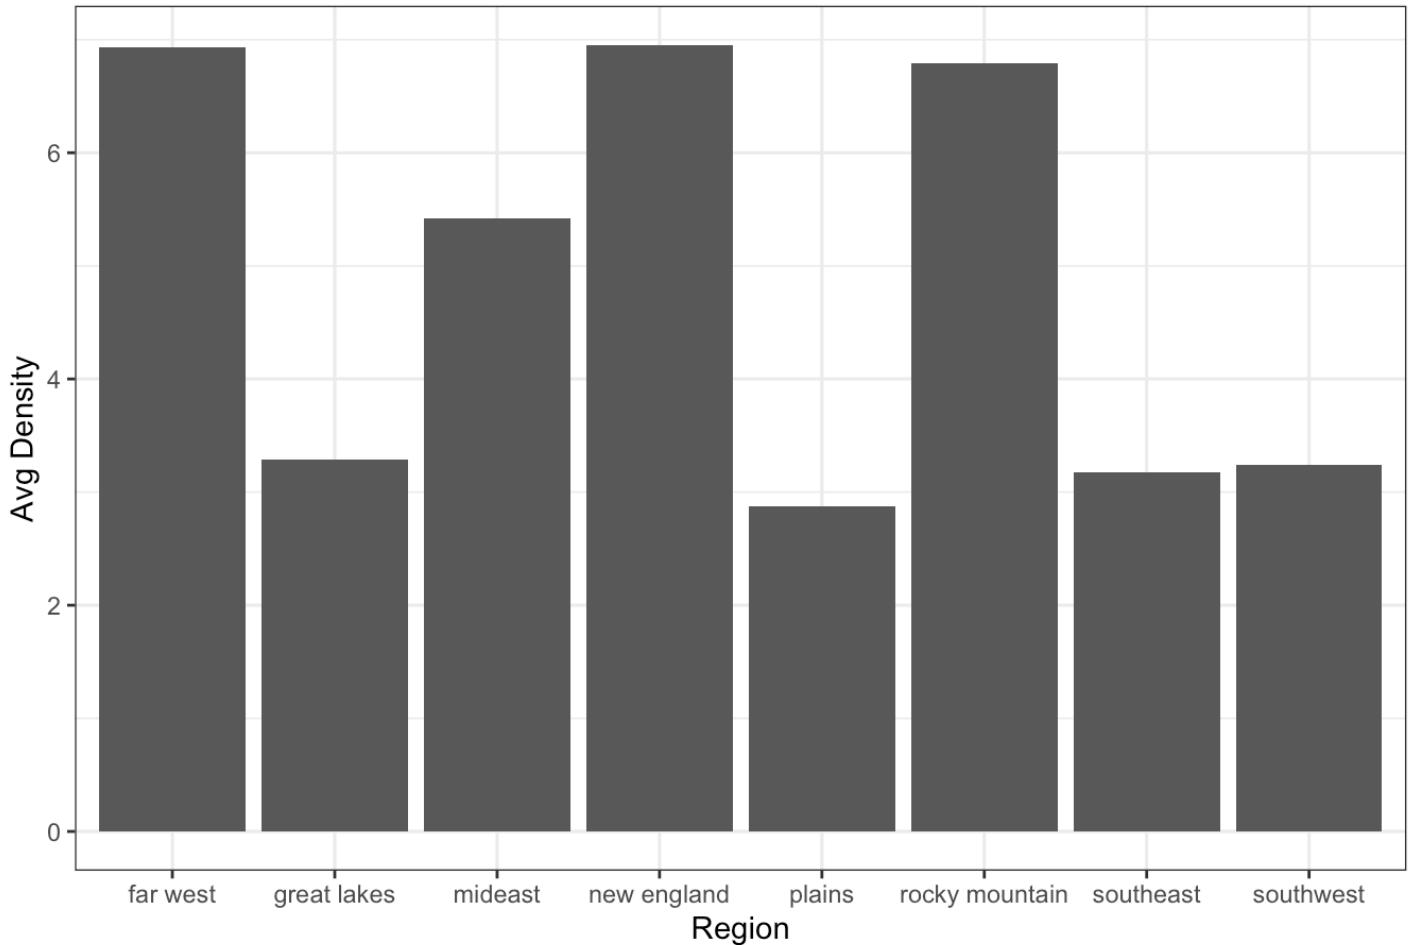
# Create bar plot of average density by region
ggplot(aes(x = region, y = avg_density)) +
  geom_bar(stat = "identity") +
  theme_bw()

# Add plot title and axis labels
labs(title = "Average Microbusiness Density Per Region",
     x = "Region", y = "Avg Density") +
  theme_bw()

# Apply a black and white theme to the plot
theme_bw()

```

Average Microbusiness Density Per Region



According to the above plot *New England* has the highest average microbusiness density, followed by *Farwest* and *Rocky Mountain* respectively, with a tiny difference, valuing more than 6.75. In contrast, *plains* has the lowest average microbusiness density, followed by *Southeast* and *Southwest*, all valued under 3.25. We can use a choropleth map to get a better view on the above information.

```

# choropleth map
# Group merged_df by region and calculate average microbusiness density
avg_density <- merged_df %>%
  group_by(region) %>%
  summarize(avg_density = mean(microbusiness_density))

# Create a lookup table for state abbreviations and their corresponding full names
state_names <- data.frame(state = state.abb, name = tolower(state.name))

# Lowercase region column of region_map
region_map <- region_map %>%
  mutate(region = tolower(region))

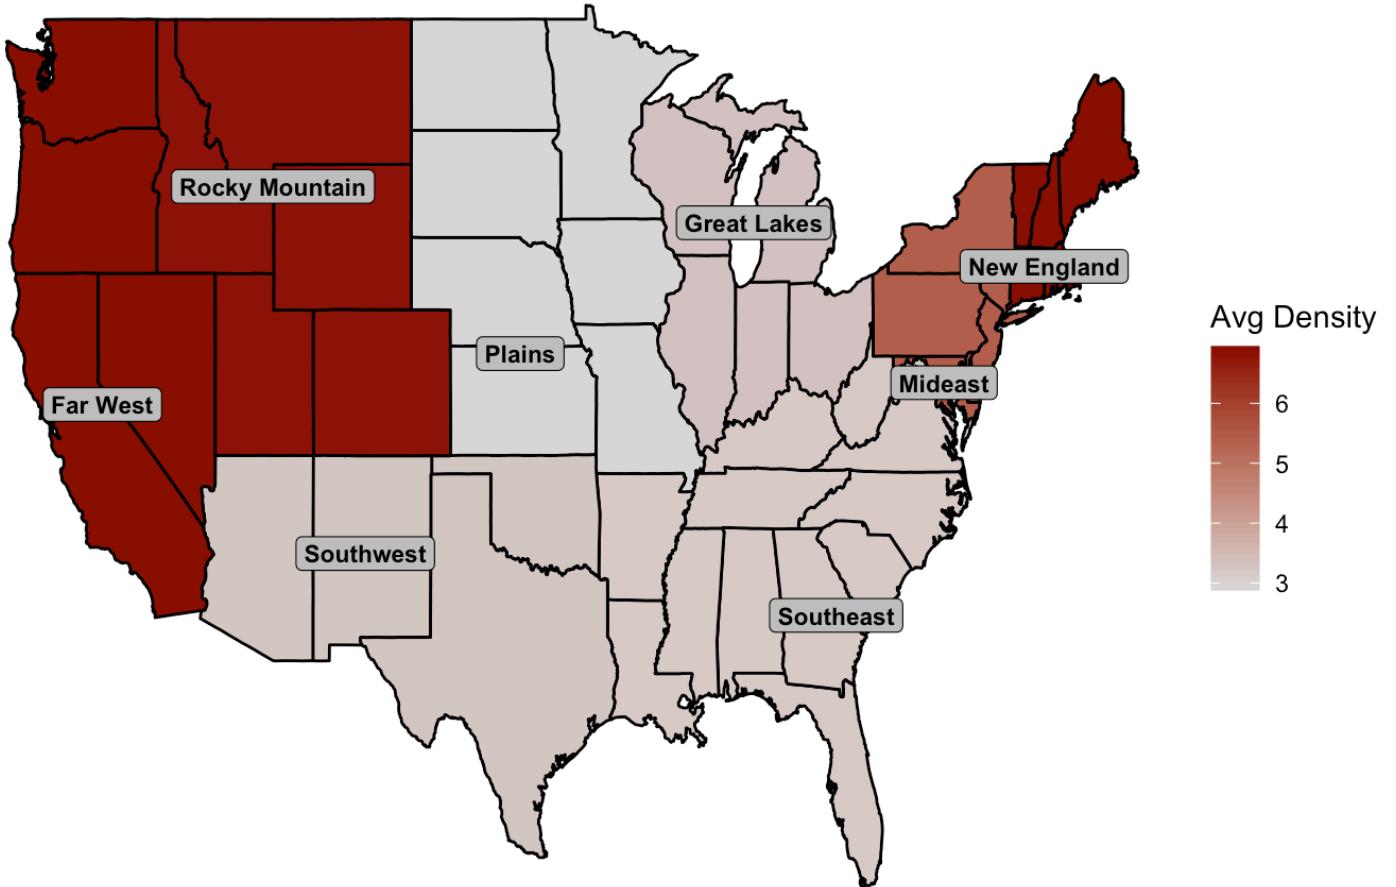
# Merge the average density data with the region_map data
plot_data <- merge(region_map, avg_density, by = "region") %>%
  arrange(order)

# Coordinates of the center of regions
bea_regions <- data.frame(
  group = c("New England", "Mideast", "Great Lakes", "Plains",
            "Southeast", "Southwest", "Rocky Mountain", "Far West"),
  x = c(-71.8, -76.9, -86.6, -98.5, -82.4, -106.4, -111.1, -119.8),
  y = c(42.2, 39, 43.4, 39.8, 32.6, 34.3, 44.4, 38.4)
)

# Create the plot
ggplot(plot_data, aes(x = long, y = lat, group = group, fill = avg_density)) +
  geom_polygon(color = "black") +
  geom_label(data = bea_regions,
             aes(x = x, y = y, label = group),
             size = 3, fontface = "bold",
             label.padding = unit(0.2, "lines"),
             label.size = 0.2,
             fill = "gray75", color = "black") +
  scale_fill_gradient(low = "gray85", high = "darkred") +
  # scale_fill_viridis(name = "Avg Density", na.value = "gray") +
  labs(title = "Average Microbusiness Density Per Region", fill = "Avg Density") +
  theme_void()

```

Average Microbusiness Density Per Region

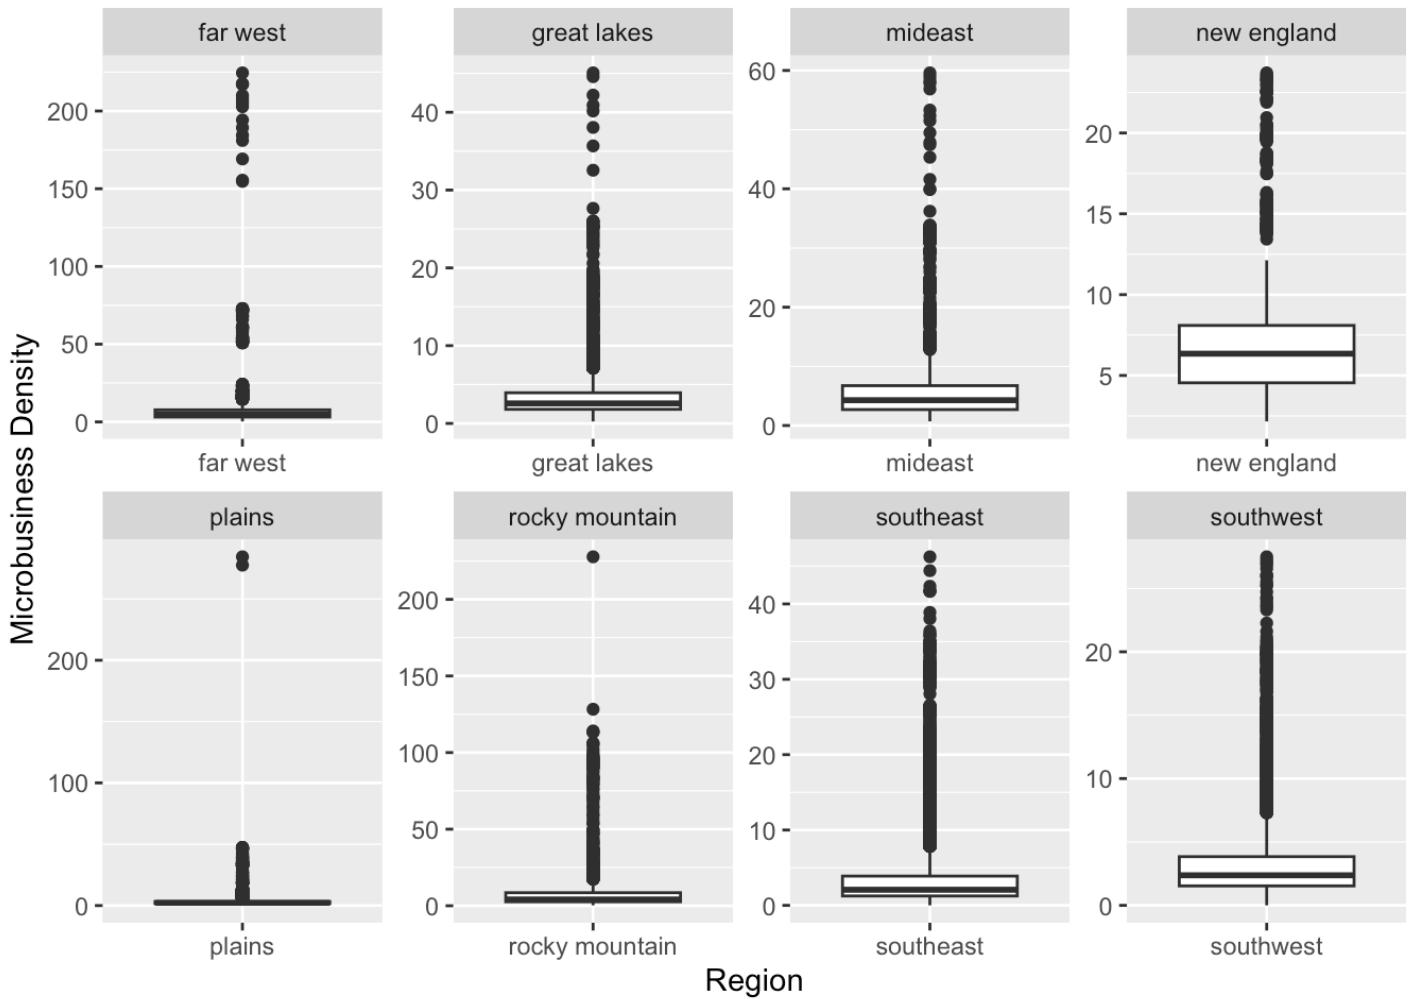


```
# theme(panel.background = element_rect(fill = "gray90", color = NA))
```

Although, we can only see only one parameter on above map. To have a better look on the distribution, central tendency, spread, and variability of the `microbusiness_density` variable, we can use boxplots.

```
# Aggregate data by region
df_by_region <- aggregate(microbusiness_density ~ region, merged_df, median)

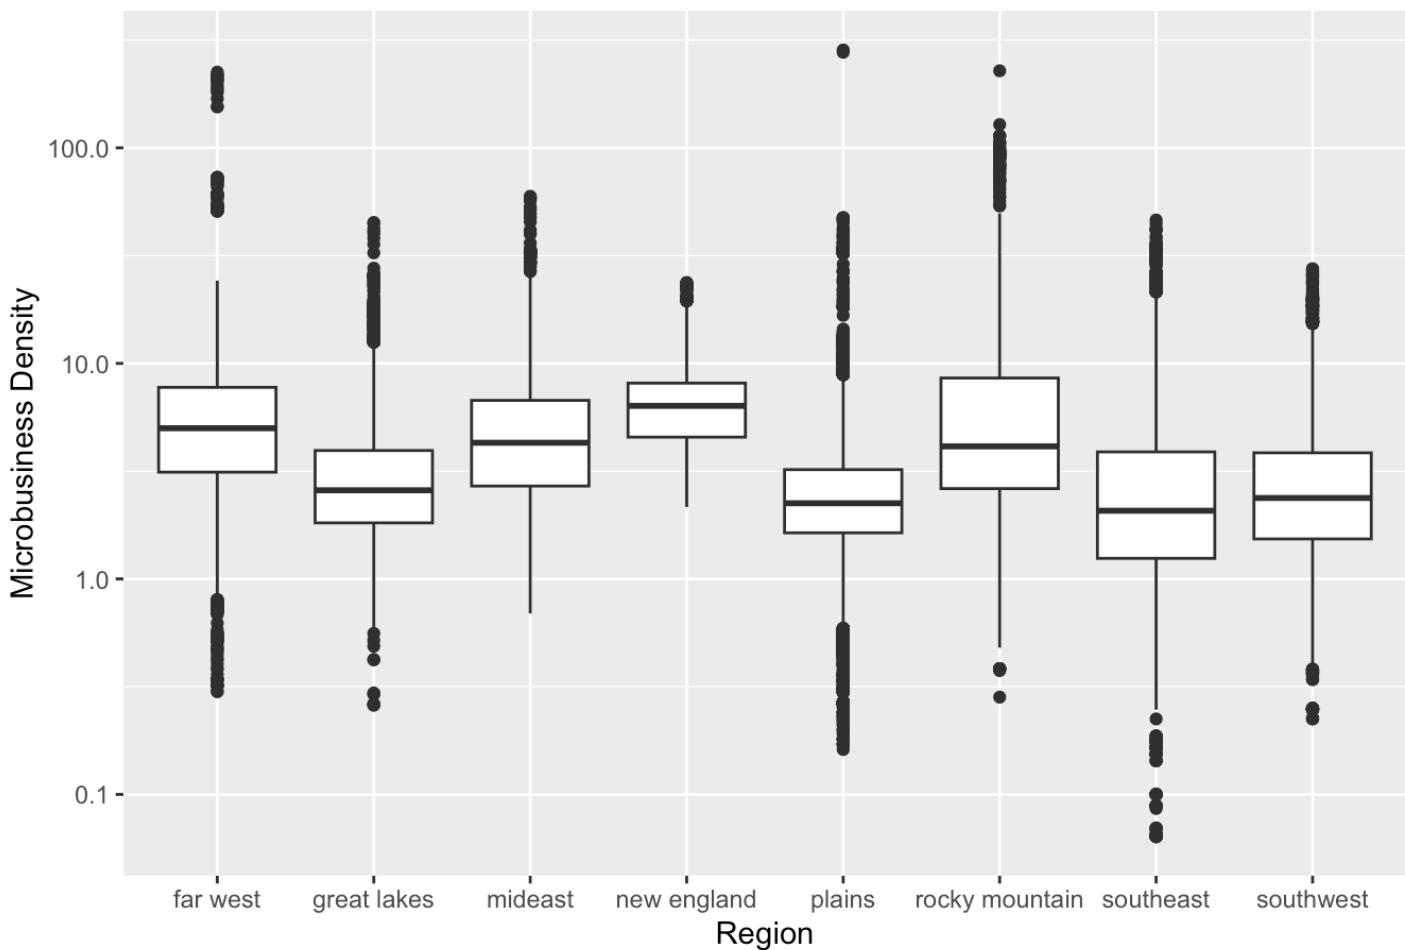
# Create boxplot
ggplot(merged_df, aes(x = region, y = microbusiness_density)) +
  geom_boxplot() +
  labs(x = "Region", y = "Microbusiness Density") +
  # Arrange plots in grid
  facet_wrap(~ region, scales = "free", nrow = 2)
```



Above plots are not very informative because the small values are obscured by the larger ones. Therefore, using a logarithmic scale on the Y-axis can help to reduce this distortion and provide a more informative visualization of the data.

```
# Create a grid of box plots
ggplot(merged_df, aes(x=region, y=microbusiness_density)) +
  geom_boxplot() +
  scale_y_log10() +
  labs(x="Region", y="Microbusiness Density") +
  ggtitle("Microbusiness Density by Region")
```

Microbusiness Density by Region



These boxplots are more informative, because using a logarithmic scale on the Y-axis helps to better reveal the differences and similarities between regions that helps to highlight any potential patterns or trends in the data. Some of the points that can be inferred from this boxplot include:

- All regions have at least some microbusiness activity. Because minimum microbusiness density is greater than 0 in all regions.
- The median microbusiness density is highest in the *New England* region, followed by the *Far West* and the *Mideast* regions.
- The 3rd quartile microbusiness density is highest in the *Rocky Mountain* region, followed by the *New England* and the *Far West* regions.
- The maximum microbusiness density is highest in the *Plains* region, followed by the *Rocky Mountain* and the *Far West* regions.
- The mean microbusiness density is highest in the *New England* region, followed by the *Far West* and the *Rocky Mountain* regions.
- The interquartile range (IQR = the difference between the 1st and 3rd quartiles) of microbusiness density is widest in the *Rocky Mountain* region, indicating that there is a greater range of microbusiness density in that region. In contrast, the IQR is narrowest in the *Plains* region.

```
# by(merged_df$microbusiness_density, merged_df$region, summary)
```

```

# Aggregate microbusiness density by state
state_avg <- aggregate(microbusiness_density ~ state, data = merged_df, FUN = mean
)

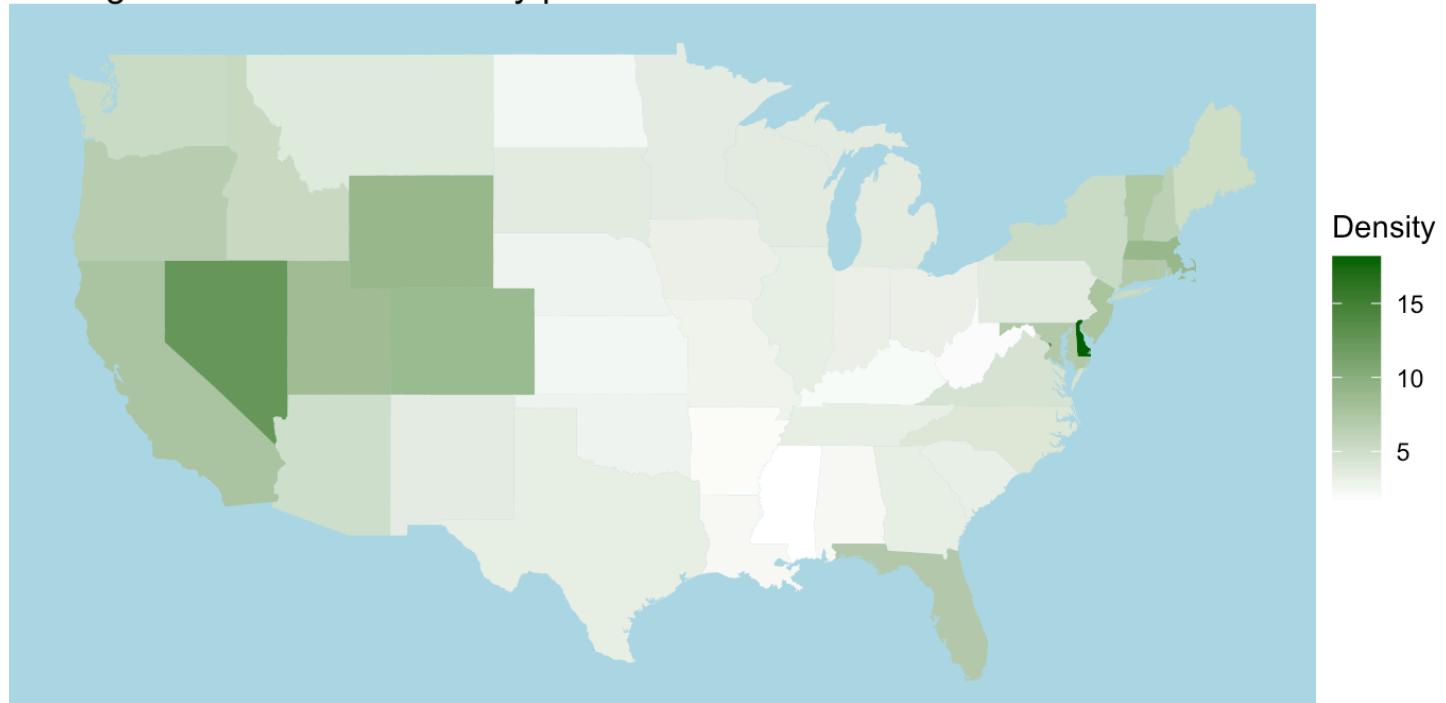
# Load US map data
us_map <- map_data("state")

# Merge state_avg with us_map based on region and state
map_data <- merge(us_map, state_avg, by.x = "region", by.y = "state")

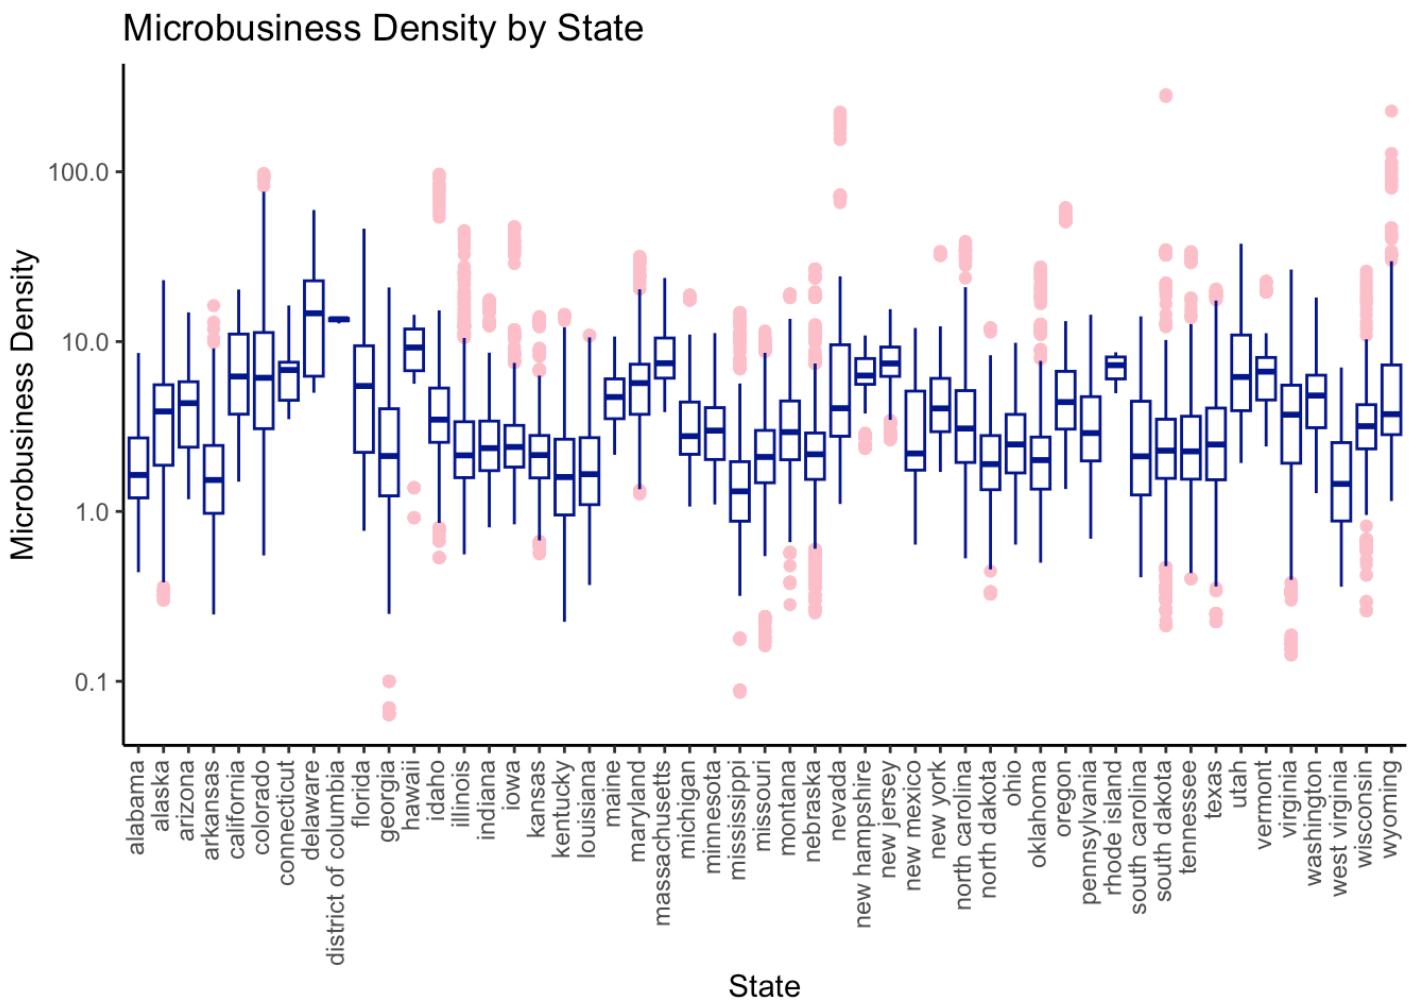
# Create a heatmap of microbusiness density by state
ggplot(map_data, aes(x = long, y = lat, group = group, fill = microbusiness_densit
y)) +
  geom_polygon() +
  scale_fill_gradient(low = "white", high = "darkgreen") +
  coord_map() +
  labs(title = "Average Microbusiness Density per State", fill = "Density") +
  theme_void() +
  theme(panel.background = element_rect(fill = "lightblue", color = NA))

```

Average Microbusiness Density per State



```
# Create a grid of box plots
ggplot(merged_df, aes(x=state, y=microbusiness_density)) +
  geom_boxplot(colour = "darkblue", outlier.colour = "pink") +
  scale_y_log10() +
  labs(x="State", y="Microbusiness Density") +
  ggtitle("Microbusiness Density by State") +
  theme_classic() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust = 0.3))
```



```

# {r fig.width = 10 ,fig.height = 12, out.width='100%', fig.align='center'}
# Aggregate microbusiness density by county

#county_avg <- merged_df %>%
#  group_by(cfips, county) %>%
#  summarise(microbusiness_density = mean(microbusiness_density))

county_avg <- aggregate(microbusiness_density ~ county + state, data = merged_df,
FUN = mean)

# Get rid of county, city, and parish in the end of county names
county_avg$county <- gsub(" county", "", county_avg$county)
county_avg$county <- gsub(" city", "", county_avg$county)
county_avg$county <- gsub(" parish", "", county_avg$county)

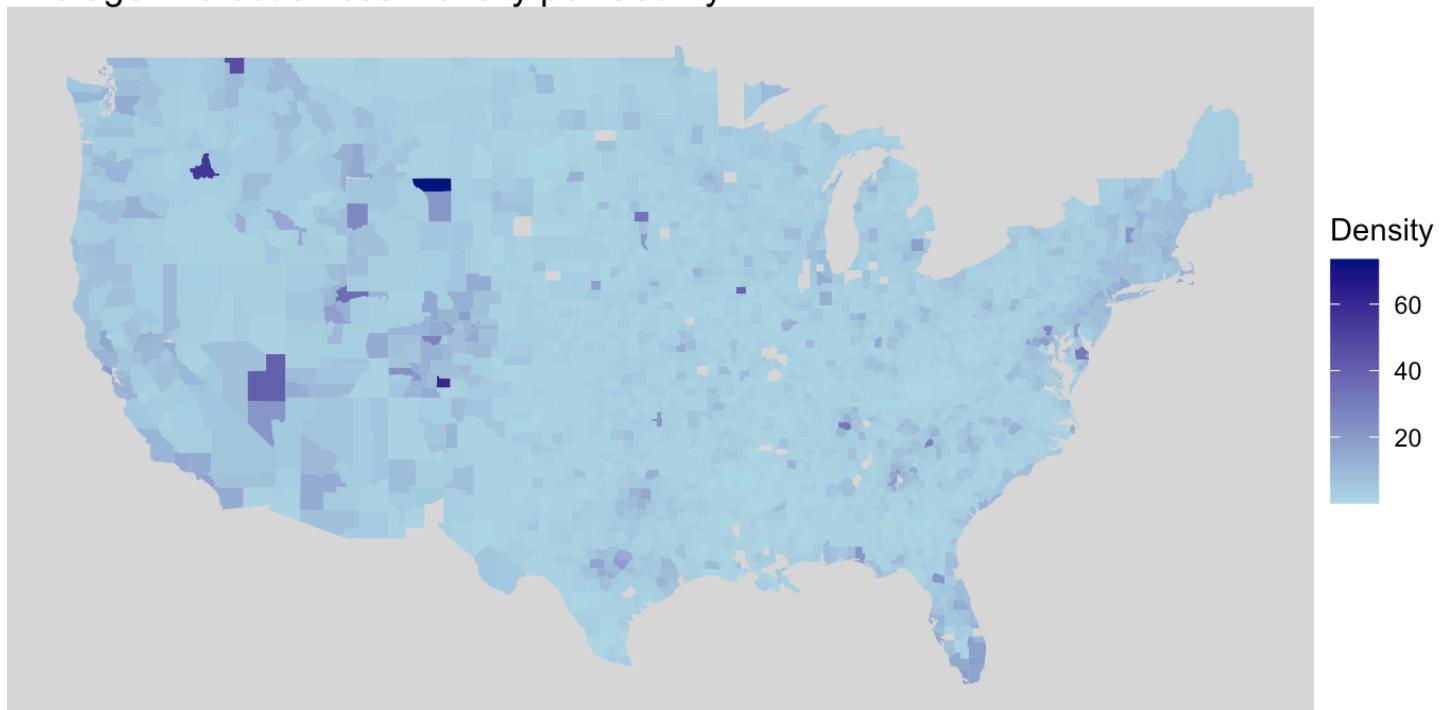
# Load US county map data
us_map <- map_data("county")

# Merge county_avg with us_map based on region and county
map_data <- merge(us_map, county_avg, by.x = c("subregion", "region"), by.y = c("c
ounty", "state")) %>%
  arrange(order)

# Create a heatmap of microbusiness density by county using ggplot2
ggplot(map_data, aes(x = long, y = lat, group = group, fill = microbusiness_densit
y)) +
  geom_polygon() +
  scale_fill_gradient(low = "lightblue", high = "navyblue") +
  coord_map() +
  labs(title = "Average Microbusiness Density per County", fill = "Density") +
  theme_void() +
  theme(panel.background = element_rect(fill = "gray85", color = NA))

```

Average Microbusiness Density per County



```
str(merged_df)
```

```

## 'data.frame': 122265 obs. of 16 variables:
## $ row_id : chr "1001_2019-08-01" "1001_2019-09-01" "1001_2019-1
0-01" "1001_2019-11-01" ...
## $ cfips : int 1001 1001 1001 1001 1001 1001 1001 1001 1001 100
1 ...
## $ county : chr "autauga county" "autauga county" "autauga count
y" "autauga county" ...
## $ state : chr "alabama" "alabama" "alabama" "alabama" ...
## $ first_day_of_month : Date, format: "2019-08-01" "2019-09-01" ...
## $ microbusiness_density: num 3.01 2.88 3.06 2.99 2.99 ...
## $ active : int 1249 1198 1269 1243 1243 1242 1217 1227 1255 125
7 ...
## $ year_month : chr "2019-08" "2019-09" "2019-10" "2019-11" ...
## $ year : num 2019 2019 2019 2019 2019 ...
## $ month : int 8 9 10 11 12 1 2 3 4 5 ...
## $ pct_bb : num 76.6 76.6 76.6 76.6 76.6 78.9 78.9 78.9 78.
9 ...
## $ pct_college : num 14.5 14.5 14.5 14.5 14.5 15.9 15.9 15.9 15.
9 ...
## $ pct_foreign_born : num 2.1 2.1 2.1 2.1 2 2 2 2 2 ...
## $ pct_it_workers : num 1.3 1.3 1.3 1.3 1.3 1.1 1.1 1.1 1.1 1.1 ...
## $ median_hh_inc : num 55317 55317 55317 55317 55317 ...
## $ region : chr "southeast" "southeast" "southeast" "southeast"
...
## - attr(*, "na.action")= 'omit' Named int [1:28551] 40 41 42 43 44 45 46 47 48
88 ...
## ... attr(*, "names")= chr [1:28551] "40" "41" "42" "43" ...

```

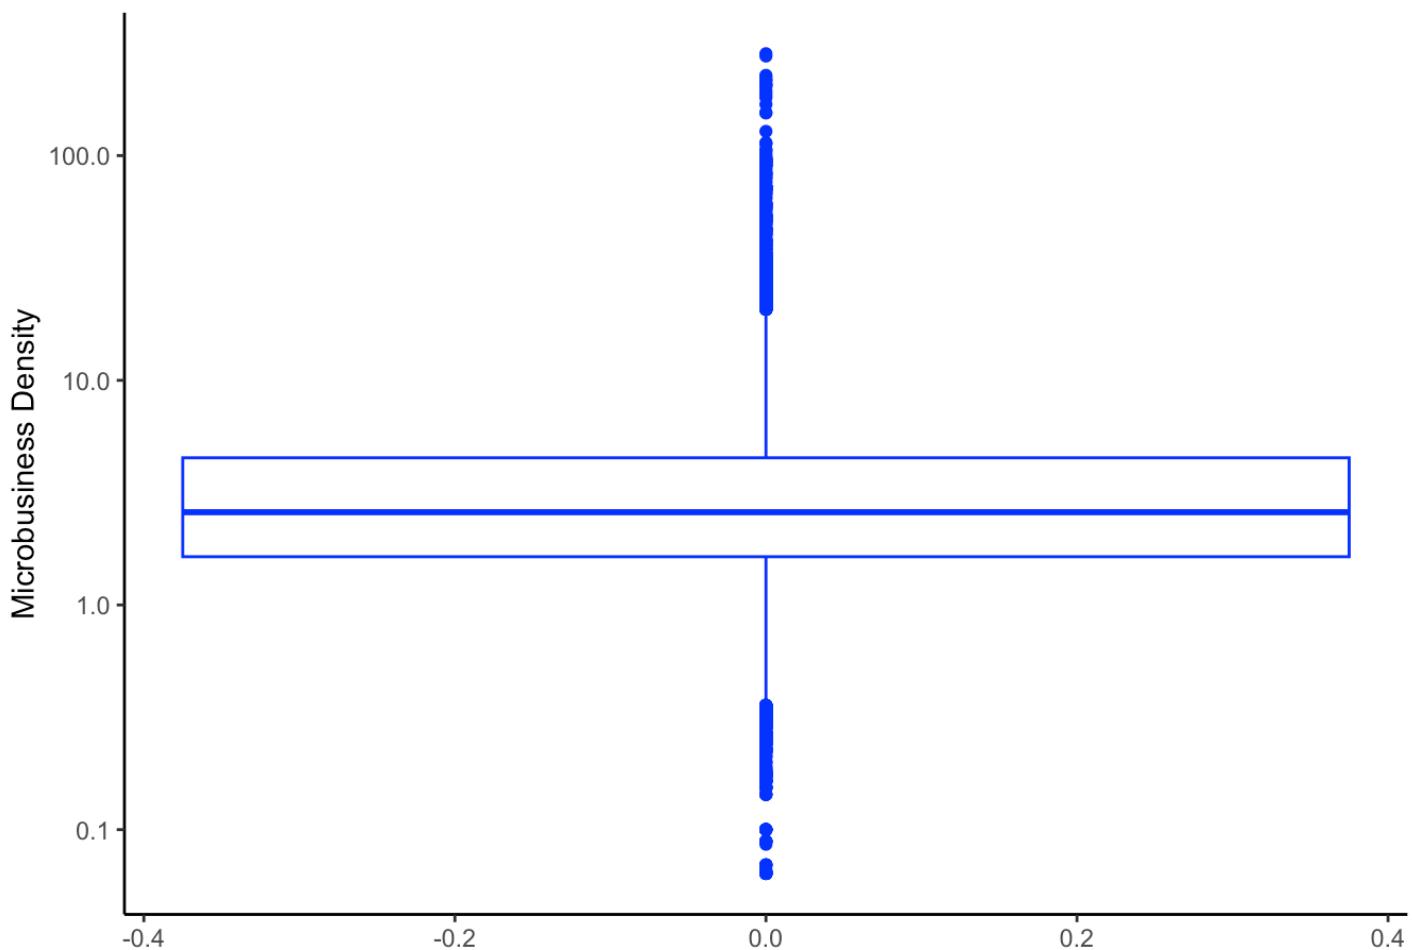
Boxplots are a visualization tool that provide insights into the central tendency and spread of a dataset, as well as identify outliers and skewness. They are useful for detecting anomalies and comparing variable distributions in a dataset, providing valuable insights into data distribution for exploratory data analysis.

```

ggplot(merged_df, aes(y = microbusiness_density)) +
# ggplot(merged_df, aes(x=region, y=microbusiness_density)) +
  geom_boxplot(colour = "blue") +
  scale_y_log10() +
  labs(y="Microbusiness Density") +
  ggtitle("Microbusiness Density") +
  theme_classic()

```

Microbusiness Density



```
# Create a grid of box plots
par(mfrow=c(3,2))
p1 <- ggplot(merged_df, aes(x=region, y=median_hh_inc)) +
  geom_boxplot(outlier.colour = "orange") +
  scale_y_log10() +
  labs(x="Region", y="Median Household Income") +
  ggtitle("Median Household Income by Region") +
  theme_classic() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

```
p2 <- ggplot(merged_df, aes(x=region, y=pct_college)) +
  geom_boxplot(outlier.colour = "orange") +
  scale_y_log10() +
  labs(x="Region", y="Percentage with College Education") +
  ggtitle("Percentage with College Education by Region") +
  theme_classic() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

```
p3 <- ggplot(merged_df, aes(x=region, y=pct_foreign_born)) +
  geom_boxplot(outlier.colour = "orange") +
  scale_y_log10() +
  labs(x="Region", y="Percentage of Foreign-born Residents") +
  ggtitle("Percentage of Foreign-born Residents by Region") +
  theme_classic() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

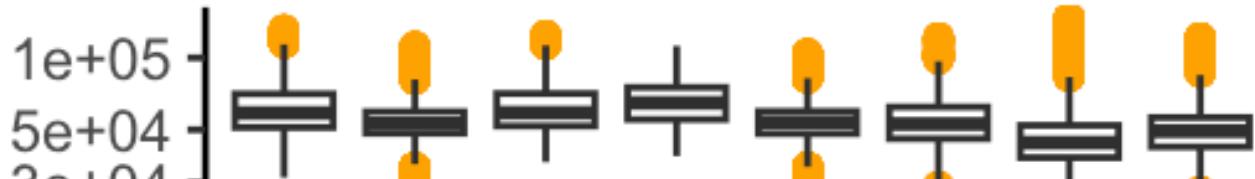
```
p4 <- ggplot(merged_df, aes(x=region, y=pct_it_workers)) +
  geom_boxplot(outlier.colour = "orange") +
  scale_y_log10() +
  labs(x="Region", y="Percentage of IT Workers") +
  ggtitle("Percentage of IT Workers by Region") +
  theme_classic() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

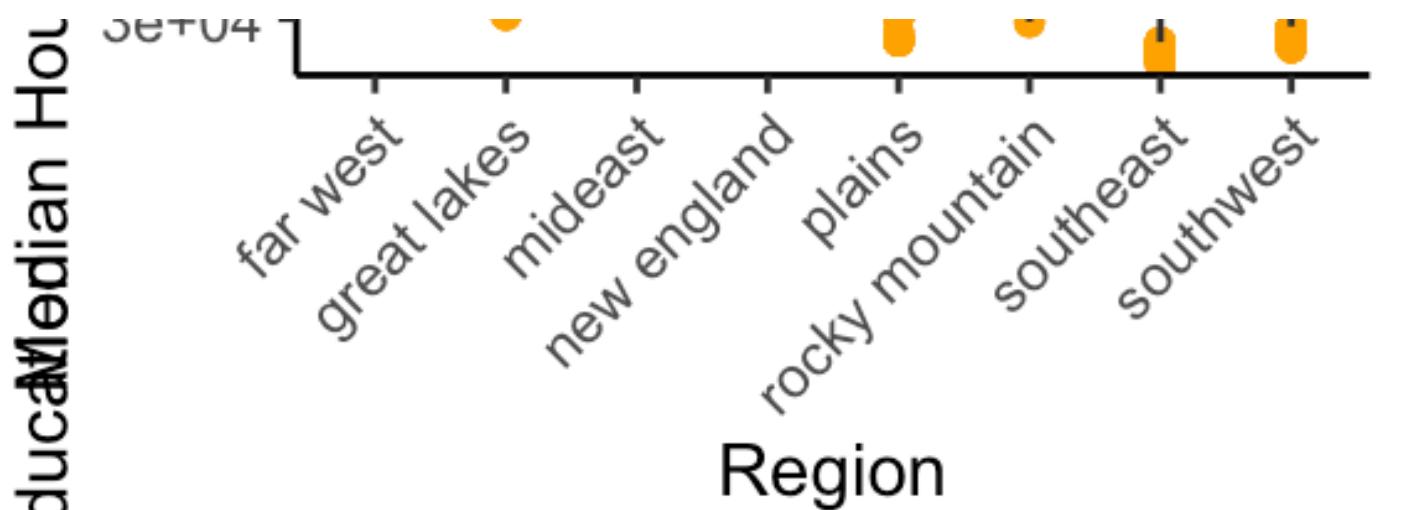
```
p5 <- ggplot(merged_df, aes(x=region, y=pct_bb)) +
  geom_boxplot(outlier.colour = "orange") +
  scale_y_log10() +
  labs(x="Region", y="Percentage of Broadband Access") +
  ggtitle("Percentage of Broadband Access by Region") +
  theme_classic() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

```
grid.arrange(p1, p2, p3, p4, p5, nrow = 5)
```

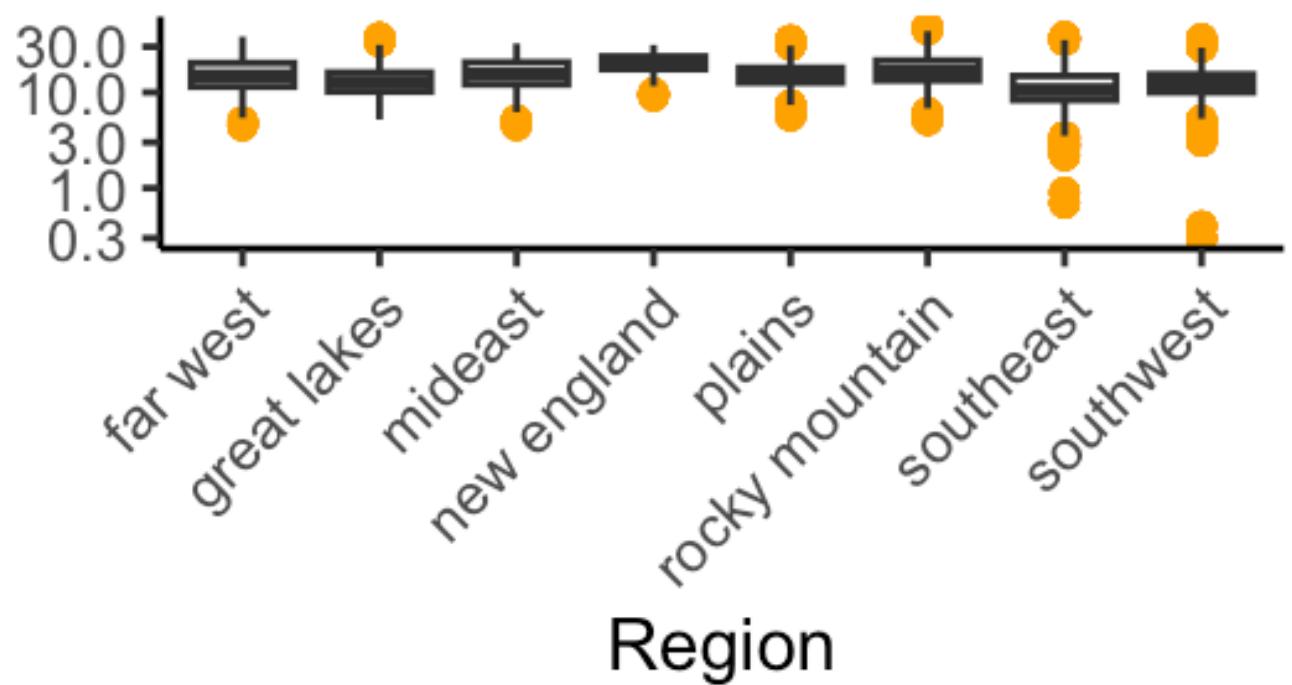
Household Inc

Median Household Income by Region

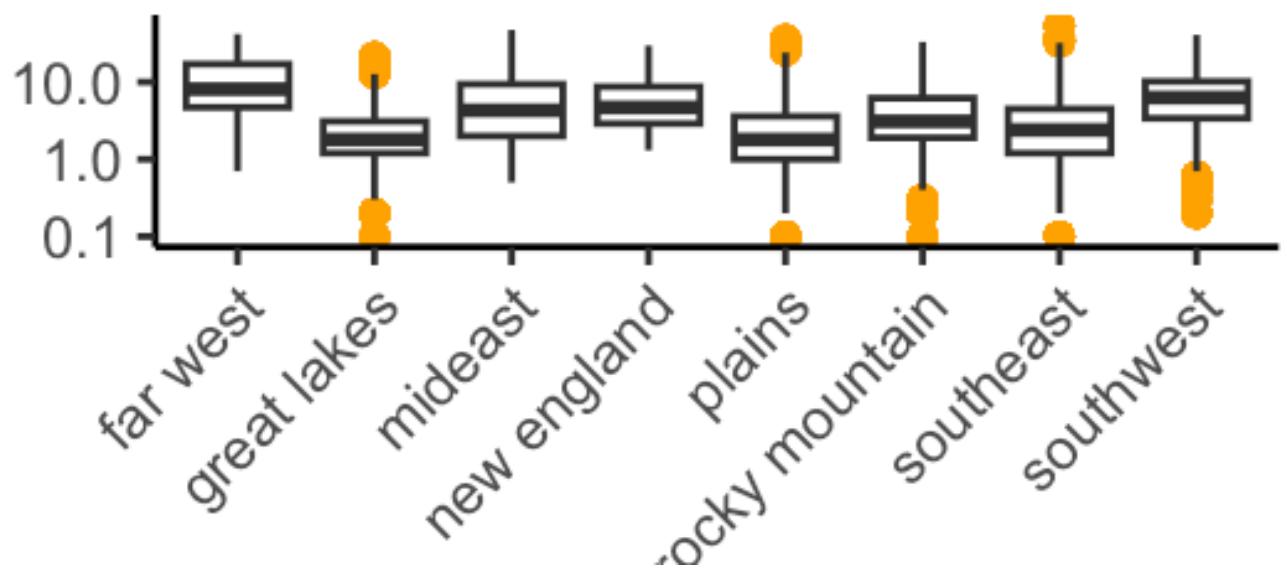


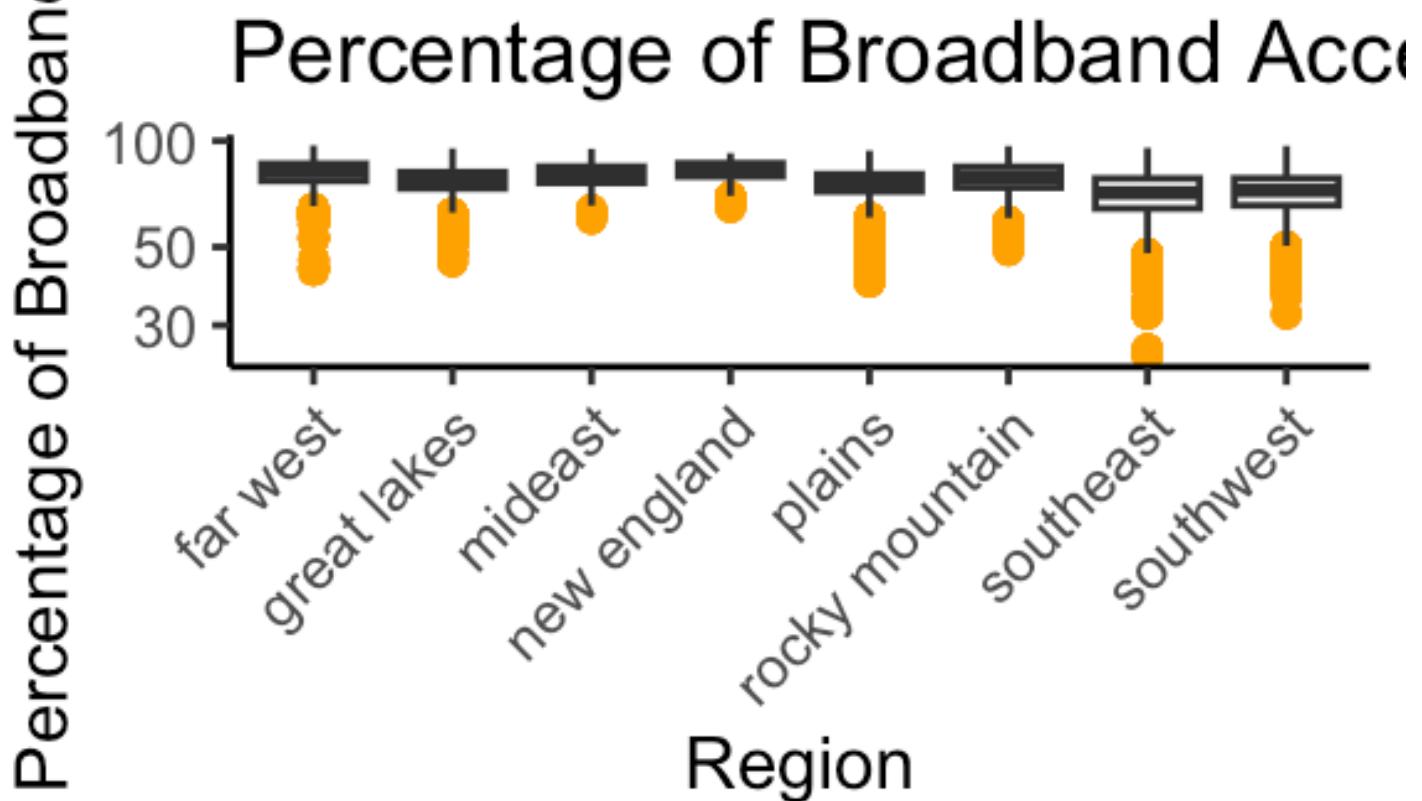
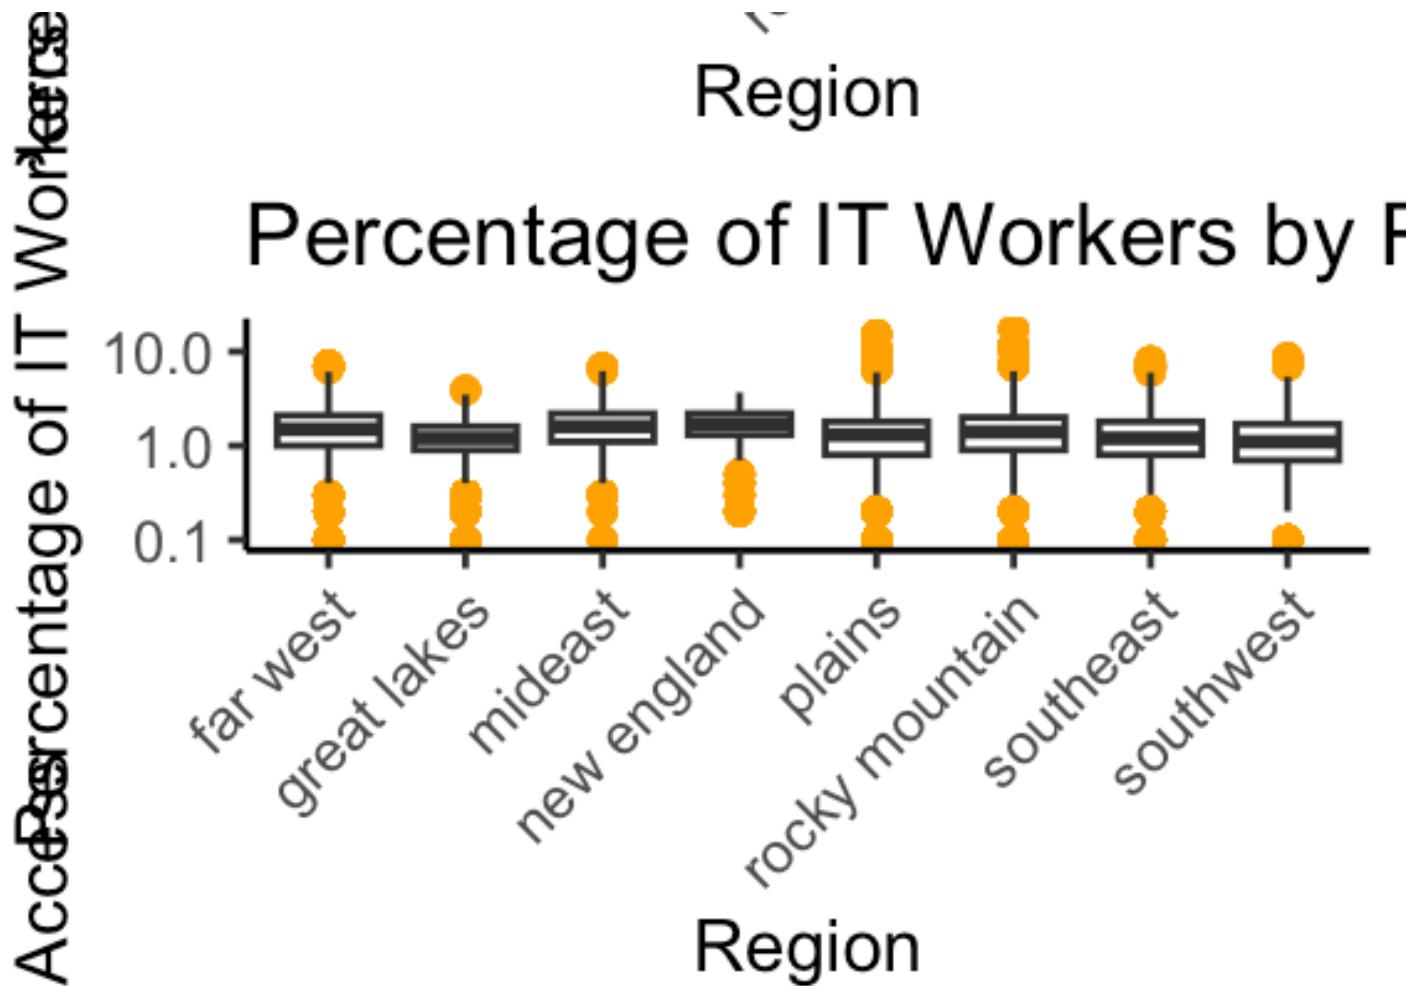


Percentage with College Education



Percentage of Foreign-born Residents





```

# Create a grid of box plots

p1 <- ggplot(merged_df, aes(x=state, y=median_hh_inc)) +
  geom_boxplot(outlier.colour = "orange") +
  scale_y_log10() +
  labs(x="State", y="Median Household Income") +
  ggtitle("Median Household Income by State") +
  theme_classic() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust = 0.3))

p2 <- ggplot(merged_df, aes(x=state, y=pct_college)) +
  geom_boxplot(outlier.colour = "orange") +
  scale_y_log10() +
  labs(x="State", y="Percentage with College Education") +
  ggtitle("Percentage with College Education by State") +
  theme_classic() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust = 0.3))

p3 <- ggplot(merged_df, aes(x=state, y=pct_foreign_born)) +
  geom_boxplot(outlier.colour = "orange") +
  scale_y_log10() +
  labs(x="State", y="Percentage of Foreign-born Residents") +
  ggtitle("Percentage of Foreign-born Residents by State") +
  theme_classic() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust = 0.3))

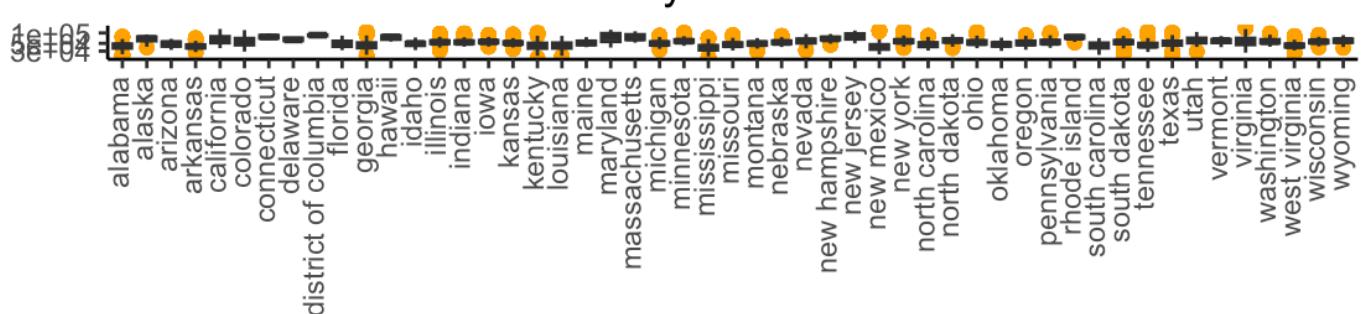
p4 <- ggplot(merged_df, aes(x=state, y=pct_it_workers)) +
  geom_boxplot(outlier.colour = "orange") +
  scale_y_log10() +
  labs(x="State", y="Percentage of IT Workers") +
  ggtitle("Percentage of IT Workers by State") +
  theme_classic() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust = 0.3))

p5 <- ggplot(merged_df, aes(x=state, y=pct_bb)) +
  geom_boxplot(outlier.colour = "orange") +
  scale_y_log10() +
  labs(x="State", y="Percentage of Broadband Access") +
  ggtitle("Percentage of Broadband Access by State") +
  theme_classic() +
  theme(axis.text.x = element_text(angle = 90, hjust = 1, vjust = 0.3))

grid.arrange(p1, p2, p3, p4, p5, nrow = 5)

```

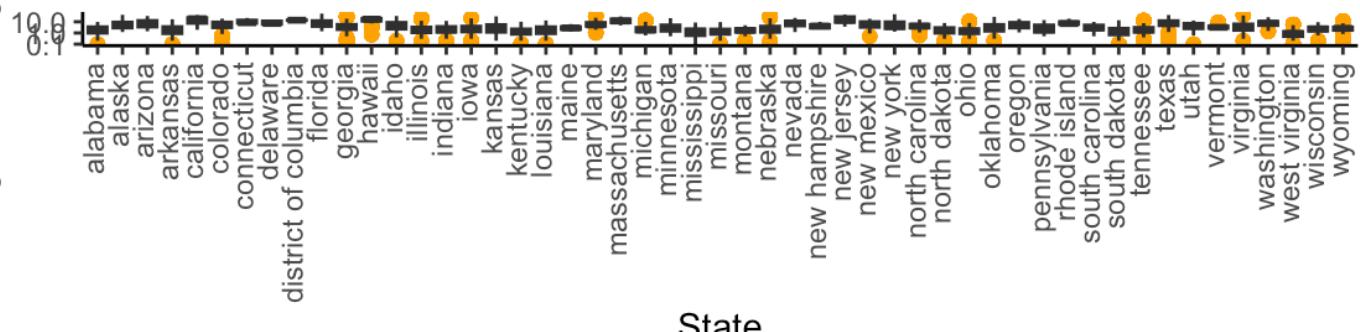
Median Household Income by State



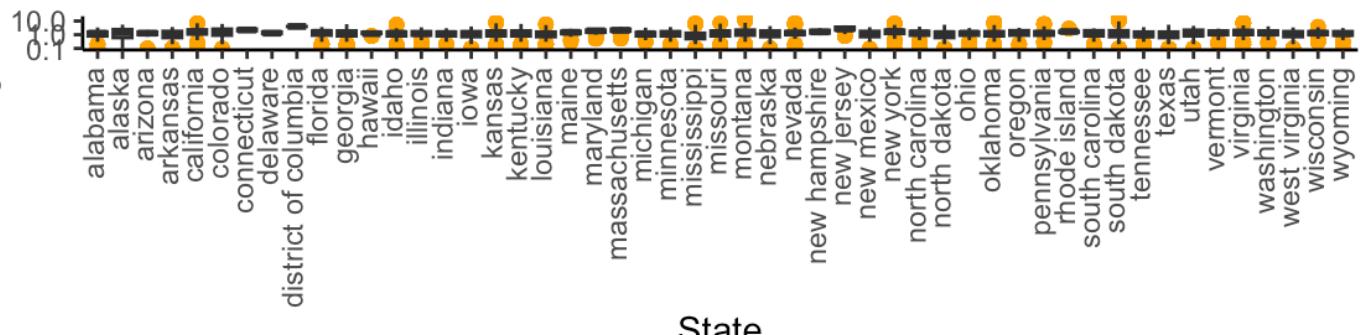
Percentage of Foreign-born Residents with College Education by State



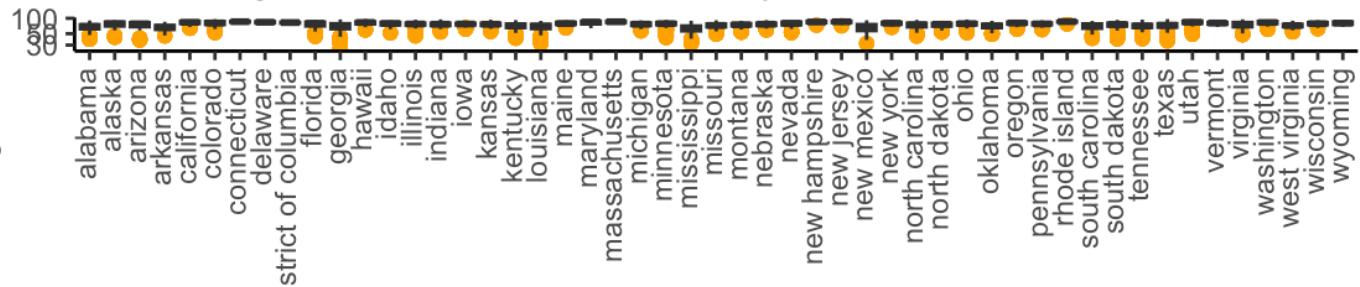
Percentage of Foreign-born Residents by State



Percentage of IT Workers by State



Percentage of Broadband Access by State



```
# choropleth map with 5 variables
```

```
# Group merged_df by region and calculate average variables
```

```
avg_variables <- merged_df %>%
```

```
group_by(region) %>%
```

```
summarize(avg_pct_bb = mean(pct_bb),
```

```
avg_pct_college = mean(pct_college),
```

```
avg_pct_foreign_born = mean(pct_foreign_born),
```

```

avg_pct_it_workers = mean(pct_it_workers),
avg_median_hh_inc = mean(median_hh_inc))

# Create a lookup table for state abbreviations and their corresponding full names
state_names <- data.frame(state = state.abb, name = tolower(state.name))

# Lowercase region column of region_map
region_map <- region_map %>%
mutate(region = tolower(region))

# Merge the average density data with the region_map data
plot_data <- merge(region_map, avg_variables, by = "region") %>%
arrange(order)

# Coordinates of the center of regions
bea_regions <- data.frame(
group = c("New England", "Mideast", "Great Lakes", "Plains",
"Southeast", "Southwest", "Rocky Mountain", "Far West"),
x = c(-71.8, -76.9, -86.6, -98.5, -82.4, -106.4, -111.1, -119.8),
y = c(42.2, 39, 43.4, 39.8, 32.6, 34.3, 44.4, 38.4)
)

# Create the plot with a grid of 5 rows and 1 column
p1 <- ggplot(plot_data, aes(x = long, y = lat, group = group, fill = avg_pct_bb)) +
geom_polygon(color = "black") +
geom_label(data = bea_regions,
aes(x = x, y = y, label = group),
size = 3, fontface = "bold",
label.padding = unit(0.2, "lines"),
label.size = 0.2,
fill = "gray75", color = "black") +
scale_fill_gradient(low = "gray85", high = "darkred") +
labs(title = "Average Percent of Broadband Access Per Region", fill = "Avg Density") +
theme_void()

p2 <- ggplot(plot_data, aes(x = long, y = lat, group = group, fill = avg_pct_college)) +
geom_polygon(color = "black") +
geom_label(data = bea_regions,
aes(x = x, y = y, label = group),
size = 3, fontface = "bold",
label.padding = unit(0.2, "lines"),
label.size = 0.2,
fill = "gray75", color = "black") +
scale_fill_gradient(low = "gray85", high = "darkred") +
labs(title = "Average Percent of College Graduates Per Region", fill = "Avg Percent") +
theme_void()

p3 <- ggplot(plot_data, aes(x = long, y = lat, group = group, fill = avg_pct_foreign_born)) +
geom_polygon(color = "black") +

```

```

geom_label(data = bea_regions,
           aes(x = x, y = y, label = group),
           size = 3, fontface = "bold",
           label.padding = unit(0.2, "lines"),
           label.size = 0.2,
           fill = "gray75", color = "black") +
scale_fill_gradient(low = "gray85", high = "darkred") +
labs(title = "Average Percent of Foreign-Born Population Per Region", fill = "Avg Percent") +
theme_void()

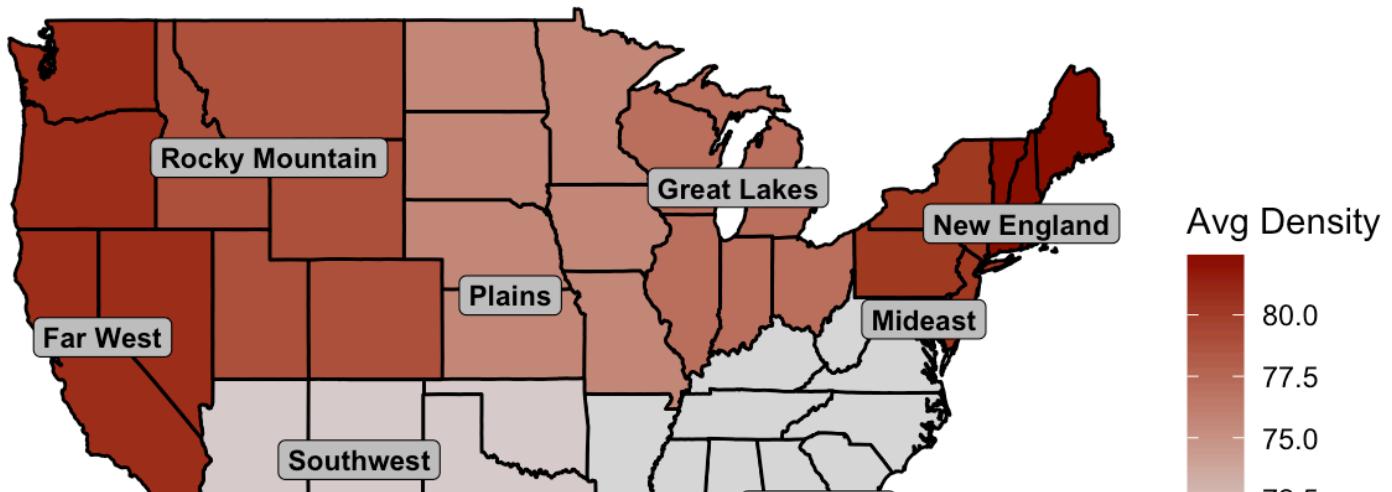
p4 <- ggplot(plot_data, aes(x = long, y = lat, group = group, fill = avg_pct_it_workers)) +
geom_polygon(color = "black") +
geom_label(data = bea_regions,
           aes(x = x, y = y, label = group),
           size = 3, fontface = "bold",
           label.padding = unit(0.2, "lines"),
           label.size = 0.2,
           fill = "gray75", color = "black") +
scale_fill_gradient(low = "gray85", high = "darkred") +
labs(title = "Average Percent of IT Workers Per Region", fill = "Avg Percent") +
theme_void()

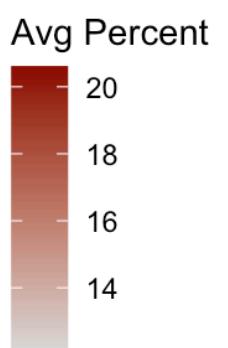
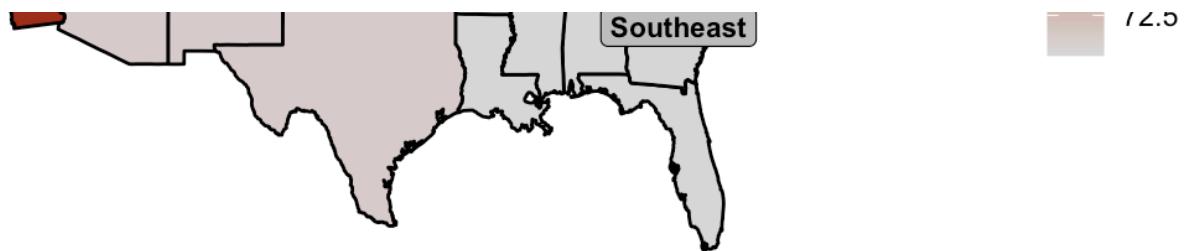
p5 <- ggplot(plot_data, aes(x = long, y = lat, group = group, fill = avg_median_hh_inc)) +
geom_polygon(color = "black") +
geom_label(data = bea_regions,
           aes(x = x, y = y, label = group),
           size = 3, fontface = "bold",
           label.padding = unit(0.2, "lines"),
           label.size = 0.2,
           fill = "gray75", color = "black") +
scale_fill_gradient(low = "gray85", high = "darkred") +
labs(title = "Average Median Household inc Per Region", fill = "Avg inc") +
theme_void()

grid.arrange(p1, p2, p3, p4, p5, nrow = 5)

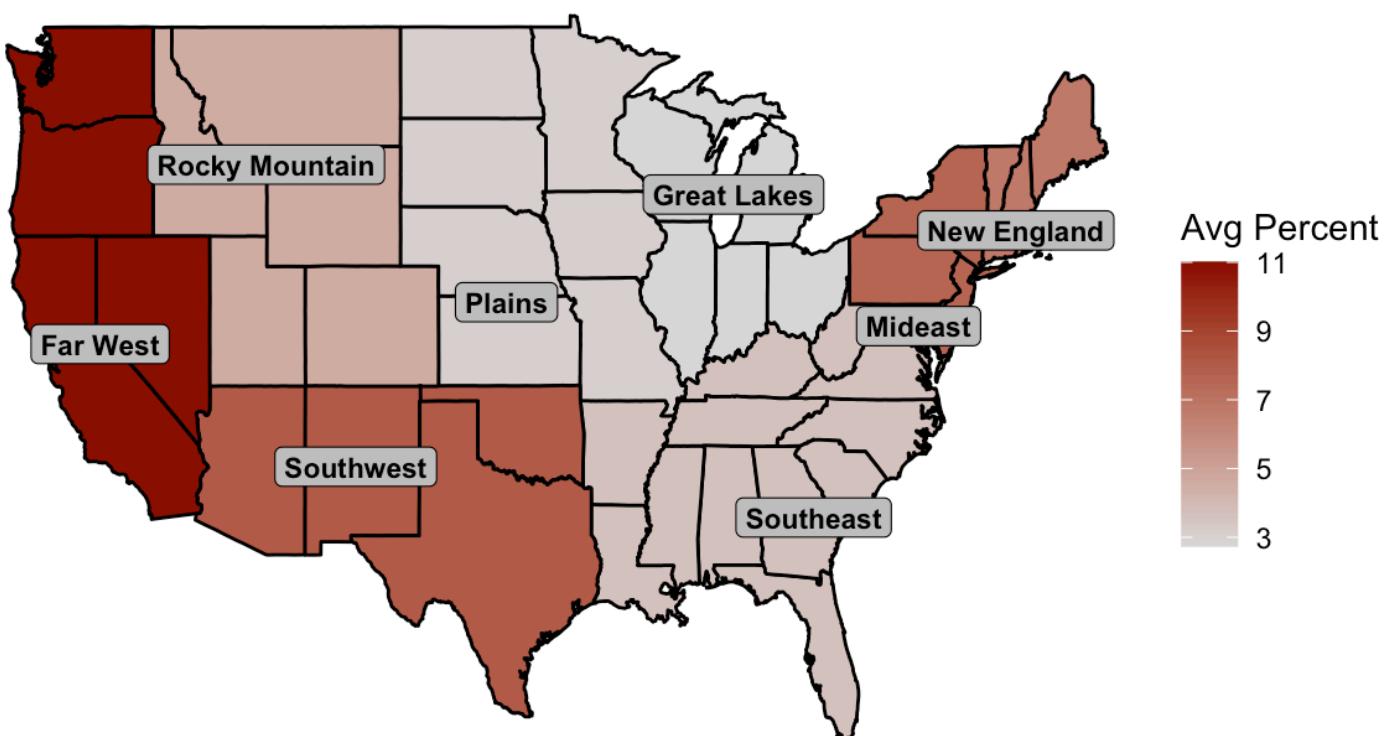
```

Average Percent of Broadband Access Per Region

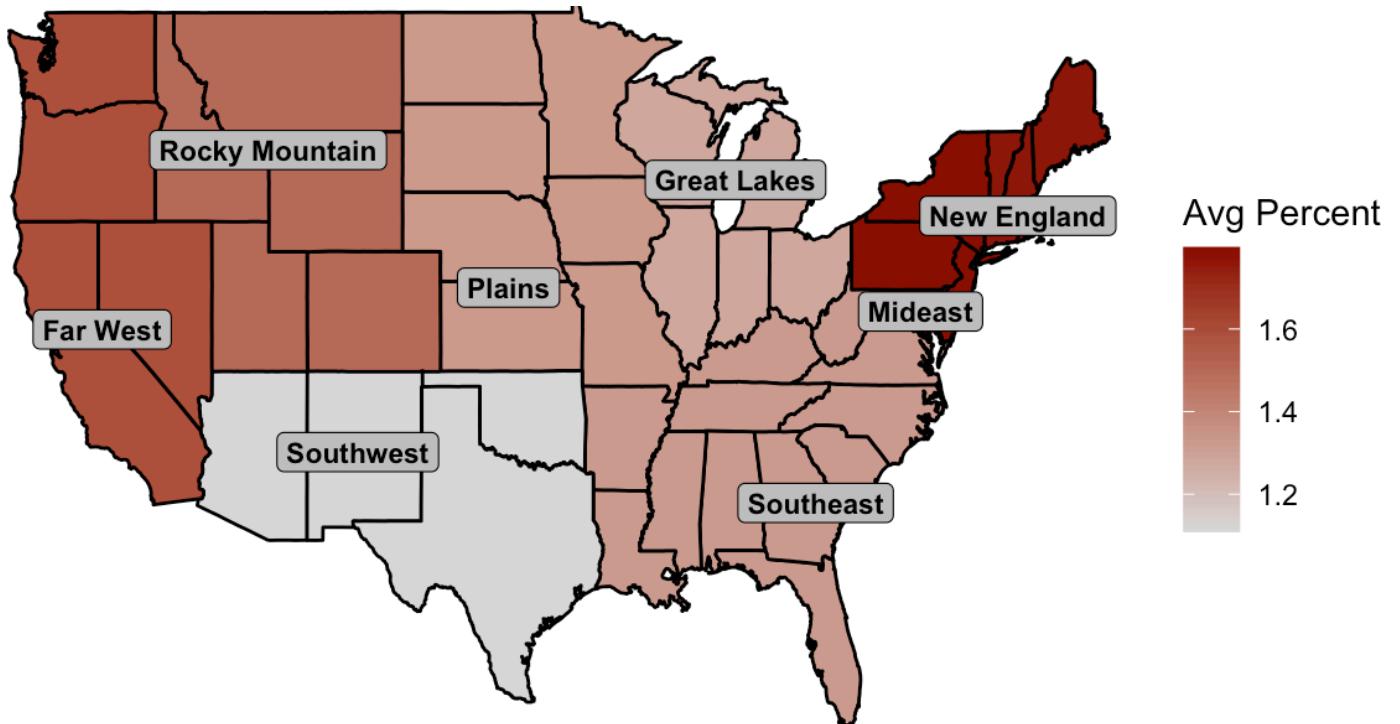




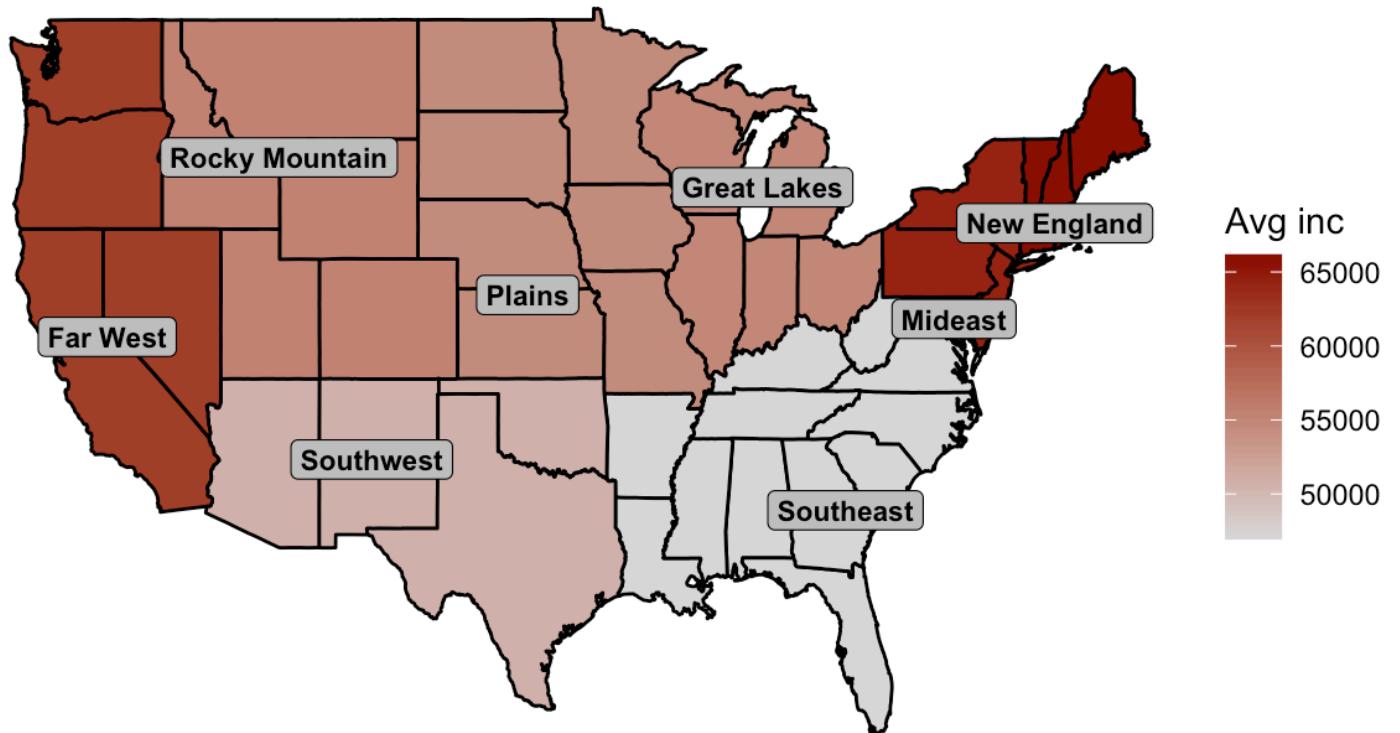
Average Percent of Foreign-Born Population Per Region



Average Percent of IT Workers Per Region



Average Median Household inc Per Region



```
state_avg <- aggregate(cbind(pct_bb, pct_college, pct_foreign_born, pct_it_workers
, median_hh_inc) ~ state, data = merged_df, FUN = mean)

#Load US state map data
us_map <- map_data("state")

#Merge state_avg with us_map based on region and state
map_data <- merge(us_map, state_avg, by.x = c("region"), by.y = c("state")) %>%
arrange(order)
```

```

#create a grid of heatmaps for each variable
grid_arrange_shared_legend <- function(...) {
  plots <- list(...)
  g <- ggplotGrob(plots[[1]] + theme(legend.position="bottom"))$grobs
  legend <- g[[which(sapply(g, function(x) x$name) == "guide-box")]]
  lheight <- sum(legend$height)
  grid.arrange(
    do.call(arrangeGrob, lapply(plots, function(x)
      x + theme(legend.position="none") + theme(panel.background = element_rect(fill =
        "gray85", color = NA)))
  ),
  bottom = legend,
  ncol = 5,
  heights = rep((unit(1, "npc") - lheight) / length(plots), length(plots))
)
}

heatmap_bb <- ggplot(map_data, aes(x = long, y = lat, group = group, fill = pct_bb)) +
  geom_polygon() +
  scale_fill_gradient(low = "white", high = "darkgreen") +
  coord_map() +
  labs(title = "Percent of Broadband Access per State", fill = "Percent") +
  theme_void()

heatmap_college <-
  ggplot(map_data, aes(x = long, y = lat, group = group, fill = pct_college)) +
  geom_polygon() +
  scale_fill_gradient(low = "white", high = "darkgreen") +
  coord_map() +
  labs(title = "Percent of Population with College Education per State", fill = "Percent") +
  theme_void()

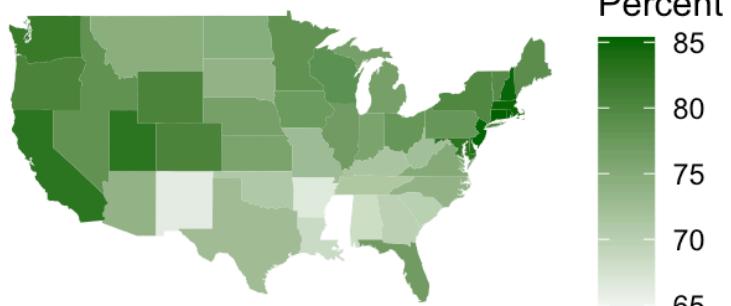
heatmap_foreign_born <-
  ggplot(map_data, aes(x = long, y = lat, group = group, fill = pct_foreign_born)) +
  geom_polygon() +
  scale_fill_gradient(low = "white", high = "darkgreen") +
  coord_map() +
  labs(title = "Percent of Foreign-born Population per State", fill = "Percent") +
  theme_void()

heatmap_it_workers <-
  ggplot(map_data, aes(x = long, y = lat, group = group, fill = pct_it_workers)) +
  geom_polygon() +
  scale_fill_gradient(low = "white", high = "darkgreen") +
  coord_map() +
  labs(title = "Percent of IT Workers per State", fill = "Percent") +
  theme_void()
heatmap_inc <-
  ggplot(map_data, aes(x = long, y = lat, group = group, fill = median_hh_inc)) +
  geom_polygon() +
  scale_fill_gradient(low = "white", high = "darkgreen") +

```

```
coord_map() +  
  labs(title = "Median Household Income per State", fill = "Dollars") +  
  theme_void()  
  
#Plot the grid of heatmaps  
#gridArrangeSharedLegend(heatmap_bb, heatmap_college, heatmap_foreign_born, hea  
tmap_it_workers, heatmap_inc)  
grid.arrange(heatmap_bb, heatmap_college, heatmap_foreign_born, heatmap_it_workers  
, heatmap_inc, nrow = 5)
```

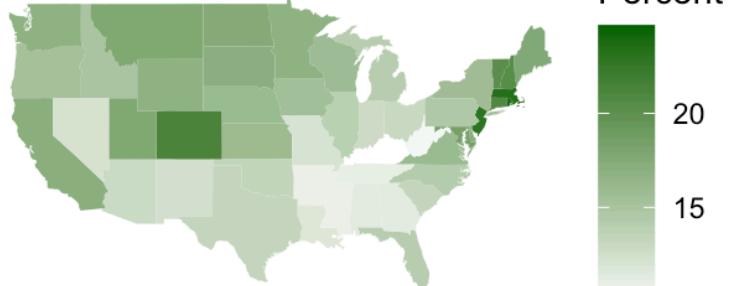
Percent of Broadband Access per State



Percent

85
80
75
70
65

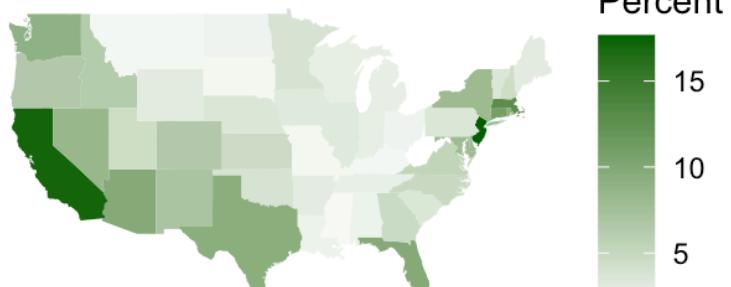
Percent of Population with College Education per State



Percent

20
15
10

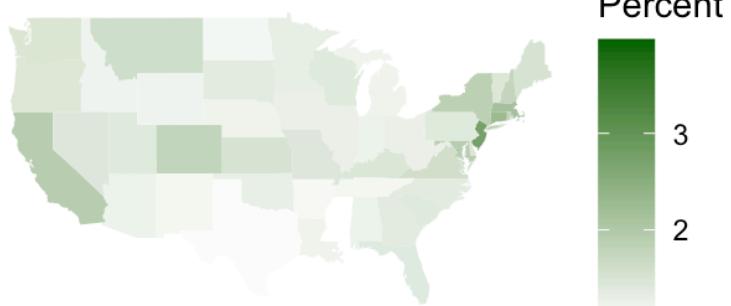
Percent of Foreign-born Population per State



Percent

15
10
5

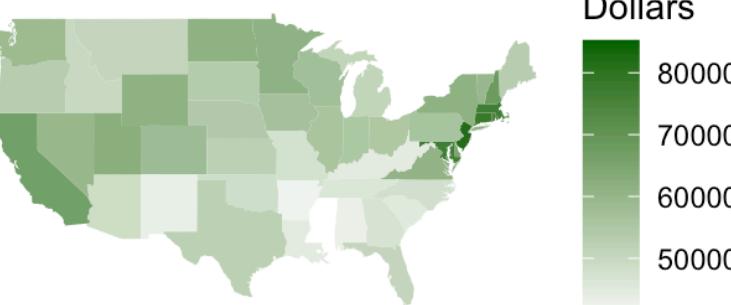
Percent of IT Workers per State



Percent

3
2

Median Household Income per State



Dollars

80000
70000
60000
50000
40000

```
county_avg <- aggregate(cbind(pct_bb, pct_college, pct_foreign_born, pct_it_worker  
s, median_hh_inc) ~ county + state, data = merged_df, FUN = mean)
```

```

#Get rid of county, city, and parish in the end of county names
county_avg$county <- gsub(" county", "", county_avg$county)
county_avg$county <- gsub(" city", "", county_avg$county)
county_avg$county <- gsub(" parish", "", county_avg$county)

#Load US county map data
us_map <- map_data("county")

#Merge county_avg with us_map based on region and county
map_data <- merge(us_map, county_avg, by.x = c("subregion", "region"), by.y = c("county", "state")) %>%
  arrange(order)

#Create a grid of heatmaps for each variable
grid_arrange_shared_legend <- function(...) {
  plots <- list(...)
  g <- ggplotGrob(plots[[1]] + theme(legend.position="bottom"))$grobs
  legend <- g[[which(sapply(g, function(x) x$name) == "guide-box")]]
  lheight <- sum(legend$height)
  grid.arrange(
    do.call(arrangeGrob, lapply(plots, function(x)
      x + theme(legend.position="none") + theme(panel.background = element_rect(fill =
        gray85, color = NA)))
    ),
    bottom = legend,
    ncol = 5,
    heights = rep((unit(1, "npc") - lheight) / length(plots), length(plots))
  )
}

heatmap_bb <- ggplot(map_data, aes(x = long, y = lat, group = group, fill = pct_bb)) +
  geom_polygon() +
  scale_fill_gradient(low = "lightblue", high = "navyblue") +
  coord_map() +
  labs(title = "Percent of Broadband Access per County", fill = "Percent") +
  theme_void()

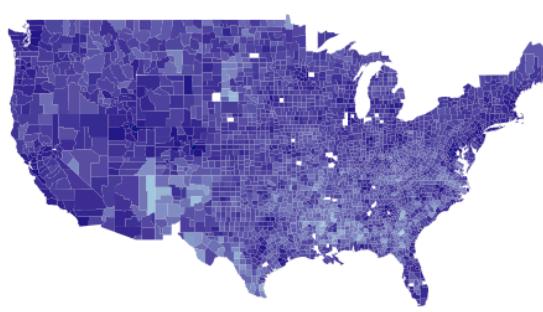
heatmap_college <-
  ggplot(map_data, aes(x = long, y = lat, group = group, fill = pct_college)) +
  geom_polygon() +
  scale_fill_gradient(low = "lightblue", high = "navyblue") +
  coord_map() +
  labs(title = "Percent of Population with College Education per County", fill = "Percent") +
  theme_void()

heatmap_foreign_born <-
  ggplot(map_data, aes(x = long, y = lat, group = group, fill = pct_foreign_born)) +
  geom_polygon() +
  scale_fill_gradient(low = "lightblue", high = "navyblue") +
  coord_map() +
  labs(title = "Percent of Foreign-born Population per County", fill = "Percent")

```

```
+  
theme_void()  
  
heatmap_it_workers <-  
  ggplot(map_data, aes(x = long, y = lat, group = group, fill = pct_it_workers)) +  
  geom_polygon() +  
  scale_fill_gradient(low = "lightblue", high = "navyblue") +  
  coord_map() +  
  labs(title = "Percent of IT Workers per County", fill = "Percent") +  
  theme_void()  
heatmap_inc <-  
  ggplot(map_data, aes(x = long, y = lat, group = group, fill = median_hh_inc)) +  
  geom_polygon() +  
  scale_fill_gradient(low = "lightblue", high = "navyblue") +  
  coord_map() +  
  labs(title = "Median Household Income per County", fill = "Dollars") +  
  theme_void()  
  
#Plot the grid of heatmaps  
#gridArrange_shared_legend(heatmap_bb, heatmap_college, heatmap_foreign_born, hea  
tmap_it_workers, heatmap_inc)  
grid.arrange(heatmap_bb, heatmap_college, heatmap_foreign_born, heatmap_it_workers  
, heatmap_inc, nrow = 5)
```

Percent of Broadband Access per County



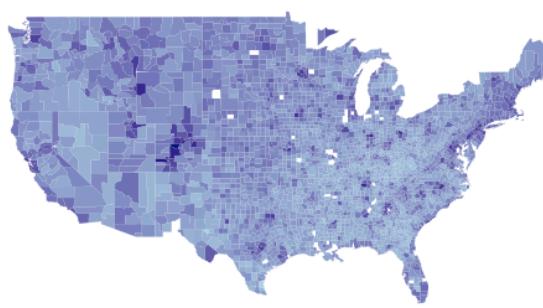
Percent

80

60

40

Percent of Population with College Education per County



Percent

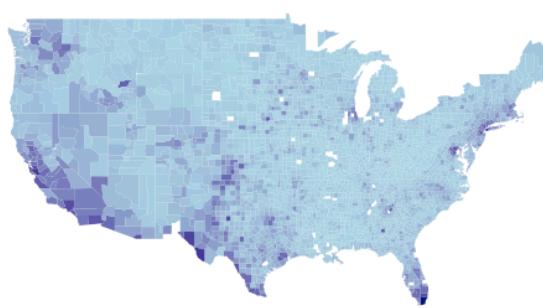
40

30

20

10

Percent of Foreign-born Population per County



Percent

50

40

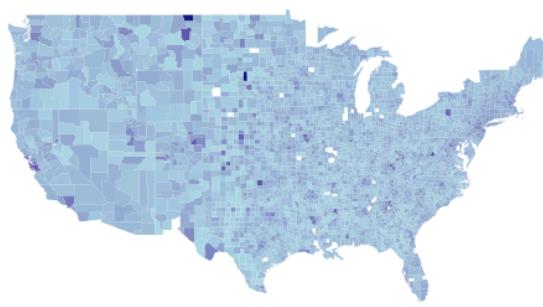
30

20

10

0

Percent of IT Workers per County



Percent

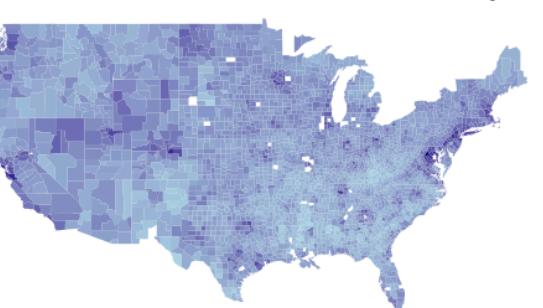
9

6

3

0

Median Household Income per County



Dollars

140000

120000

100000

80000

60000

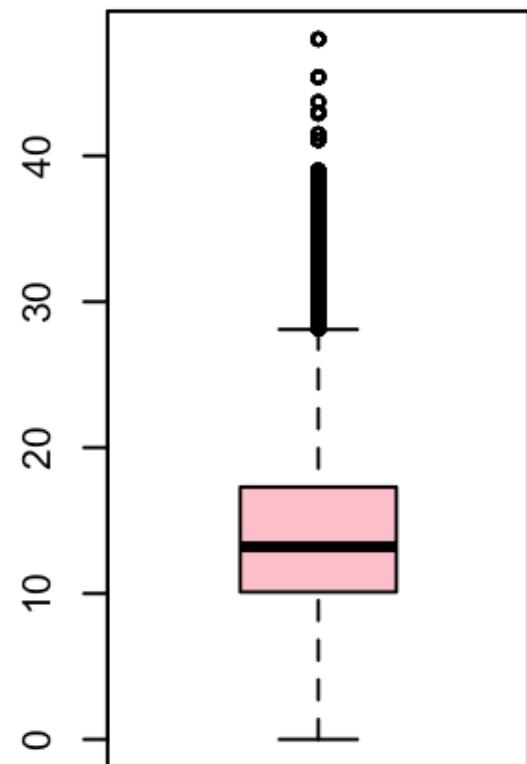
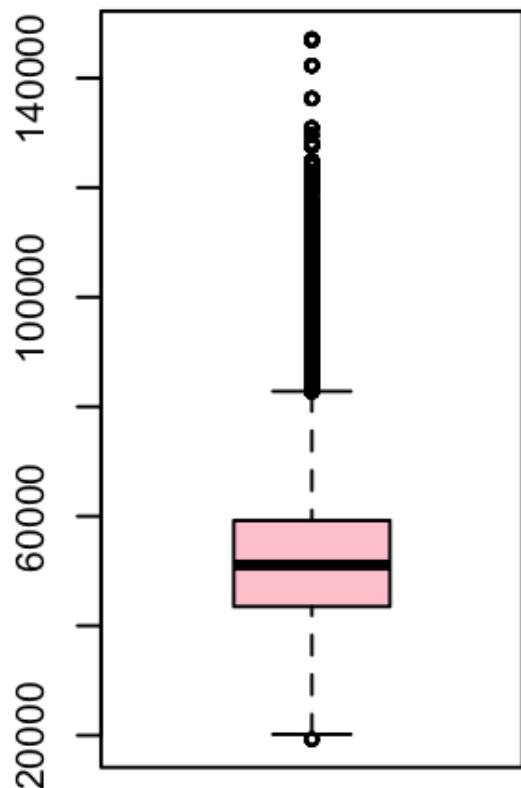
40000

```

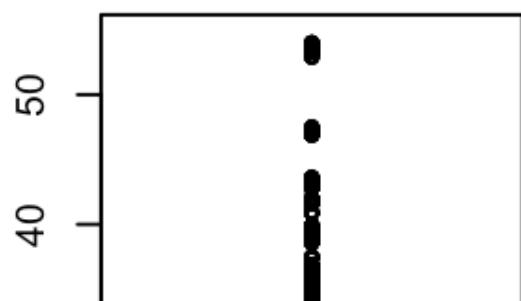
par(mfrow=c(3,2)) # set plot layout to 3 rows and 2 columns
boxplot(merged_df$median_hh_inc, col = "pink", main = "Median Household Income")
boxplot(merged_df$pct_college, col = "pink", main = "Percentage with College Education")
boxplot(merged_df$pct_foreign_born, col = "pink", main = "Percentage of Foreign-born Residents")
boxplot(merged_df$pct_it_workers, col = "pink", main = "Percentage of IT Workers")
boxplot(merged_df$pct_bb, col = "pink", main = "Percentage of Broadband Access")
#boxplot(merged_df$active, col = "pink", main = "Active Microbusiness Count")

```

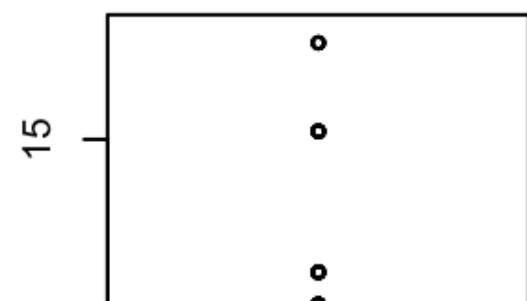
Median Household Income Percentage with College Education

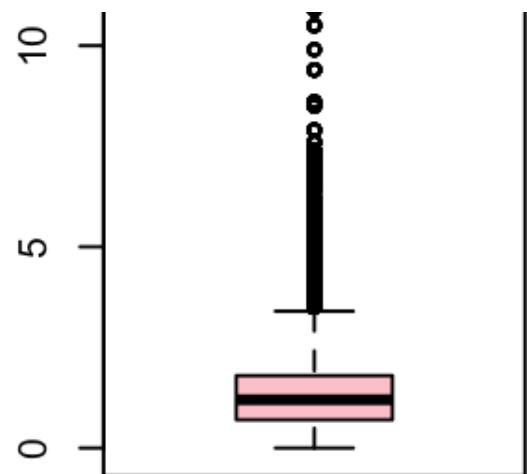
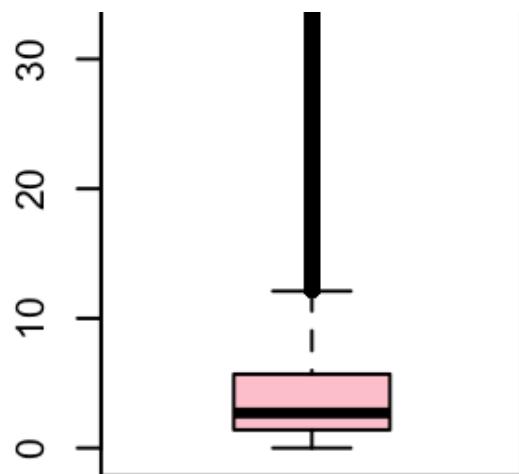


Percentage of Foreign-born Residents

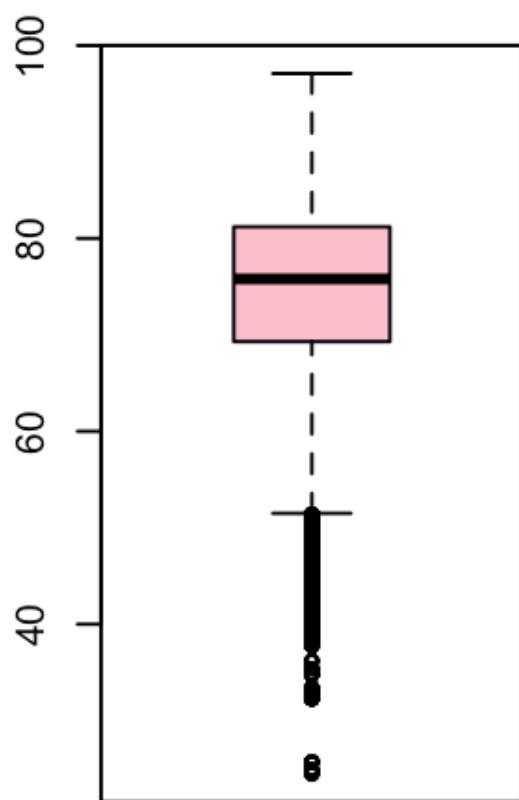


Percentage of IT Workers





Percentage of Broadband Access



3.8. Outlier Detection

Outlier detection is an important step in data analysis, as outliers can significantly affect the results of statistical analyses. There are several methods to detect outliers depending on the distribution of the data.

The Shapiro-Wilk normality test is a statistical test used to determine if a given dataset follows a normal distribution. We will perform a Shapiro-Wilk normality test on a random sample of 5000 observations from the `microbusiness_density` column of `merged_df` dataframe. To run this test we first set the seed value using `set.seed()` function to a specific random seed to ensure that the results are reproducible if the code is run again. Then, we will use the `shapiro.test()` function to perform the Shapiro-Wilk normality test on the `sample_data` object. The function returns the test statistic (W) and the p-value. A W value closer to 1 indicates that the data is more normally distributed, while a W value closer to 0 indicates greater deviation from normality. If the p-value is less than the significance level (typically 0.05), then the null hypothesis (that the sample data is normally distributed) is rejected in favor of the alternative hypothesis (that the sample data is not normally distributed).

```
# Sample 5000 observations from microbusiness_density column
set.seed(92) # Set seed for reproducibility
sample_data <- sample(merged_df$microbusiness_density, 5000)

# Perform Shapiro-Wilk test on sample_data
shapiro.test(sample_data)
```

```
## 
##  Shapiro-Wilk normality test
##
## data: sample_data
## W = 0.55179, p-value < 2.2e-16
```

The test resulted in a W statistic of 0.55179 and a p-value of less than 2.2e-16. Based on the results of the Shapiro-Wilk normality test, it can be concluded that the `sample_data` is not normally distributed.

To have a better visual on distribution of the data we will use a bell curve, a boxplot, and a Q-Q plot on `microbusiness_density`.

```

# Calculate the mean and standard deviation of 'microbusiness_density'
mean_density <- mean(merged_df$microbusiness_density)
sd_density <- sd(merged_df$microbusiness_density)

# Create a range of values for the x-axis
x_values <- seq(mean_density - 3*sd_density, mean_density + 3*sd_density, length.out = 1000)

# Create a bell curve with mean and standard deviation calculated above
y_values <- dnorm(x_values, mean = mean_density, sd = sd_density)

# Combine the 'x_values' and 'y_values' into a data frame
density_df <- data.frame(x = x_values, y = y_values)

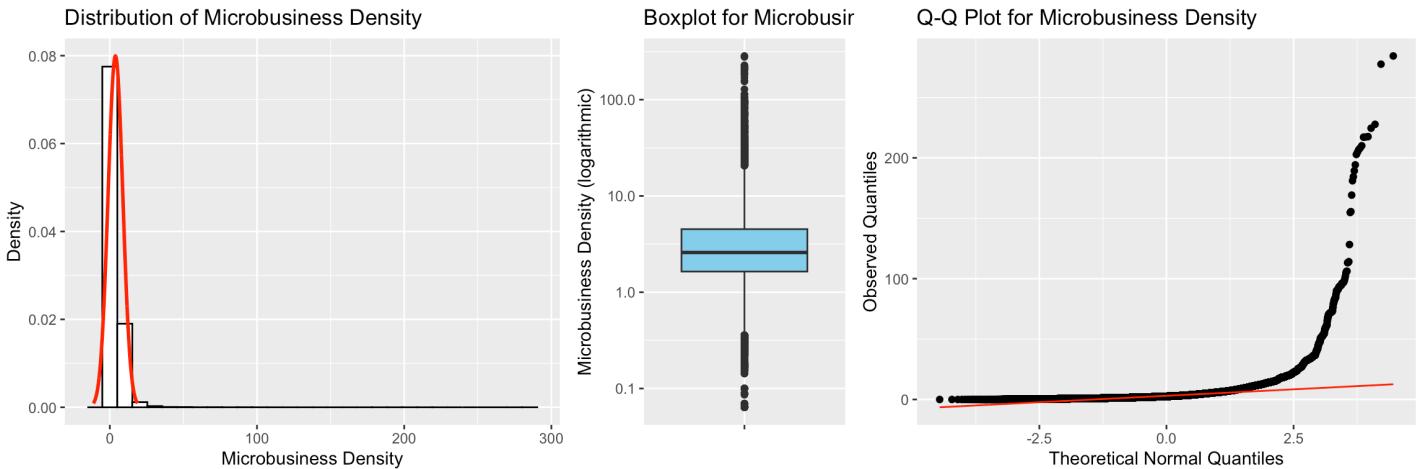
# Create a boxplot
boxplot <- ggplot(data = merged_df, aes(x = "", y = merged_df$microbusiness_density)) +
  geom_boxplot(fill = "skyblue") +
  scale_y_log10() +
  labs(x = "", y = "Microbusiness Density (logarithmic)") +
  ggtitle("Boxplot for Microbusiness Density")

# Create a Q-Q plot and add a diagonal line
qqplot <- ggplot(data = merged_df, aes(sample = microbusiness_density)) +
  stat_qq() +
  stat_qq_line(colour = "red") +
  labs(x = "Theoretical Normal Quantiles", y = "Observed Quantiles") +
  ggtitle("Q-Q Plot for Microbusiness Density")

# Create a histogram and add the bell curve
densityplot <- ggplot(data = merged_df, aes(x = microbusiness_density)) +
  geom_histogram(aes(y = after_stat(density)), bins = 30, colour = "black", fill = "white") +
  geom_line(data = density_df, aes(x = x, y = y), colour = "red", linewidth = 1) +
  labs(x = "Microbusiness Density", y = "Density") +
  ggtitle("Distribution of Microbusiness Density")

# Arrange the plots in one row using the 'grid.arrange' function from the 'gridExtra' package
grid.arrange(densityplot, boxplot, qqplot, ncol = 3, widths = c(2, 1, 2))

```



The above distribution plot explains that the dataset is right-skewed. The boxplot shows some data points away from the upper whisker; hence outliers are present in `microbusiness_density`. Q-Q plot's alignment is away from the 45-degree angle depicting outliers in the dataset.

Since the data is right-skewed and not normally distributed, a common approach to detecting outliers is to use the interquartile range (IQR) method. Now, we have to find the boundary of minimum and maximum values, out of which data would be considered an outlier.

The decision range approach involves setting a range of values outside of which any observations are considered outliers. One common approach is to use the interquartile range (IQR) to define the decision range. The IQR is calculated as the difference between the third quartile (Q3) and the first quartile (Q1) of the data.

The decision range is then defined as the range from $Q1 - 1.5 * IQR$ to $Q3 + 1.5 * IQR$. Any observations that fall outside of this range are considered outliers. This method is useful for identifying potential outliers in a dataset and can help to ensure that statistical analyses are robust and accurate.

```
quartiles <- quantile(merged_df$microbusiness_density, probs = seq(0, 1, 0.25), na.rm = FALSE,
names = TRUE, type = 7, digits = 6)
quartiles
```

```
##          0%        25%        50%        75%       100%
## 0.000000  1.639344  2.586543  4.519231 284.340030
```

```
# Calculate IQR of microbusiness_density column
q <- quantile(merged_df$microbusiness_density, c(0.25, 0.75))
iqr <- q[2] - q[1]

# Calculate lower and upper bounds for outliers
lower_bound <- q[1] - 1.5*iqr
upper_bound <- q[2] + 1.5*iqr

# Count number of outliers
num_outliers <- sum(merged_df$microbusiness_density < lower_bound | merged_df$microbusiness_density > upper_bound)

# Calculate percent of outliers
percent_outliers <- num_outliers / length(merged_df$microbusiness_density) * 100

# Print results
cat("Number of outliers:", num_outliers, "\n")
```

```
## Number of outliers: 8746
```

```
cat("Percent of outliers:", percent_outliers, "%\n")
```

```
## Percent of outliers: 7.153315 %
```

```

# Create new dataframe without outliers
merged_df_new <- merged_df[merged_df$microbusiness_density >= lower_bound & merged
                           _df$microbusiness_density <= upper_bound,]

# Print number of rows removed
cat("Number of rows removed:", nrow(merged_df) - nrow(merged_df_new), "\n")

```

```
## Number of rows removed: 8746
```

```

# Calculate the mean and standard deviation of 'microbusiness_density'
mean_density <- mean(merged_df_new$microbusiness_density)
sd_density <- sd(merged_df_new$microbusiness_density)

# Create a range of values for the x-axis
x_values <- seq(mean_density - 3*sd_density, mean_density + 3*sd_density, length.o
ut = 1000)

# Create a bell curve with mean and standard deviation calculated above
y_values <- dnorm(x_values, mean = mean_density, sd = sd_density)

# Combine the 'x_values' and 'y_values' into a data frame
density_df <- data.frame(x = x_values, y = y_values)

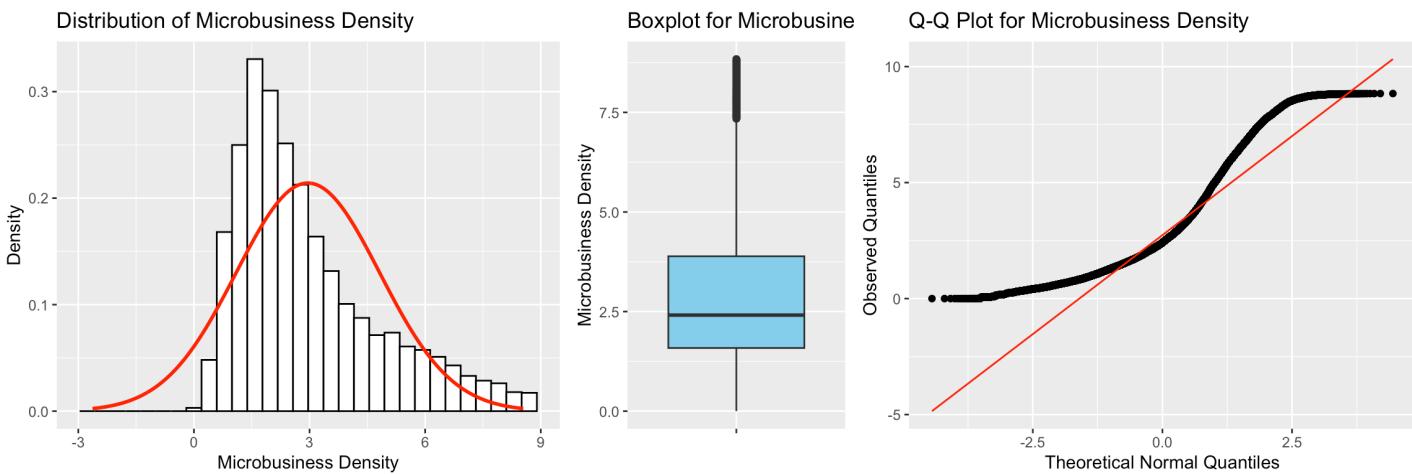
# Create a boxplot
boxplot <- ggplot(data = merged_df_new, aes(x = "", y = merged_df_new$microbusines
s_density)) +
  geom_boxplot(fill = "skyblue") +
  # scale_y_log10() +
  labs(x = "", y = "Microbusiness Density") +
  ggtitle("Boxplot for Microbusiness Density")

# Create a Q-Q plot and add a diagonal line
qqplot <- ggplot(data = merged_df_new, aes(sample = microbusiness_density)) +
  stat_qq() +
  stat_qq_line(colour = "red") +
  labs(x = "Theoretical Normal Quantiles", y = "Observed Quantiles") +
  ggtitle("Q-Q Plot for Microbusiness Density")

# Create a bell curve and add the bell curve
densityplot <- ggplot(data = merged_df_new, aes(x = microbusiness_density)) +
  geom_histogram(aes(y = after_stat(density)), bins = 30, colour = "black", fill =
"white") +
  geom_line(data = density_df, aes(x = x, y = y), colour = "red", linewidth = 1) +
  labs(x = "Microbusiness Density", y = "Density") +
  ggtitle("Distribution of Microbusiness Density")

# Arrange the plots in one row
grid.arrange(densityplot, boxplot, qqplot, ncol = 3, widths = c(2, 1, 2))

```



```
# Calculate IQR of microbusiness_density column
q <- quantile(merged_df_new$microbusiness_density, c(0.25, 0.75))
iqr <- q[2] - q[1]

# Calculate lower and upper bounds for outliers
lower_bound <- q[1] - 1.5*iqr
upper_bound <- q[2] + 1.5*iqr

# Count number of outliers
num_outliers <- sum(merged_df_new$microbusiness_density < lower_bound | merged_df_new$microbusiness_density > upper_bound)

# Calculate percent of outliers
percent_outliers <- num_outliers / length(merged_df_new$microbusiness_density) * 1
00

# Print results
cat("Number of outliers:", num_outliers, "\n")
```

```
## Number of outliers: 3946
```

```
cat("Percent of outliers:", percent_outliers, "%\n")
```

```
## Percent of outliers: 3.47607 %
```

```
# Create new dataframe without outliers
merged_df_clean <- merged_df_new[merged_df_new$microbusiness_density >= lower_bound & merged_df_new$microbusiness_density <= upper_bound,]

# Print number of rows removed
cat("Number of rows removed:", nrow(merged_df_new) - nrow(merged_df_clean), "\n")
```

```
## Number of rows removed: 3946
```

```
cat("Total rows removed:", nrow(merged_df) - nrow(merged_df_clean))
```

```
## Total rows removed: 12692
```

```
# Calculate the mean and standard deviation of 'microbusiness_density'
mean_density <- mean(merged_df_clean$microbusiness_density)
sd_density <- sd(merged_df_clean$microbusiness_density)

# Create a range of values for the x-axis
x_values <- seq(mean_density - 3*sd_density, mean_density + 3*sd_density, length.out = 1000)

# Create a bell curve with mean and standard deviation calculated above
y_values <- dnorm(x_values, mean = mean_density, sd = sd_density)

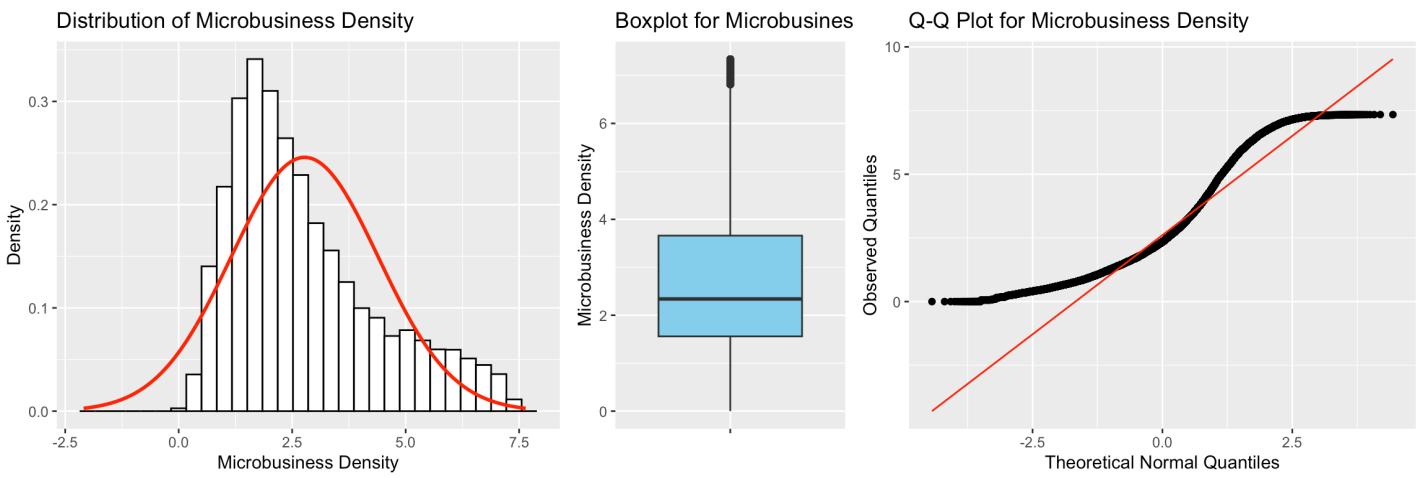
# Combine the 'x_values' and 'y_values' into a data frame
density_df <- data.frame(x = x_values, y = y_values)

# Create a boxplot
boxplot <- ggplot(data = merged_df_clean, aes(x = "", y = merged_df_clean$microbusiness_density)) +
  geom_boxplot(fill = "skyblue") +
# scale_y_log10() +
  labs(x = "", y = "Microbusiness Density") +
  ggtitle("Boxplot for Microbusiness Density")

# Create a Q-Q plot and add a diagonal line
qqplot <- ggplot(data = merged_df_clean, aes(sample = microbusiness_density)) +
  stat_qq() +
  stat_qq_line(colour = "red") +
  labs(x = "Theoretical Normal Quantiles", y = "Observed Quantiles") +
  ggtitle("Q-Q Plot for Microbusiness Density")

# Create a histogram and add the bell curve
densityplot <- ggplot(data = merged_df_clean, aes(x = microbusiness_density)) +
  geom_histogram(aes(y = after_stat(density)), bins = 30, colour = "black", fill = "white") +
  geom_line(data = density_df, aes(x = x, y = y), colour = "red", linewidth = 1) +
  labs(x = "Microbusiness Density", y = "Density") +
  ggtitle("Distribution of Microbusiness Density")

# Arrange the plots in one row using the 'grid.arrange' function from the 'gridExtra' package
grid.arrange(densityplot, boxplot, qqplot, ncol = 3, widths = c(2, 1, 2))
```



```
# Calculate the mean and standard deviation of 'pct_bb'
mean_pct_bb <- mean(merged_df$pct_bb)
sd_pct_bb <- sd(merged_df$pct_bb)

# Create a range of values for the x-axis
x_values <- seq(mean_pct_bb - 3*sd_pct_bb, mean_pct_bb + 3*sd_pct_bb, length.out = 1000)

# Create a bell curve with mean and standard deviation calculated above
y_values <- dnorm(x_values, mean = mean_pct_bb, sd = sd_pct_bb)

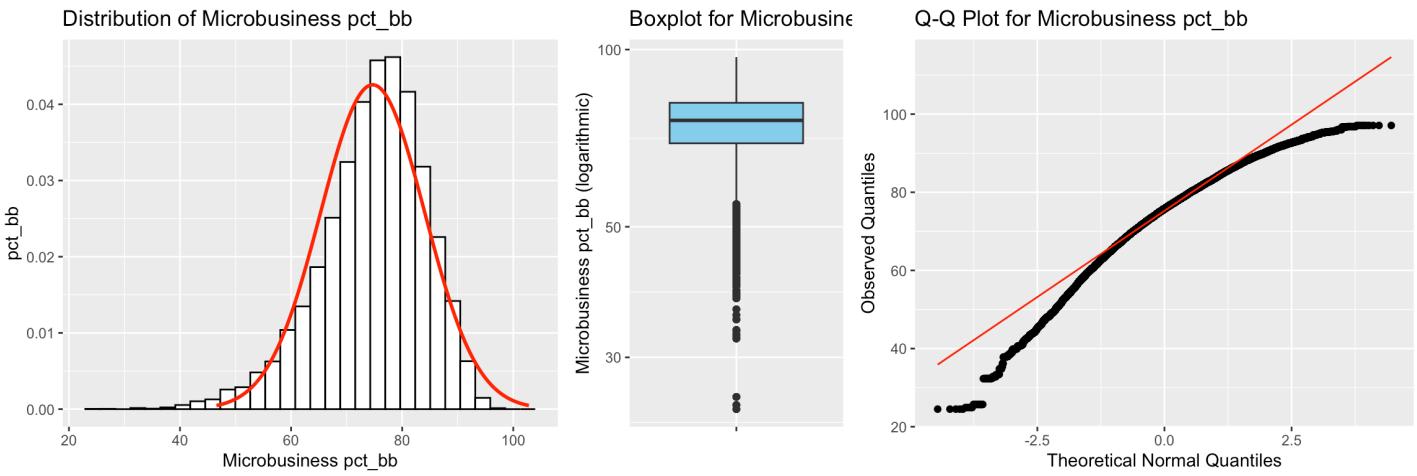
# Combine the 'x_values' and 'y_values' into a data frame
pct_bb_df <- data.frame(x = x_values, y = y_values)

# Create a boxplot
boxplot <- ggplot(data = merged_df, aes(x = "", y = merged_df$pct_bb)) +
  geom_boxplot(fill = "skyblue") +
  scale_y_log10() +
  labs(x = "", y = "Microbusiness pct_bb (logarithmic)") +
  ggtitle("Boxplot for Microbusiness pct_bb")

# Create a Q-Q plot and add a diagonal line
qqplot <- ggplot(data = merged_df, aes(sample = pct_bb)) +
  stat_qq() +
  stat_qq_line(colour = "red") +
  labs(x = "Theoretical Normal Quantiles", y = "Observed Quantiles") +
  ggtitle("Q-Q Plot for Microbusiness pct_bb")

# Create a histogram and add the bell curve
densityplot <- ggplot(data = merged_df, aes(x = pct_bb)) +
  geom_histogram(aes(y = after_stat(density)), bins = 30, colour = "black", fill = "white") +
  geom_line(data = pct_bb_df, aes(x = x, y = y), colour = "red", linewidth = 1) +
  labs(x = "Microbusiness pct_bb", y = "pct_bb") +
  ggtitle("Distribution of Microbusiness pct_bb")

# Arrange the plots in one row using the 'grid.arrange' function from the 'gridExt
# ra' package
grid.arrange(densityplot, boxplot, qqplot, ncol = 3, widths = c(2, 1, 2))
```



```

# Specify the columns to analyze
cols_to_analyze <- c("pct_bb", "pct_college", "pct_it_workers", "pct_foreign_born",
, "median_hh_inc")

# Loop through each column and detect outliers using the IQR method
for (col in cols_to_analyze) {
  # Calculate the interquartile range (IQR) of the column
  q1 <- quantile(merged_df_new[[col]], 0.25)
  q3 <- quantile(merged_df_new[[col]], 0.75)
  iqr <- q3 - q1

  # Calculate the upper and lower bounds for outliers
  upper_bound <- q3 + (1.5 * iqr)
  lower_bound <- q1 - (1.5 * iqr)

  # Identify the outliers in the column
  outliers <- merged_df_new[[col]][merged_df_new[[col]] > upper_bound | merged_df_
new[[col]] < lower_bound]

  # Print the results
  # cat(paste("Outliers in", col, ":", toString(outliers), "\n"))
  # Count number of outliers
  num_outliers <- sum(merged_df_new[[col]] < lower_bound | merged_df_new[[col]] >
upper_bound)
  # Calculate percent of outliers
  percent_outliers <- num_outliers / length(merged_df_new[[col]]) * 100

  # Print results
  cat("Number of outliers in", col, ":", num_outliers, "\n")
  cat("Percent of outliers in", col, ":", percent_outliers, "%\n")
}

```

```
## Number of outliers in pct_bb : 2507
## Percent of outliers in pct_bb : 2.208441 %
## Number of outliers in pct_college : 1428
## Percent of outliers in pct_college : 1.257939 %
## Number of outliers in pct_it_workers : 2556
## Percent of outliers in pct_it_workers : 2.251605 %
## Number of outliers in pct_foreign_born : 8329
## Percent of outliers in pct_foreign_born : 7.337098 %
## Number of outliers in median_hh_inc : 3410
## Percent of outliers in median_hh_inc : 3.003902 %
```

```
# Create a copy of the original dataframe
merged_df_clean <- merged_df_new

# Loop through each column and detect outliers using the IQR method
for (col in cols_to_analyze) {
  # Calculate the interquartile range (IQR) of the column
  q1 <- quantile(merged_df_clean[[col]], 0.25)
  q3 <- quantile(merged_df_clean[[col]], 0.75)
  iqr <- q3 - q1

  # Calculate the upper and lower bounds for outliers
  upper_bound <- q3 + (1.5 * iqr)
  lower_bound <- q1 - (1.5 * iqr)

  # Identify the outliers in the column
  outliers <- merged_df_clean[[col]][merged_df_clean[[col]] > upper_bound | merged_df_clean[[col]] < lower_bound]

  # Remove the outliers from the dataframe
  merged_df_clean <- merged_df_clean[!(merged_df_clean[[col]] %in% outliers), ]
}

# Print the dimensions of the original and cleaned dataframes
cat("Original dataframe dimensions:", nrow(merged_df_new), "\n")
```

```
## Original dataframe dimensions: 113519
```

```
cat("Cleaned dataframe dimensions:", nrow(merged_df_clean), "\n")
```

```
## Cleaned dataframe dimensions: 97098
```

```
cat("Total rows removed:", nrow(merged_df) - nrow(merged_df_clean), "\n")
```

```
## Total rows removed: 25167
```

```
cat("Total percent removed:", 100 * (nrow(merged_df) - nrow(merged_df_clean)) / nrow(merged_df), "%", "\n")
```

```
## Total percent removed: 20.58398 %
```

3.9. Non-Parametric Test

```
# Specify a list of column pairs to compare
col_pairs <- list(c("microbusiness_density", "pct_bb"), c("microbusiness_density",
"pct_college"), c("microbusiness_density", "pct_foreign_born"), c("microbusiness_d
ensity", "pct_it_workers"), c("microbusiness_density", "median_hh_inc"))

# Loop through each column pair and perform the Wilcoxon signed-rank test
for (pair in col_pairs) {
  test_result <- wilcox.test(merged_df_new[[pair[1]]], merged_df_new[[pair[2]]], p
aired = TRUE)
  cat(paste("Wilcoxon signed-rank test results for", pair[1], "and", pair[2], ":\n"
"))
  print(test_result)
  cat("\n")
}
```

```
## Wilcoxon signed-rank test results for microbusiness_density and pct_bb :  
##  
## Wilcoxon signed rank test with continuity correction  
##  
## data: merged_df_new[[pair[1]]] and merged_df_new[[pair[2]]]  
## V = 0, p-value < 2.2e-16  
## alternative hypothesis: true location shift is not equal to 0  
##  
##  
## Wilcoxon signed-rank test results for microbusiness_density and pct_college :  
##  
## Wilcoxon signed rank test with continuity correction  
##  
## data: merged_df_new[[pair[1]]] and merged_df_new[[pair[2]]]  
## V = 6780, p-value < 2.2e-16  
## alternative hypothesis: true location shift is not equal to 0  
##  
##  
## Wilcoxon signed-rank test results for microbusiness_density and pct_foreign_bor  
n :  
##  
## Wilcoxon signed rank test with continuity correction  
##  
## data: merged_df_new[[pair[1]]] and merged_df_new[[pair[2]]]  
## V = 2843324130, p-value < 2.2e-16  
## alternative hypothesis: true location shift is not equal to 0  
##  
##  
## Wilcoxon signed-rank test results for microbusiness_density and pct_it_workers :  
##  
## Wilcoxon signed rank test with continuity correction  
##  
## data: merged_df_new[[pair[1]]] and merged_df_new[[pair[2]]]  
## V = 5996103713, p-value < 2.2e-16  
## alternative hypothesis: true location shift is not equal to 0  
##  
##  
## Wilcoxon signed-rank test results for microbusiness_density and median_hh_inc :  
##  
## Wilcoxon signed rank test with continuity correction  
##  
## data: merged_df_new[[pair[1]]] and merged_df_new[[pair[2]]]  
## V = 0, p-value < 2.2e-16  
## alternative hypothesis: true location shift is not equal to 0
```

```
# Specify a list of column pairs to compare
col_pairs <- list(c("microbusiness_density", "pct_bb"), c("microbusiness_density",
"pct_college"), c("microbusiness_density", "pct_foreign_born"), c("microbusiness_d
ensity", "pct_it_workers"), c("microbusiness_density", "median_hh_inc"))

# Loop through each column pair and perform the Wilcoxon rank sum test
for (pair in col_pairs) {
  test_result <- wilcox.test(merged_df_new[[pair[1]]], merged_df_new[[pair[2]]], p
aired = FALSE)
  cat(paste("Wilcoxon rank-sum test results for", pair[1], "and", pair[2], ":\n"))
  print(test_result)
  cat("\n")
}
```

```

## Wilcoxon rank-sum test results for microbusiness_density and pct_bb :
##
## Wilcoxon rank sum test with continuity correction
##
## data: merged_df_new[[pair[1]]] and merged_df_new[[pair[2]]]
## W = 0, p-value < 2.2e-16
## alternative hypothesis: true location shift is not equal to 0
##
##
## Wilcoxon rank-sum test results for microbusiness_density and pct_college :
##
## Wilcoxon rank sum test with continuity correction
##
## data: merged_df_new[[pair[1]]] and merged_df_new[[pair[2]]]
## W = 112644499, p-value < 2.2e-16
## alternative hypothesis: true location shift is not equal to 0
##
##
## Wilcoxon rank-sum test results for microbusiness_density and pct_foreign_born :
##
## Wilcoxon rank sum test with continuity correction
##
## data: merged_df_new[[pair[1]]] and merged_df_new[[pair[2]]]
## W = 6244818672, p-value < 2.2e-16
## alternative hypothesis: true location shift is not equal to 0
##
##
## Wilcoxon rank-sum test results for microbusiness_density and pct_it_workers :
##
## Wilcoxon rank sum test with continuity correction
##
## data: merged_df_new[[pair[1]]] and merged_df_new[[pair[2]]]
## W = 1.0468e+10, p-value < 2.2e-16
## alternative hypothesis: true location shift is not equal to 0
##
##
## Wilcoxon rank-sum test results for microbusiness_density and median_hh_inc :
##
## Wilcoxon rank sum test with continuity correction
##
## data: merged_df_new[[pair[1]]] and merged_df_new[[pair[2]]]
## W = 0, p-value < 2.2e-16
## alternative hypothesis: true location shift is not equal to 0

```

The results show that for each pair of columns, the p-value is less than the significance level of 0.05, which means that we reject the null hypothesis that the median difference between the two columns is zero. Instead, we conclude that there is a statistically significant difference between the two columns.

The results suggest that there is evidence to support the hypothesis that the population median of each column is different from the population median of microbusiness_density.

kernel regression

```
# create a vector of region names and their corresponding codes
region_codes <- c("new england" = 1,
                 "mideast" = 2,
                 "great lakes" = 3,
                 "plains" = 4,
                 "southeast" = 5,
                 "southwest" = 6,
                 "rocky mountain" = 7,
                 "far west" = 8)

# use the `match()` function to find the region code for each row in `merged_df_new
w$region`
merged_df_new$region_code <- match(merged_df_new$region, names(region_codes))
merged_df_clean$region_code <- match(merged_df_clean$region, names(region_codes))

# create a vector of state names and their corresponding codes
state_codes <- sort(unique(merged_df_new$state))
state_codes <- setNames(1:length(state_codes), state_codes)

# use the `match()` function to find the state code for each row in `merged_df_new
\$state`
merged_df_new$state_code <- match(merged_df_new$state, names(state_codes))

state_codes <- sort(unique(merged_df_clean$state))
state_codes <- setNames(1:length(state_codes), state_codes)
merged_df_clean$state_code <- match(merged_df_clean$state, names(state_codes))

# print head of the updated dataframe
head(merged_df_new)
```

```

##           row_id cfips      county state first_day_of_month
## 1 1001_2019-08-01 1001 autauga county alabama 2019-08-01
## 2 1001_2019-09-01 1001 autauga county alabama 2019-09-01
## 3 1001_2019-10-01 1001 autauga county alabama 2019-10-01
## 4 1001_2019-11-01 1001 autauga county alabama 2019-11-01
## 5 1001_2019-12-01 1001 autauga county alabama 2019-12-01
## 6 1001_2020-01-01 1001 autauga county alabama 2020-01-01
##   microbusiness_density active year_month year month pct_bb pct_college
## 1             3.007682    1249 2019-08 2019     8 76.6 14.5
## 2             2.884870    1198 2019-09 2019     9 76.6 14.5
## 3             3.055843    1269 2019-10 2019    10 76.6 14.5
## 4             2.993233    1243 2019-11 2019    11 76.6 14.5
## 5             2.993233    1243 2019-12 2019    12 76.6 14.5
## 6             2.969090    1242 2020-01 2020     1 78.9 15.9
##   pct_foreign_born pct_it_workers median_hh_inc region region_code
## 1            2.1          1.3       55317 southeast 5
## 2            2.1          1.3       55317 southeast 5
## 3            2.1          1.3       55317 southeast 5
## 4            2.1          1.3       55317 southeast 5
## 5            2.1          1.3       55317 southeast 5
## 6            2.0          1.1       58786 southeast 5
##   state_code
## 1            1
## 2            1
## 3            1
## 4            1
## 5            1
## 6            1

```

```

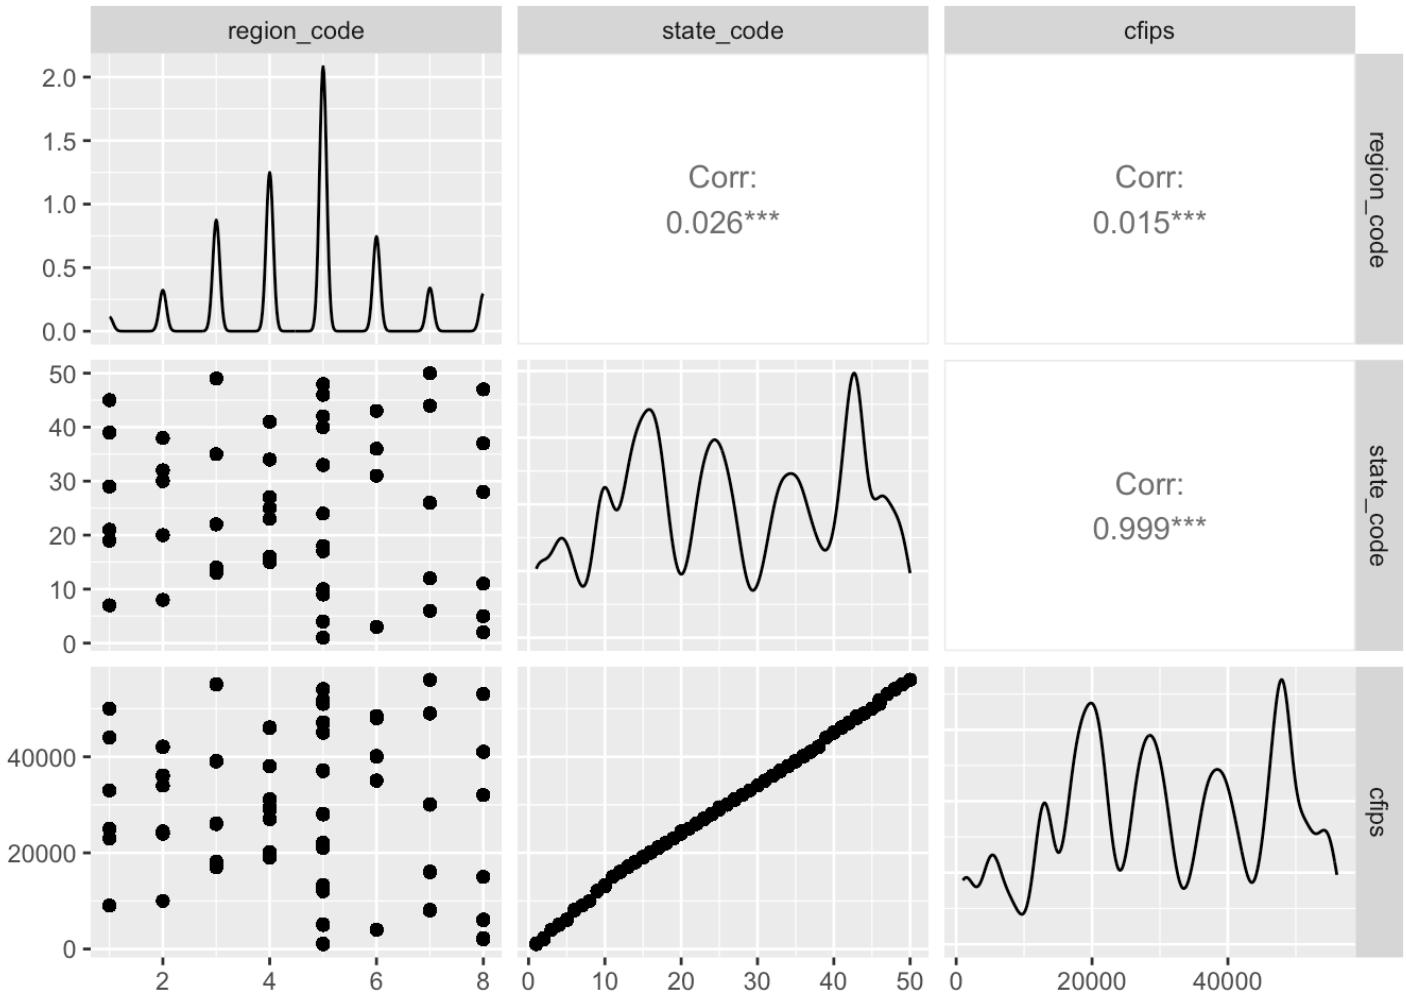
merged_df_new |>
  GGally::ggpairs(columns = c(17,18,2))

```

```

## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg   ggplot2

```



```
# Select the columns with numeric data
numeric_cols <- c("cfips", "state_code", "region_code", "microbusiness_density",
active", "year", "month", "pct_bb", "pct_college", "pct_foreign_born", "pct_it_workers",
"median_hh_inc")

# Subset the dataframe with the selected columns
merged_df_numeric <- merged_df_new[, numeric_cols]

# Calculate the correlation matrix
cor_matrix <- cor(merged_df_numeric, use="pairwise.complete.obs")

# Plot the correlation matrix using ggcorrplot
ggcorrplot(cor_matrix,
            hc.order = TRUE,
            type = "lower",
            method = "square",
            lab = TRUE,
            lab_size = 3,
            title = "Correlation Plot of merged_df_new",
            colors = c("#6D9EC1", "#FAC200", "#FA5252"),
            ggtheme = ggplot2::theme_gray,
            show.legend = TRUE)
```

Correlation Plot of merged_df_new



```
# Select the columns with numeric data
numeric_cols <- c("cfips", "state_code", "region_code", "microbusiness_density", "active", "year", "month", "pct_bb", "pct_college", "pct_foreign_born", "pct_it_workers", "median_hh_inc")

# Subset the dataframe with the selected columns
merged_df_numeric <- merged_df_clean[, numeric_cols]

# Calculate the correlation matrix
cor_matrix <- cor(merged_df_numeric, use="pairwise.complete.obs")

# Plot the correlation matrix using ggcorrplot
ggcorrplot(cor_matrix,
            hc.order = TRUE,
            type = "lower",
            method = "square",
            lab = TRUE,
            lab_size = 3,
            title = "Correlation Plot of merged_df_clean",
            colors = c("#6D9EC1", "#FAC200", "#FA5252"),
            ggtheme = ggplot2::theme_gray,
            show.legend = TRUE)
```

Correlation Plot of merged_df_clean



The `cfips` column and `state_code` have a completely positive correlation. So, we will remove the redundant feature `state_code`. Also, `pct_bb` strongly correlates with `median_hh_inc` and `pct_college`. Also, `pct_college` strongly correlates with `microbusiness_density`, `median_hh_inc` and `pct_bb`. As features with moderate strong and strong correlation are redundant, we will remove these features (`pct_bb` and `pct_college`) before further analysis. Now, we draw the new correlation plot.

```

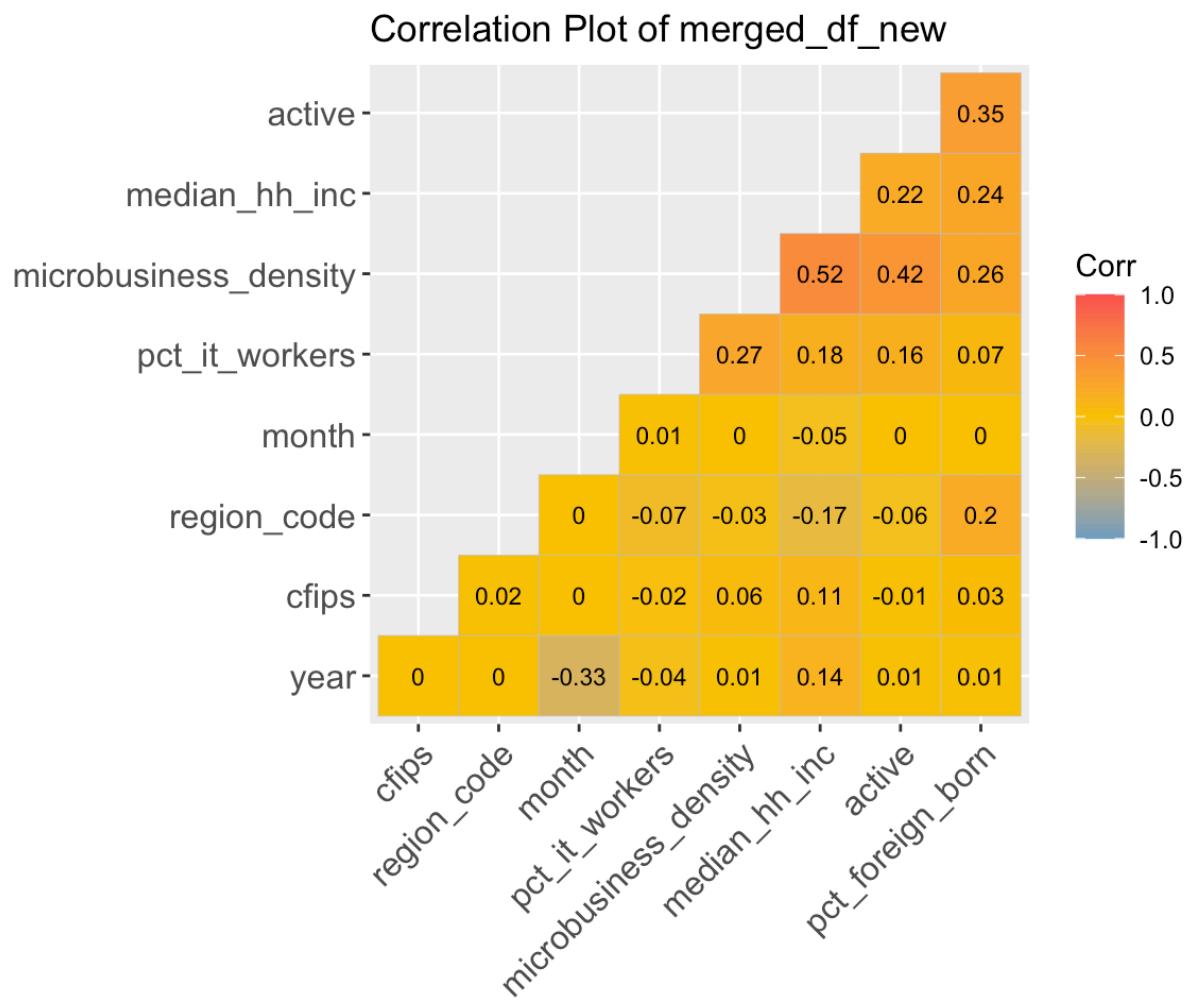
# Select the columns with numeric data (not including "state_code", "pct_bb", "pct_college")
numeric_cols <- c("cfips", "region_code", "microbusiness_density", "active", "year", "month", "pct_foreign_born", "pct_it_workers", "median_hh_inc")

# Subset the dataframe with the selected columns
merged_df_numeric <- merged_df_new[, numeric_cols]

# Calculate the correlation matrix
cor_matrix <- cor(merged_df_numeric, use="pairwise.complete.obs")

# Plot the correlation matrix using ggcorrplot
ggcorrplot(cor_matrix,
            hc.order = TRUE,
            type = "lower",
            method = "square",
            lab = TRUE,
            lab_size = 3,
            title = "Correlation Plot of merged_df_new",
            colors = c("#6D9EC1", "#FAC200", "#FA5252"),
            ggtheme = ggplot2::theme_gray,
            show.legend = TRUE)

```



3.10. Feature Importance Using Random Forest

```

# Split the dataset into training and testing sets
set.seed(92) # for reproducibility
train_index <- sample(nrow(merged_df_new), 0.7 * nrow(merged_df_new))
train_data <- merged_df_new[train_index, ]
test_data <- merged_df_new[-train_index, ]

# Build the random forest model
model <- randomForest(microbusiness_density ~ cfips + region_code + active + year
+ month + pct_foreign_born + pct_it_workers + median_hh_inc,
                        data = train_data)

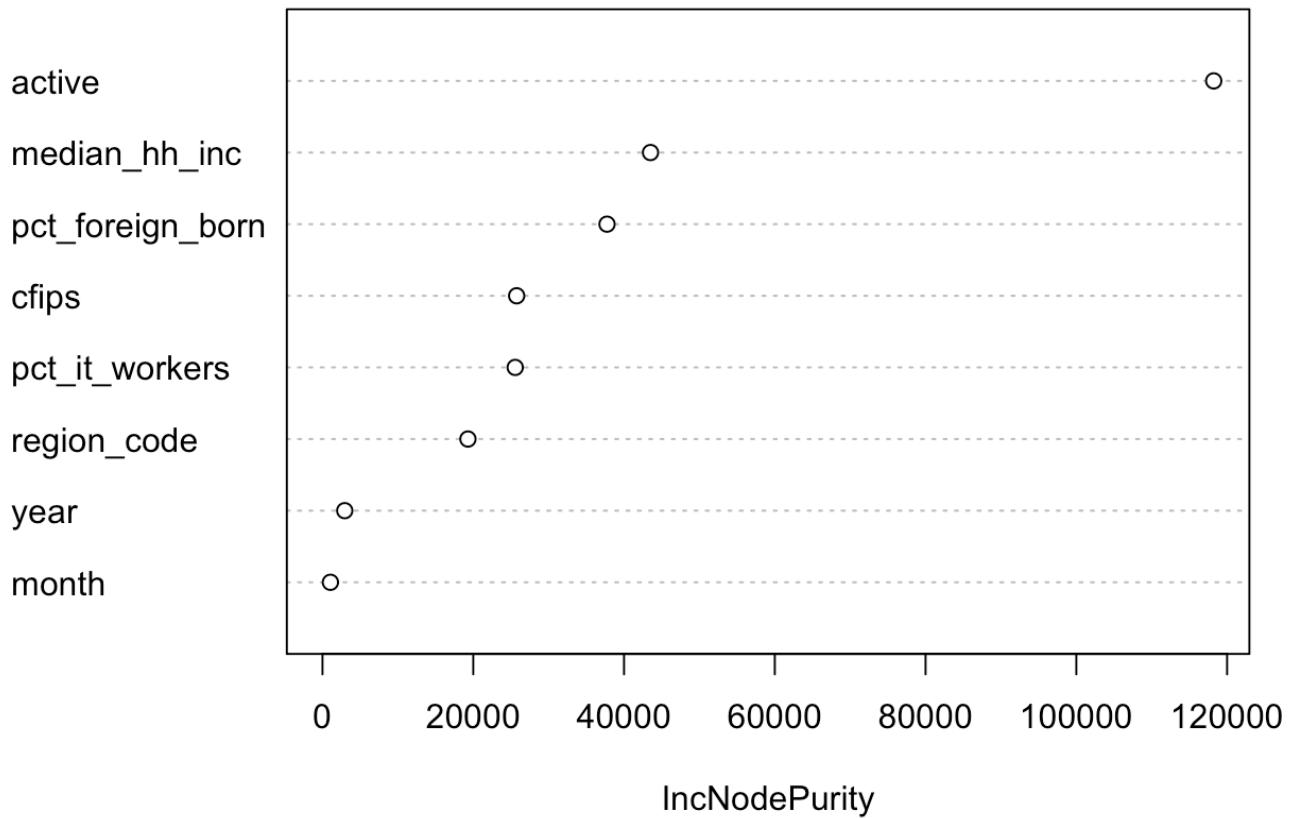
# Make predictions on the test set
predictions <- predict(model, test_data)

# Check the accuracy of the model
# accuracy <- sum(predictions == test_data$microbusiness_density) / nrow(test_data)
# cat("Accuracy:", accuracy, "\n")

# Plot the variable importance
varImpPlot(model, main = "Variable Importance Plot")

```

Variable Importance Plot



```

# Load revealed_test.csv into a dataframe
revealed_test_df <- read.csv("./datasets/revealed_test.csv")

# Change first_day_of_month format in "revealed_test_df" to Date
revealed_test_df$first_day_of_month <- as.Date(revealed_test_df$first_day_of_month)

```

4. Time Series Forecasting

4.1 Exponential Forecasting

Forecasts produced using exponential smoothing methods are weighted averages of past observations, with the weights decaying exponentially as the observations get older. In other words, the more recent the observation the higher the associated weight. This framework generates reliable forecasts quickly and for a wide range of time series, which is a great advantage and of major importance to applications in industry.

4.1.1. Simple Exponential Smoothing Forecast

Exponential smoothing is a general technique for smoothing time series data by giving more weight to recent observations. The simplest of the exponentially smoothing methods is naturally called simple exponential smoothing (SES). This method is suitable for forecasting data with no clear trend or seasonal pattern. (There is a grow in the last few years, which might suggest a trend. We will consider whether a trended method would be better for this series later.)

```

# Select relevant predictor columns
# predictors <- merged_df_new[, c("microbusiness_density", "cfips")]

merged_df_ts <- merged_df_clean %>%
  group_by(first_day_of_month) %>%
  summarise(microbusiness_density = mean(microbusiness_density))

# Convert data to time series format
ts_data <- ts(merged_df_ts$microbusiness_density, start = c(2019, 8), frequency =
12)

```

```

# fit a simple exponential smoothing model to the time series data
fit <- ets(ts_data, model = "ANN")

```

```

# generate forecasts for the time series data using the fitted model
forecast_values <- forecast::forecast(fit, h = 8)

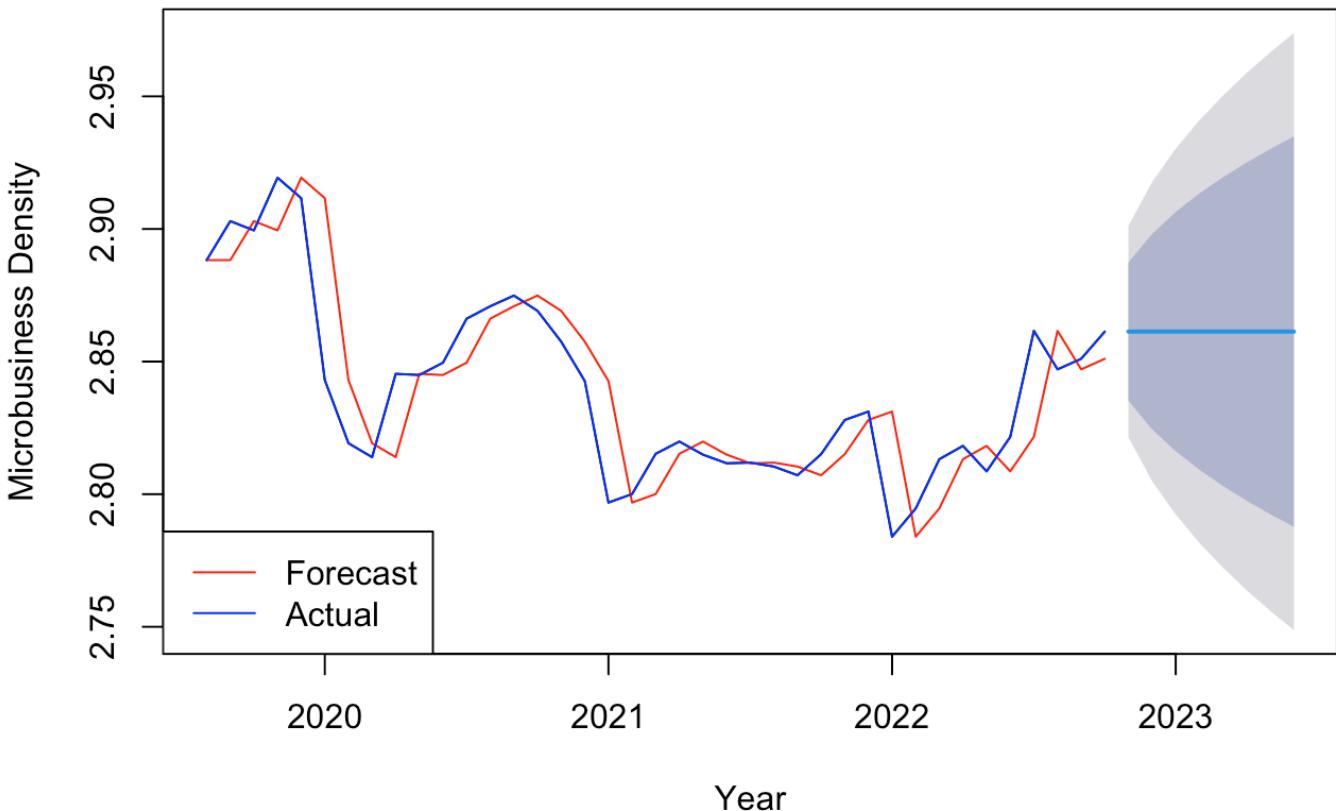
```

```

plot(forecast_values, main = "Simple Exponential Smoothing Forecast", xlab = "Year",
", ylab = "Microbusiness Density")
lines(forecast_values$fitted, col = "red")
lines(ts_data, col = "blue")
legend("bottomleft", legend = c("Forecast", "Actual"), col = c("red", "blue"), lty
= 1)

```

Simple Exponential Smoothing Forecast



The forecasts for the period 11/2022 to 06/2023 are plotted in above. Also, plotted are one-step-ahead fitted values alongside the data over the period 08/2019 to 10/2022.

```
# Split data into training and testing sets
train_data <- window(ts_data, end = c(2022, 2))
test_data <- window(ts_data, start = c(2022, 3))

# Train neural network autoregression model
nnetar_model <- nnetar(ts_data)

# Make forecasts using the trained model
forecast_data <- forecast::forecast(nnetar_model, h = length(test_data))

# Convert data to time series format
ts_data <- ts(merged_df_ts$microbusiness_density, start = c(2019, 8), frequency =
12)

# Split data into training and testing sets
train_data <- window(ts_data, end = c(2022, 2))
test_data <- window(ts_data, start = c(2022, 3))

# Fit a neural network autoregression model with error("A"), trend("N"), season("N")
# preferences
nnetar_model <- nnetar(train_data, lambda = 0, P = 12, size = 10, repeats = 10,
decay = 0.5, maxit = 1000, trace = TRUE,
entropy = FALSE, parallel = FALSE, xreg = NULL)
```

```
## # weights: 41
## initial value 21.283645
## iter 10 value 11.004232
## iter 20 value 10.950528
## iter 30 value 10.940432
## iter 40 value 10.939499
## final value 10.939493
## converged
## # weights: 41
## initial value 59.692453
## iter 10 value 11.964849
## iter 20 value 10.974613
## iter 30 value 10.949670
## iter 40 value 10.945596
## iter 50 value 10.941198
## iter 60 value 10.939534
## final value 10.939493
## converged
## # weights: 41
## initial value 70.392818
## iter 10 value 11.771520
## iter 20 value 10.980551
## iter 30 value 10.943868
## iter 40 value 10.939699
## iter 50 value 10.939518
## final value 10.939493
## converged
## # weights: 41
## initial value 26.849765
## iter 10 value 10.989322
## iter 20 value 10.955714
## iter 30 value 10.951777
## iter 40 value 10.951562
## final value 10.951541
## converged
## # weights: 41
## initial value 32.025911
## iter 10 value 11.050677
## iter 20 value 10.979900
## iter 30 value 10.942657
## iter 40 value 10.939626
## iter 50 value 10.939493
## final value 10.939493
## converged
## # weights: 41
## initial value 42.111868
## iter 10 value 11.083042
## iter 20 value 10.953153
## iter 30 value 10.951628
## iter 40 value 10.951561
## iter 50 value 10.951540
## iter 50 value 10.951540
## iter 50 value 10.951540
```

```

## final value 10.951540
## converged
## # weights: 41
## initial value 51.452718
## iter 10 value 13.681090
## iter 20 value 10.979947
## iter 30 value 10.958008
## iter 40 value 10.954194
## iter 50 value 10.952654
## iter 60 value 10.951666
## iter 70 value 10.951540
## iter 70 value 10.951540
## iter 70 value 10.951540
## final value 10.951540
## converged
## # weights: 41
## initial value 38.123377
## iter 10 value 11.066611
## iter 20 value 10.957299
## iter 30 value 10.951926
## iter 40 value 10.951545
## final value 10.951540
## converged
## # weights: 41
## initial value 33.945638
## iter 10 value 11.005029
## iter 20 value 10.957049
## iter 30 value 10.951637
## iter 40 value 10.951559
## iter 50 value 10.951540
## final value 10.951540
## converged
## # weights: 41
## initial value 53.766483
## iter 10 value 12.197989
## iter 20 value 10.958533
## iter 30 value 10.943320
## iter 40 value 10.941039
## iter 50 value 10.939557
## iter 60 value 10.939496
## final value 10.939493
## converged

```

```

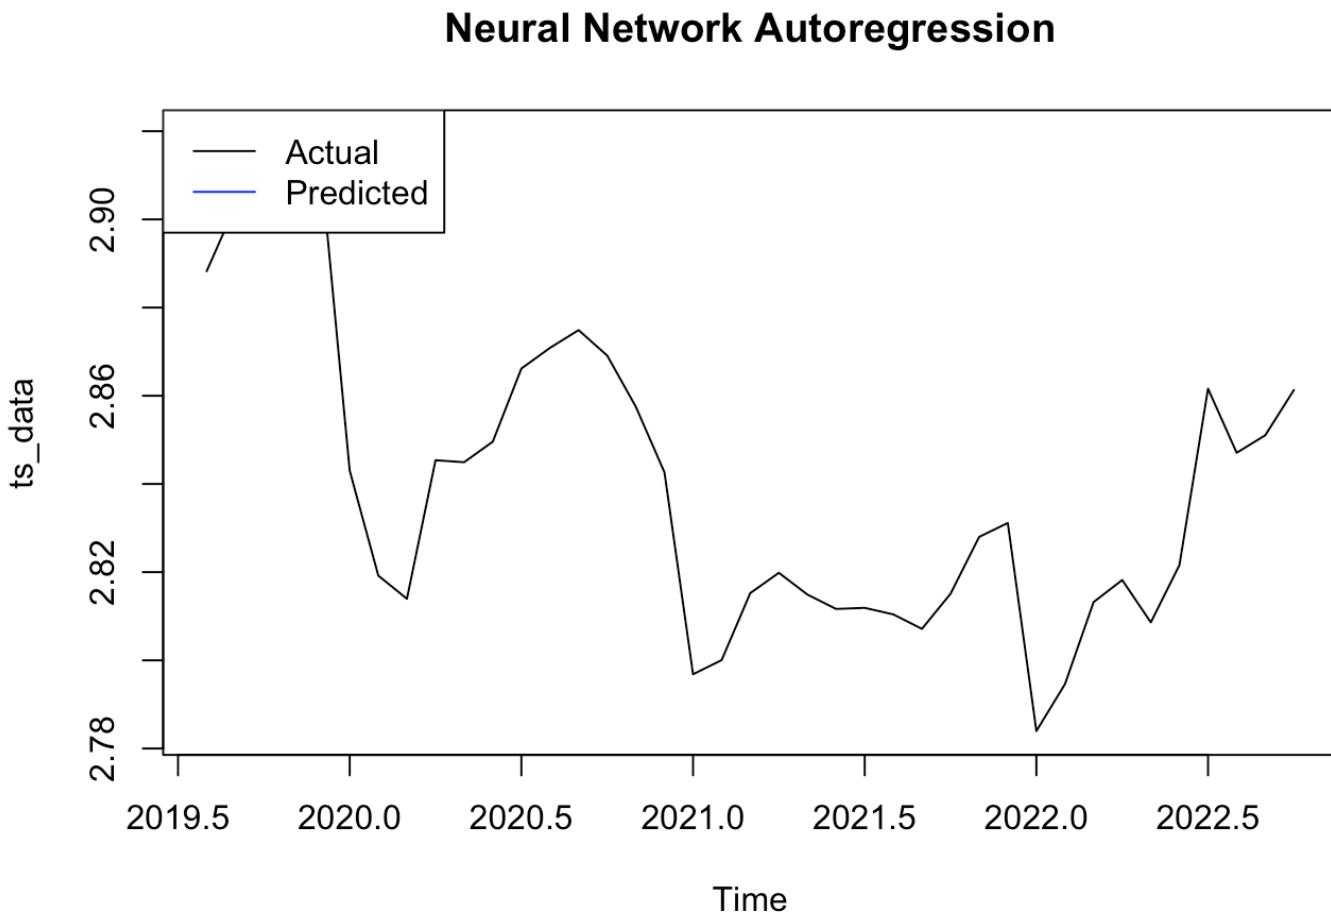
# Make predictions using the fitted model
nnetar_fc <- forecast::forecast(nnetar_model, h = length(test_data))

# Print the forecasted values
print(nnetar_fc$mean)

```

	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct
## 2022	2.804899	2.811903	2.816620	2.819850	2.822088	2.823652	2.824751	2.825525

```
# Plot actual and predicted values for comparison
plot(ts_data, type = "l", main = "Neural Network Autoregression")
#lines(nnetar_fc$mean, col = "blue")
legend("topleft", legend = c("Actual", "Predicted"), col = c("black", "blue"), lty = 1)
```



```
# Calculate accuracy measures
forecast::accuracy(nnetar_fc$mean, test_data)
```

```
##               ME      RMSE      MAE      MPE      MAPE      ACF1
## Test set 0.01666848 0.02306722 0.0186722 0.5839484 0.6552904 0.3798398
##           Theil's U
## Test set 1.367118
```

```
# Calculate forecast errors
nnetar_errors <- test_data - nnetar_fc$mean
mean(nnetar_errors)
```

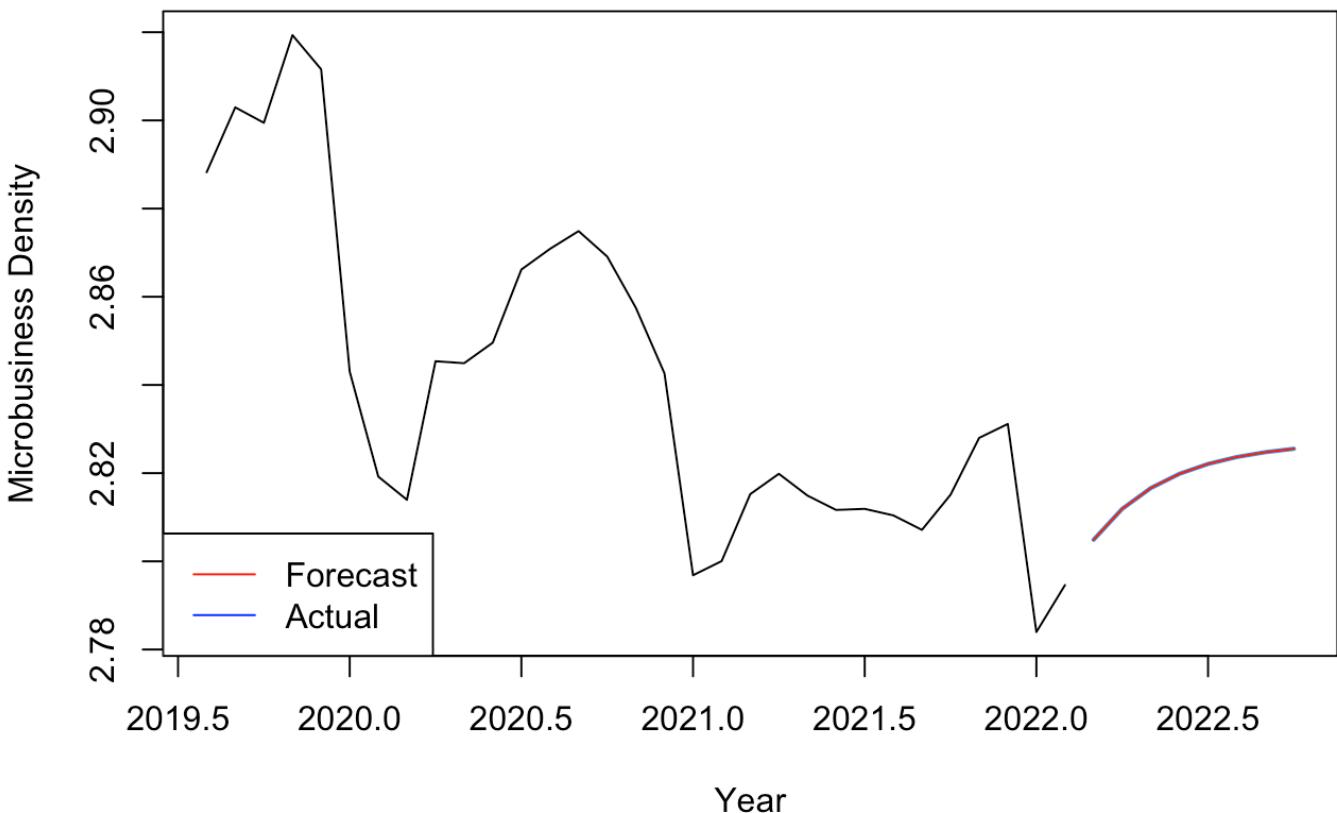
```
## [1] 0.01666848
```

```
sd(nnetar_errors)
```

```
## [1] 0.01704644
```

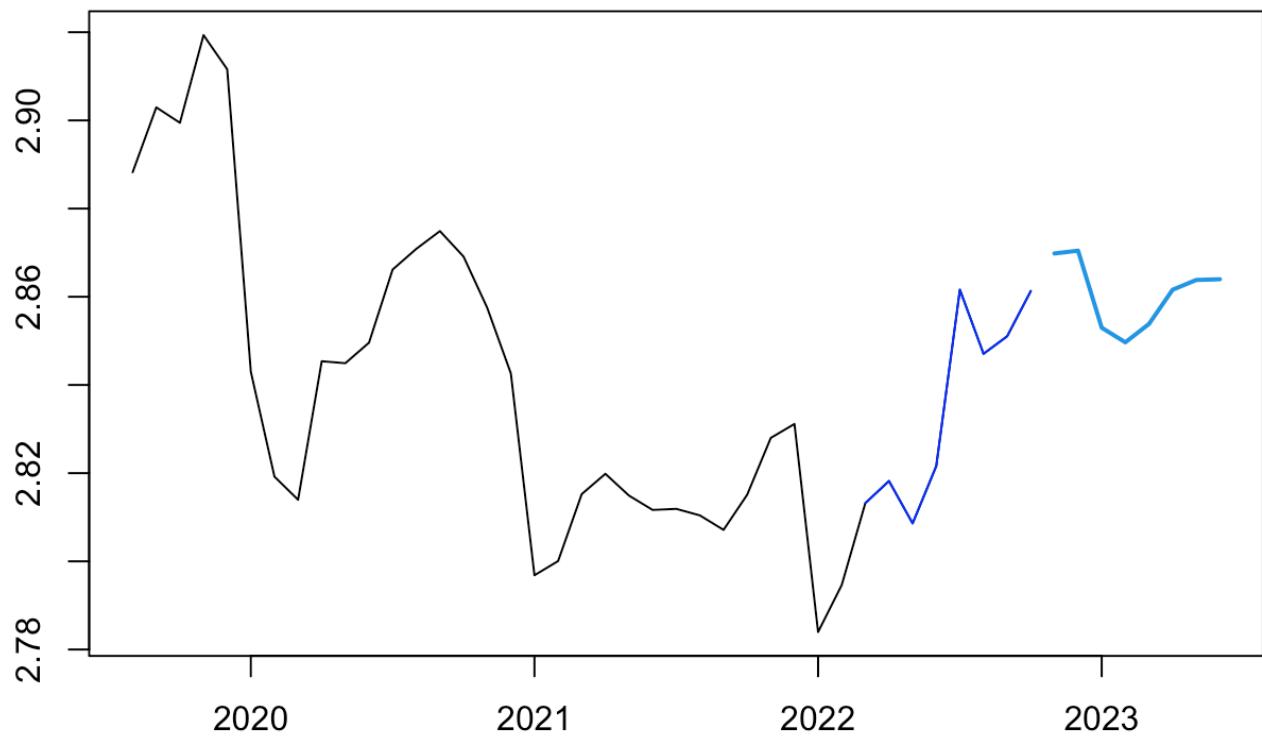
```
plot(nnetar_fc, main = "Neural network autoregression Forecast", xlab = "Year", yl-
ab = "Microbusiness Density")
lines(nnetar_fc$mean, col = "red")
#lines(ts_data, col = "blue")
legend("bottomleft", legend = c("Forecast", "Actual"), col = c("red", "blue"), lty =
= 1)
```

Neural network autoregression Forecast



```
# Plot forecasts and actual values
plot(forecast_data, main = "Neural Network Autoregression Forecasts")
lines(test_data, col = "blue")
```

Neural Network Autoregression Forecasts



-
1. <https://doi.org/10.1371/journal.pone.0256407.g001>
(<https://doi.org/10.1371/journal.pone.0256407.g001>) ↵